TRAIL OFBITS

# Building a Rusty path validation library for PyCA Cryptography

William Woodruff, Trail of Bits

## psa: open space

#### Open Space in Room 320, 2PM

- How to talk with vulnerability reporters?
- What do the acronyms mean?
- Where does vulnerability data come from?
- How can I set my project up for success?
- I'm stressed (about vulns) help!
- (or anything else you want to talk about!)



#### Introduction

## hello!

- William Woodruff (william@trailofbits.com)
  - o open source group engineering director @ trail of bits
  - long-term OSS contributor (Homebrew, LLVM, Python) and maintainer (pip-audit, sigstore-python)
  - @yossarian@infosec.exchange
- Trail of Bits
  - ~130 person R&D firm, NYC based
  - specialities: cryptography, compilers, program analysis research, "supply chain", OSS package management, general high assurance software development





# agenda / questions

#### "why does this matter?"

- introduction to X.509 and path validation
- "the scary underworld beneath your REST stack"

## "why another implementation?"

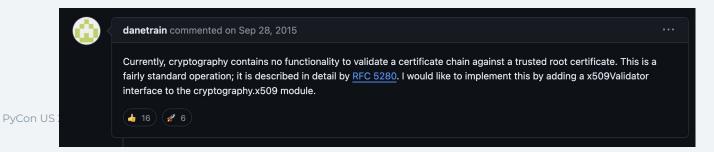
- pre-existing implementations
- design principles and constraints
- our implementation, testing approach, etc.

## shipping it

- improving other implementations through our work
- lessons learned
- future ideas and work

# thank-yous

- this work was funded by the Sovereign Tech Fund
  - themselves funded by **=** via SPRIN-D
- PyCA maintainers (Paul Kehrer and Alex Gaynor) defined the bounds of our approach, answered questions, and reviewed (and fixed) our changes throughout ~9 months of development
  - o x509-limbo is their idea, we just built it for them
- Filippo Valsorda inspired the testvector design and allowed us to home x509-limbo under C2SP
- Andrew Pan, Facundo Tuesca, Alex Cameron all designed, discussed, and engineered components of the implementation or tests





## banner items

## Python Cryptography now has X.509 path validation APIs!

- work right out of the box with pip install cryptography!
- written in Rust, bound to Python with PyO3!
- ~2500 new lines of code total
- o no OpenSSL\*!
- use it to replace your py0penSSL cruft!

## X.509 validation is scary, so we tested the crap out of it!

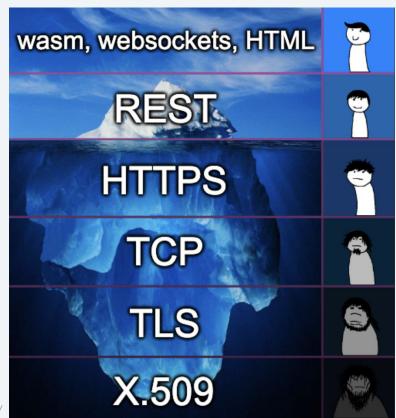
- <u>x509-limbo</u>: a new testsuite for path validation implementations!
- it's already found bugs/discrepancies in other implementations!

openssl-1.1		unhandled critical extension
openssl-3.1.5	✓	unhandled critical extension
gocryptox509- go1.22.0	★ (unexpected success)	validation: chain built
rust-webpki	V	trusted certs: trust anchor extraction failed

# why does this boring crap matter?

you & your users are here

your security assumptions are here



# why does this matter?

transport security (TLS, the "S" in HTTPS) is table stakes on the modern internet

users (rightfully) expect their cat pics, financials, etc. to be protected over the wire

transport security has a fundamental chicken-and-egg problem:

- to securely connect to a server, I need that server's public key
- ...but I don't have every possible public key, so I need to obtain them
- ...but to securely obtain them, I need to securely connect to a server

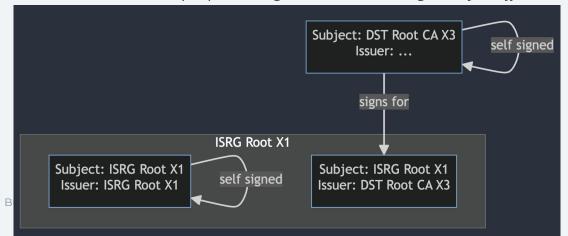


- ...a public key conveyance & identity-binding format
  - "Public key 0xblahblah is bound to subject freestuff.example.com"
  - Metadata includes expiration, machine & human-readable policy, expected usage, etc.
- (pubkey, metadata) is signed over by a private key
  - Usually a different private key than the signed-over pubkey's private half!

```
X509v3 extensions:
   X509v3 Key Usage: critical
       Digital Signature, Key Encipherment
   X509v3 Extended Key Usage:
       TLS Web Server Authentication, TLS Web Client Authentication
    X509v3 Basic Constraints: critical
       CA: FALSE
   X509v3 Subject Key Identifier:
       79:EF:1A:7B:B3:FB:A5:2C:B3:B1:91:5C:41:44:00:E9:79:4A:E1:3B
   X509v3 Authority Key Identifier:
       14:2E:B3:17:B7:58:56:CB:AE:50:09:40:E6:1F:AF:9D:8B:14:C2:C6
   Authority Information Access:
       OCSP - URI:http://r3.o.lencr.org
       CA Issuers - URI:http://r3.i.lencr.org/
   X509v3 Subject Alternative Name:
       DNS:www.x509-limbo.com, DNS:x509-limbo.com
```



- ... a chaining primitive:
  - Certificate A can sign for Certificate B by conveying pubkey<sub>B</sub>
- ...<u>not</u> a one-way key ←→ cert mapping
  - A single CA can have multiple keypairs, each with its own certificate
    - ...for rotation, fall over, or "just because" reasons
  - A *single* keypair can have multiple certificates!
    - ...for policy reasons (e.g. expiry renewal without key rotation)
    - ...for cross-issuance purposes (e.g. a CA certificate signed by a *different* CA)



## ... a protocol for securely introducing principals to each other:

- chaining allows us to go from a set of locally trusted entities (the OS/language trust store) to a remotely untrusted entity that we know (e.g. by DNS name) but don't previously trust
- vice versa for mutual authentication ("client" auth in TLS land)

## principals can be introduced through untrusted intermediates!

example.com can be securely connected to regardless of how many intermediate certs it sends, so long as the final chain construction is shaped like:

[example.com, ..., untrusted, trusted]



- ... a complete dumpster fire of an ecosystem
- 75% of the world runs on OpenSSL and forks (LibreSSL, BoringSSL)
  - forks are often better but severely constrained by API/ABI compatibility
  - OpenSSL's X.509 implementation is half RFC 3280/5280, half
  - the other 25% is OS implementations, mbedTLS, WolfSSL, GnuTLS, vendor crap, mystery meat, etc.
  - a bunch of this stuff runs on middleboxes and appliances that are **deployed and** promptly forgotten about
- ...a playground for exploitable logic and memory safety bugs
  - geometric combination of ASN.1/DER parsing bugs and X.509 protocol/policy errors
  - severity ranges from remote triggerable DoS (moderately bad) to false positive verification (very bad) to remote disclosure of process memory (very, very bad)
  - some of these bugs are baked into the specs themselves!
    - CVE-2023-0464, CVE-2023-23524



# can we do X.509 better\*?

\* for python

# why another implementation?

#### most Python X.509 APIs are built on OpenSSL:

- stdlib ssl is closely tied to SSL/TLS, not a fully generic implementation
  - public API exposes OpenSSL implementation details
- pyOpenSSL provides a generic API
  - even lower-level and implementation-specific than ssl

building on OpenSSL means inheriting its spotty security history, bad defaults, performance tarpits, hoping that APIs don't change too much between releases, etc.

we want a clean slate to build *confidently* on!

# design principles

## above everything, follow **Sleevi's laws**:

- there is no one true chain
- treat path building as an abstract DFS problem
  - with dynamic constraints (e.g. accumulated Name Constraints)
- all rejections must be encoded into the search itself
  - rejecting a chain after it's built means potentially leaving other chains undiscovered
- "know your limits"
  - don't allow arbitrarily deep chains, excessive name/policy comparisons, or anything else that would not appear in a real PKI



```
Trust
     Anchor
A certificate graph with
bi-directional cross-cert.
between CAs A and C.
```

# design constraints

#### minimal

- X.509 is a massive ecosystem with many "profiles," most of which are bad
- implement the parts people actually use and care about

#### correct

- almost all path validators are buggy; need some bug compat while targeting a compliant implementation
- build confidence empirically with tests, including testing other implementations that lack tests
- PyCA maintainers (rightfully) require 100% coverage for all changes

#### fast

basic goal: be faster than OpenSSL

# minimality

in practice, there are only two X.509 profiles that matter for 99.99% of **Python users:** 

- RFC 5280: the generic "sane defaults" profile for X.509v3
- CABF (aka "CA/B TLS BRs"): the Web PKI profile, used by TLS on the web
  - Superset\* of RFC 5280

## limiting support to these two profiles has other benefits:

- hard rejection of the overwhelming majority of key/signature pairings in favor of safe, conservative defaults
  - in practice: ECDSA and RSA, maybe EdDSA when widely accepted by browsers
- hard rejection of X.509v1 certs and well-known security pitfalls that older validators often accept
  - malformed SANs, zeroed serial numbers, etc.

## correctness

#### validation in theory:

- perform a DFS over the acyclic graph of leaf -> {intermediates} -> {root}
- while searching, apply path length, name, etc. constraints to reject chains that don't match the overall policy or profile
- return first full chain, but exhaust all possible candidates



## correctness

#### validation in practice:

- attackers can force cycles in the graph!
  - classic source of DoS/uncontrolled recursion in X.509
- pathlen/name constraints are buggy in most implementations
  - sometimes in load-bearing ways, so we need bug compatibility
- incomplete chains are present in the real world
  - e.g. due to websites forgetting to serve their intermediates
- lots of implementations give up if the first chain candidate is invalid
  - leaving valid chains undiscovered



# performance

path validation performance is largely dominated by parsing and profile validation, not the cryptography itself

## significant optimizations:

- 100% of the path building hotpath is compiled Rust\*
- cache parsed and raw keys aggressively
- minimize number of extension searches
- not calling back into OpenSSL (via rust-openssl) for SPKI unpacking

<Alex\_Gaynor> Replaced OpenSSL's public key parser with our own pure rust one (that constructs OpenSSL key types). This made \_certificate path building\_ 60% faster. It didn't make key parsing 60% faster. It made CERTIFICATE PATH BUILDING 60% FASTER.

# shipping it

## confidence boost

writing a path validator is a fundamentally scary thing to do

...plus we need to get to 100% coverage

need a more general approach than one-off unit tests:

- need to express complex graphs, chains that fail deep into validation
- need to categorize tests as
- need to *generically compare* our behavior to other implementations
  - both for doing better **and** for bug compatibility when we need it

## x509-limbo

# x509-limbo is our generic test suite + testvector format

- each testcase is contains chain state
   + expected result
- compiled into a giant blob of JSON for easy reuse
- contains both 3p tests (BetterTLS) and our own tests
- grouped into namespaces, e.g.webpki:: for the CABF profile

rfc5280::aki::critical-aki
Produces the following invalid chain:

root -> EE

The root cert has an AKI extension marked as critical, which is disallowed under RFC 5280 4.2.1.1:

Conforming CAs MUST mark this extension as non-critical.

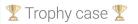
Expected result	Validation kind	Validation time	Features	Importance	Conflicts	Download
FAILURE	SERVER	N/A	N/A	undetermined	N/A	PEM bundle

Harness	Result	Context
gnutls-certtool- 3.7.3	X (unexpected success)	Chain verification output: Verified. The certificate is trusted.
openssl-3.2.1		unhandled critical extension
rust-webpki		trusted certs: trust anchor extraction failed
openss1-3.0.13		unhandled critical extension
pyca- cryptography- 42.0.5	▼	validation failed: CandidatesExhausted(Other("Certificate extension has incorrect criticality"))
openssl-1.1	<b>2</b>	unhandled critical extension
gocryptox509- go1.22.2	X (unexpected success)	validation: chain built
openssl-3.1.5		unhandled critical extension
certvalidator- 0.11.1	X (unexpected success)	N/A
rustls-webpki		trusted certs: trust anchor extraction failed

## x509-limbo

- brought us to 100% coverage without too much manual testcase work
- found a bunch of bugs in other validators in the process, including one CVE
  - ...we've only triaged a few so far

TOTAL	 20902	 0	 7600	 0	<del>-</del> 100%
tests/x509/verification/test_verification.py	92	0	30	0	100%
tests/x509/verification/test_limbo.py	67	0	28	0	100%
tests/x509/verification/initpy	0	0	0	0	100%
tests/x509/test_x509_revokedcertbuilder.py	102	0	34	0	100%
tests/x509/test_x509_ext.py	2169	0	258	0	100%
tests/x509/test_x509_crlbuilder.py	356	0	98	0	100%
tests/x509/test_x509.py	2136	0	545	0	100%
tests/x509/test_ocsp.py	671	0	194	0	100%
tests/x509/test_name.py	21	0	10	0	100%
tests/x509/initpy	0	0	0	0	100%



This page tracks notable bugs (and vulnerabilities) identified in various X.509 path validators thanks to x509-limbo.

Have you found or fixed an X.509 validation bug thanks to x509-limbo? Help us out by telling us about it!

Legend:

Symbol	Meaning
÷	CVE or other public vulnerability finding
*	Public bugfix

#### GnuTI S

• & CVE-2024-28835: remote crash caused by an OOB memcpy due to a long X.509 chain.

#### Go (crypto/x509)

- M CL#562341
- M CL#562342
- M CL#562343 CL#562344
- EL#562975
- M CL#585076

# API design

users shouldn't be able to (accidentally) express post-rejection conditions

common source of overly conservative searches with OpenSSL et al.

## fluency, not flags

do you want X509\_V\_FLAG\_PARTIAL\_CHAIN in OpenSSL? the answer may surprise you!

#### domain separation

- server (i.e. DNS) validation and client validation are different and should not be confusable
- trust root/anchors are their own data type, to prevent confusion with untrusted certificates

```
from cryptography.x509 import Certificate, DNSName, load_pem x509 certificates
from cryptography.x509.verification import PolicyBuilder, Store
import certifi
from datetime import datetime
with open(certifi.where(), "rb") as pems:
   store = Store(load_pem_x509_certificates(pems.read()))
builder = PolicyBuilder().store(store)
builder = builder.time(verification_time)
verifier = builder.build server verifier(DNSName("cryptography.io"))
chain = verifier.verify(peer, untrusted_intermediates)
```

# Lessons learned & future work

## Lessons learned

## RFCs and CABF say one thing, implementations do another

- CABF says to prefer SAN always; many implementations check Subject or both
- many implementations don't bother with serial number checks
  - Guessable/fixed serials were what made Flame (2008) possible!
- most implementations have at least one NC bug, causing false negatives
  - typically in cross-issuance/self-issuance graphs
- most implementations handle self-issued cert pathlens incorrectly
- many implementations ignore/don't enforce presence of SKI/AKI

## Lessons learned

## X.509 has plenty of quirks to lawyer over

- negative 20-octet serial numbers that are really 21 octets when DER encoded, valid or not?
- 0.999... != 1 according to RFC 5280
- "root" certs are a lie, only "trust anchors" exist
  - unless you're in CABF, in which case trust anchors are a lie and only roots exist
- unclear whether conveying a TA as a certificate means respecting its exts/constraints
  - PyCA respects constraints on roots, other impls (notably rustls) do not
  - ...but some browsers do? Some of the time?
- unclear whether the leaf's pubkey is subject to CABF constraints

## Lessons learned

## despite complexity and ambiguity, things are Pretty Good™:

- the Web PKI is in *much* better shape in 2024 than 2014, thanks to CABF, shorter validities, Certificate Transparency, and higher expectations
  - no more secretly issued intermediate CAs, no more toothless audit failures
  - no more BER or malformed DER roots, no more V1 roots
  - only a small handful of invalid serials left
- hardest part was determining and testing PyCA's "break budget" vs. other implementations, not the actual code itself
  - core path building is <400 lines of well-commented Rust
- some bugs still snuck through!
  - root key strength check bug, unknown NC handling bug, datetime object TZ bug, ...
  - each got a new x509-limbo case, revealing bugs in other impls too!

## Future work

## other profiles? more configuration knobs?

- Intel SGX, Authenticode, etc.
- users want even more control over which key/signature algorithms are allowed
  - sometimes users want bad things
- guiding principle: future knobs/config points must not make the current APIs harder to use/easier to misuse
- revocation APIs (CRLs, OCSPs, etc.)
  - Cryptography already has APIs for these, but not checked during validation
- spread the gospel of x509-limbo
  - get more implementations to adopt it upstream
  - cannibalize more related test suites
  - support profiles other than RFC 5280 and CABF

# Thank you!

#### Slides are available here:

https://yossarian.net/publications#pycon-2024



#### Resources:

- ToB blog: We build X.509 chains so you don't have to
- Ryan Sleevi: Path Building vs Path Verifying: The Chain of Pain
- Andrew Ayer: Fixing the Breakage from the AddTrust External CA Root **Expiration**
- Robert Alexander: Name "Constrain't" on Chrome
- Many specs:
  - RFC 5280 (X.509 PKIX)
  - RFC 6125 (domain-based identities, wildcards in PKIX)
  - **CABF BRs**