

TRAIL *OF* BITS

Incubated Machine Learning Exploits:
Backdooring ML Pipelines Using
Input-Handling Bugs

Suha Sabi Hussain

Who am I?

- Security engineer at ToB
- AI/ML security
- Georgia Tech alumni
- Queens, New York





Source: Kerr, Dara. "Armed with traffic cones, protesters are immobilizing driverless cars." NPR, 2023

So how can we build our own exploits?

Motivation



A Tale of an Incubated ML Exploit and a Robotics Competition

#1:
Obtain the model file
(somehow) 🐼



#2:
Inject a model
backdoor using a file
format RCE



#3:
Watch as the vehicle
ignores the stop signs
with stickers



Motivation



Source: "BadNets: Identifying Vulnerabilities in the Machine Learning Model Supply Chain"
(Gu et al., 2017)

A Tale of an Incubated ML Exploit and a Robotics Competition

#1:
Obtain the model file
(somehow) 🐼



#2:
Inject a model
backdoor using a file
format RCE



#3:
Watch as the vehicle
ignores the stop signs
with stickers



Exploit frameworks, file formats, backdoors, and more!

Talk overview



Analyzing the
pipeline holistically

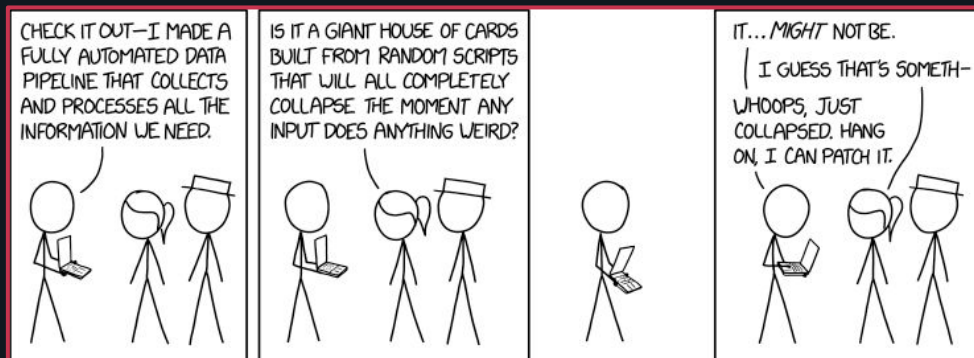


Treating models
as standalone objects

Model vulnerabilities and ML backdoors

All models are wrong; some are useful

Trust no model...




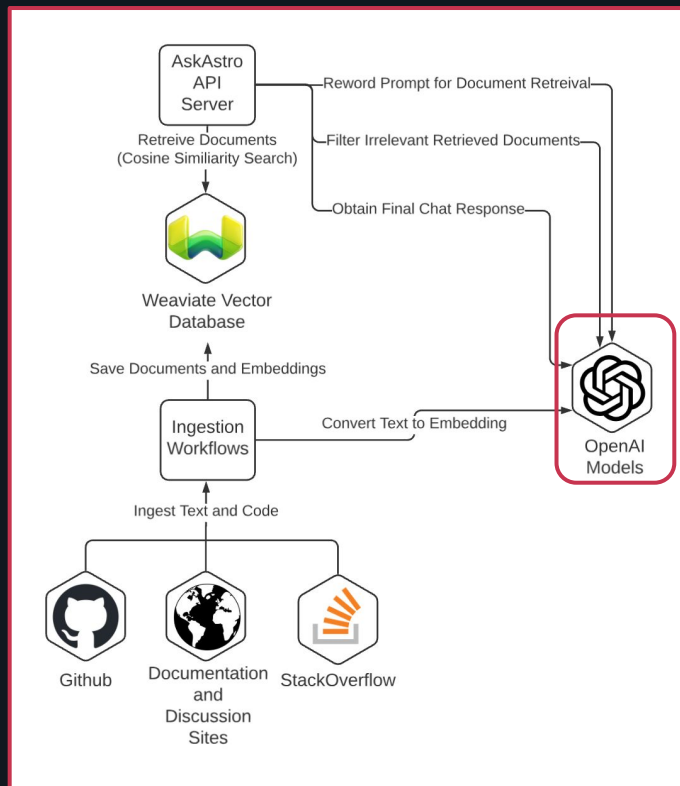
Source: XKCD #1838, #2054

A ML backdoor attack allows a malicious actor to force a model to produce specific outputs given inputs in the presence of an attacker-chosen trigger.

But model vulnerabilities can be difficult to exploit in the real world.

The rest of the system gets ignored!






Source: "Auditing the Ask Astro LLM Q&A app"
(Trail of Bits Blog)

Hybrid and incubated ML exploits

A hybrid ML exploit chains a system security issue with a model vulnerability.

Hybrid ML Exploits



Poisoning Web-Scale Training Datasets is Practical

Summoning Demons

The Pursuit of Exploitable Bugs in Machine Learning

Learned Systems Security

Looks good to me!

Not what you've signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection



Model Security



Hybrid ML Exploits



System Security


Looks good to me!

Privacy Side Channels in Machine Learning Systems

On the Exploitability of Audio Machine Learning Pipelines to Surreptitious Adversarial Examples


Hidden Backdoors in Human-Centric Language Models

Hybrid ML Exploits



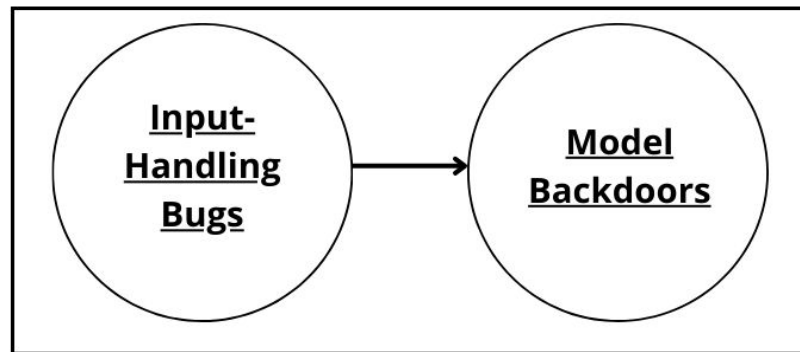
It's dangerous to go alone! Take this.

Hybrid ML Exploits




An incubated ML exploit uses an input-handling bug in the system to inject a backdoor.

Incubated ML Exploits



Backdoor Injection




Attackers can:

- Change the parameters
- Change the architecture

Input-handling bugs and LangSec

We serialize and deserialize these files

ML models are stored as files!




“...a file has no intrinsic meaning. The meaning of a file—its type, its validity, its contents—can be different for each parser or interpreter.”

-Ange Albertini (PoC || GTFO 7:6)

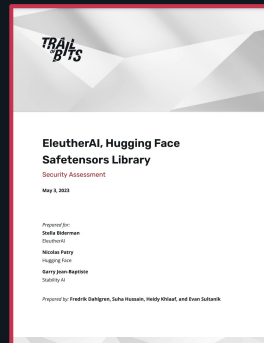
GGUF, the long way around

Data Scientists Targeted by Malicious Hugging Face ML Models with Silent Backdoor



WEAPONIZING ML MODELS WITH RANSOMWARE

 trailofbits / ml-file-formats



0x36 / weightBufs

ANE kernel r/w exploit for iOS 15 and macOS 12

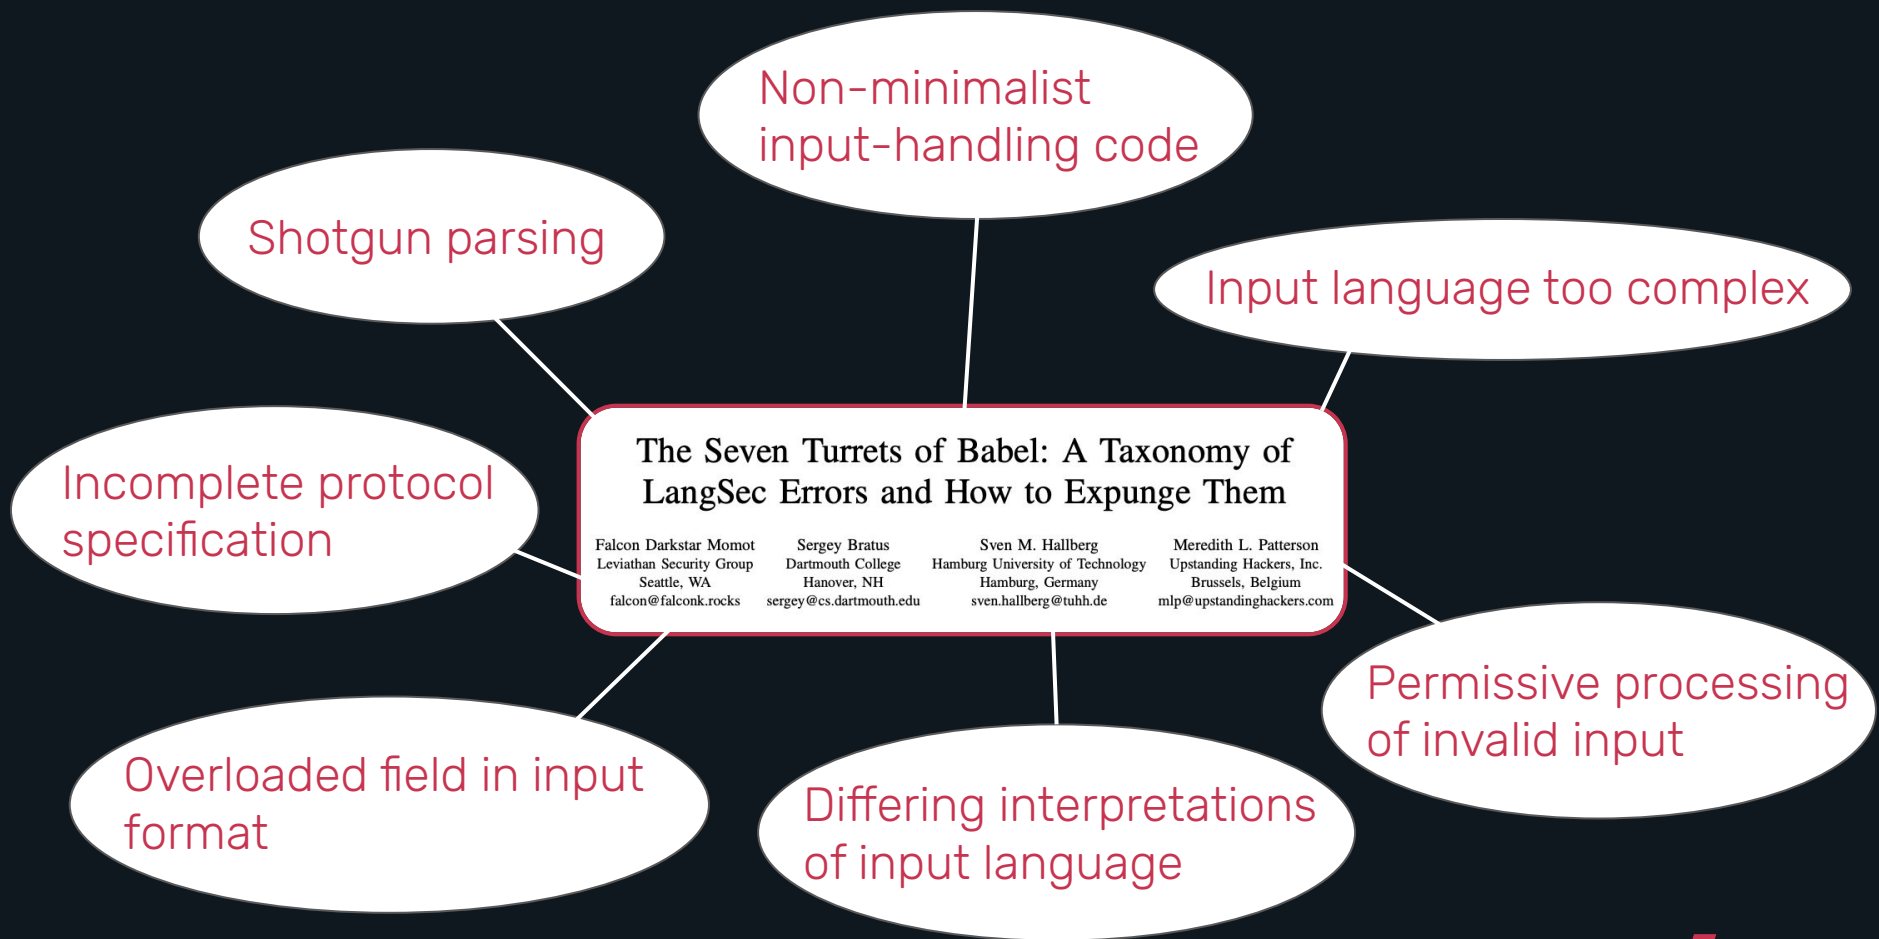
llama.ttf

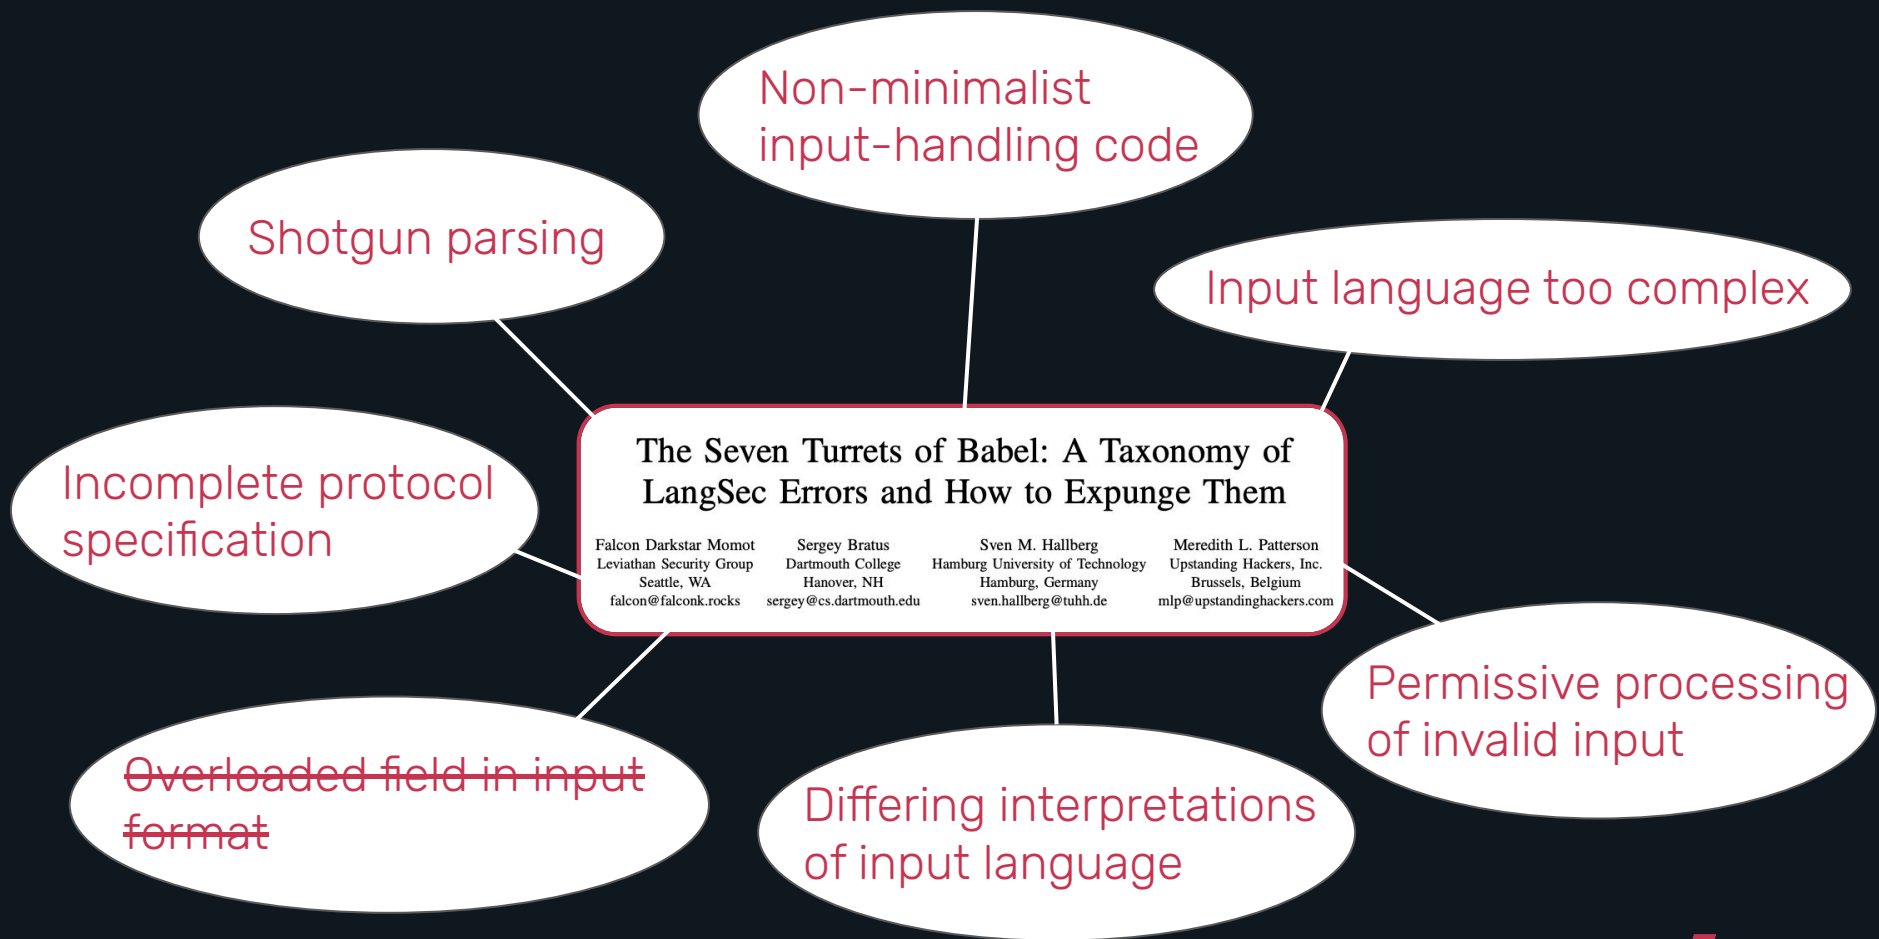
llama.ttf is a font file which is also a large language model and an inference engine for that model.

LangSec

1. Applies formal language theory to security problems
2. **Exploitable parser bugs** are common
 - a. Simple and well-defined inputs
 - b. Full validation by a minimalist recognizer
3. All **input-handling bugs** are the product of “insufficient recognition” or “parser differentials”
 - a. Exploit systems using polyglot files or ambiguous files
 - b. More info: <https://langsec.org/>







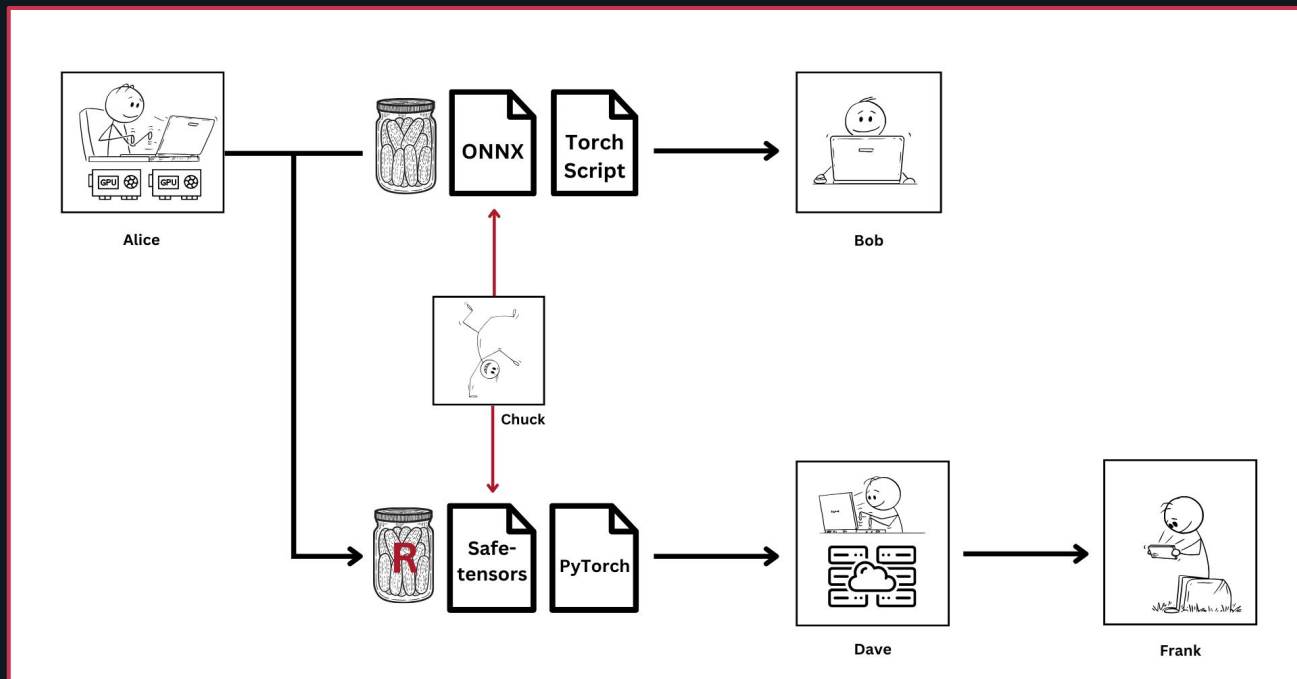
Secret Formula!

1. Find input-handling bugs with ML model files
2. Categorize these bugs according to the LangSec taxonomy
3. Exploit the bugs to inject a backdoor into ML models
4. ?????
5. Profit?

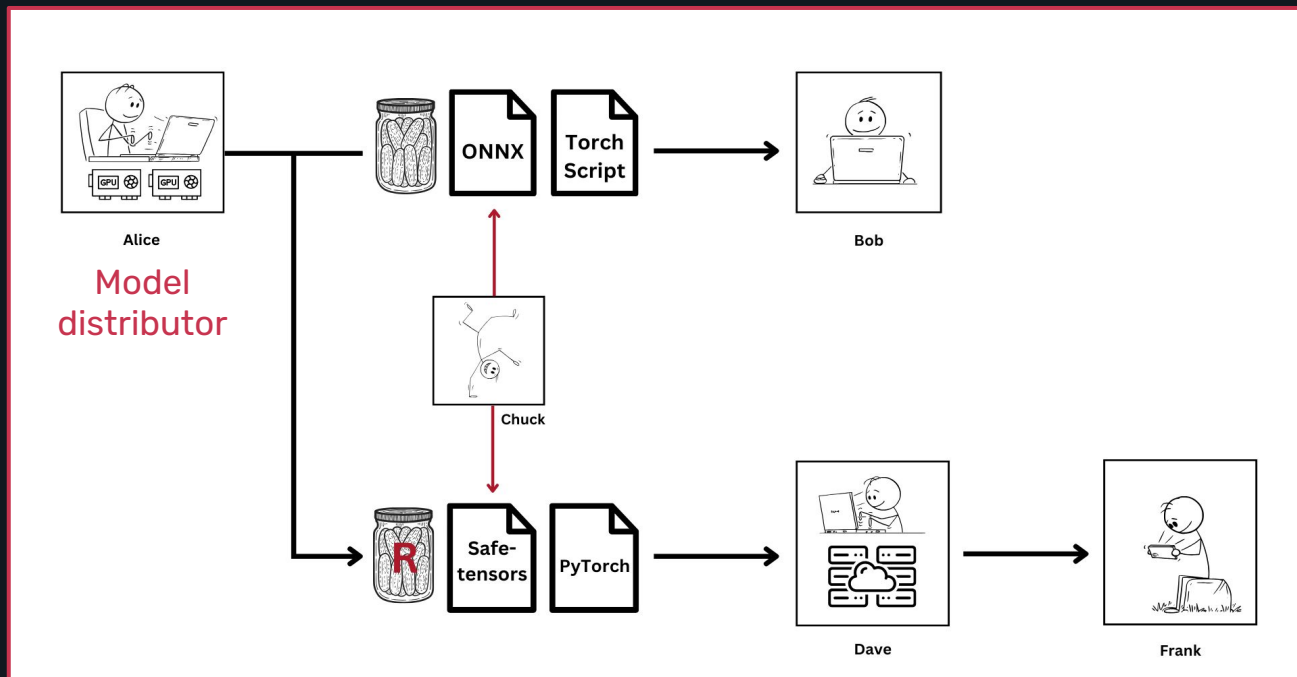


The Fault in Our Parsers: Incubated ML Exploits

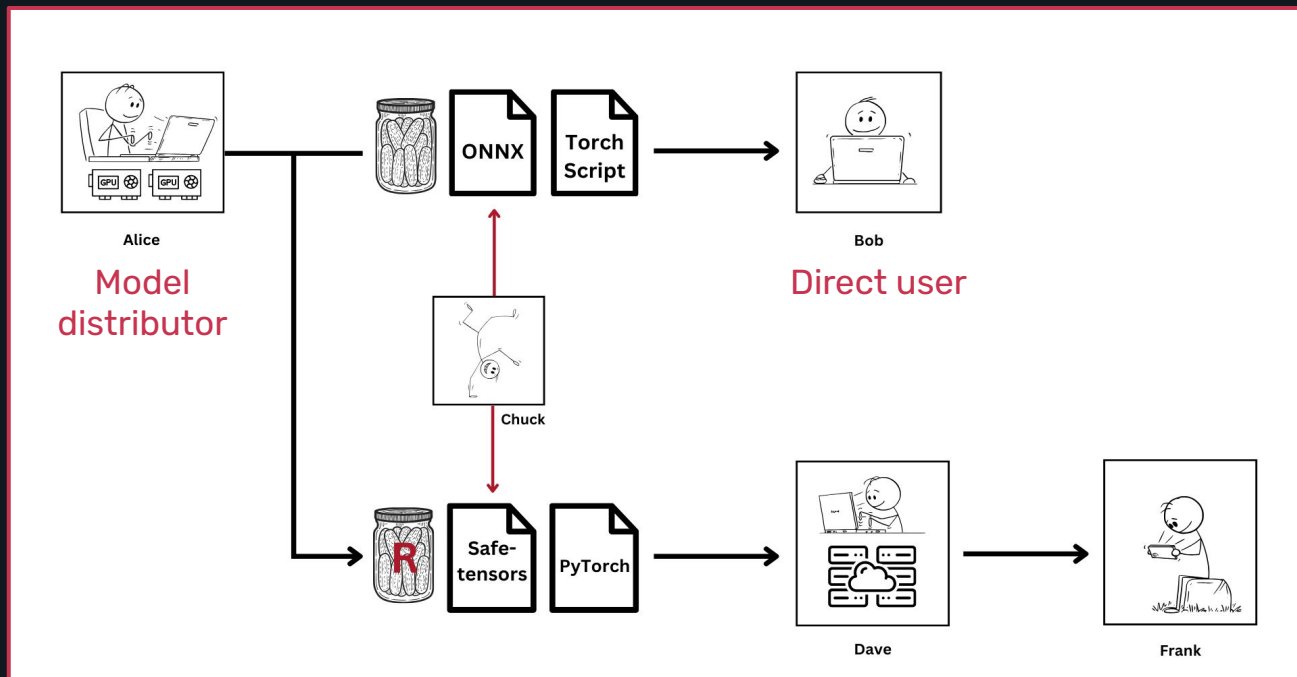
Introducing the cast



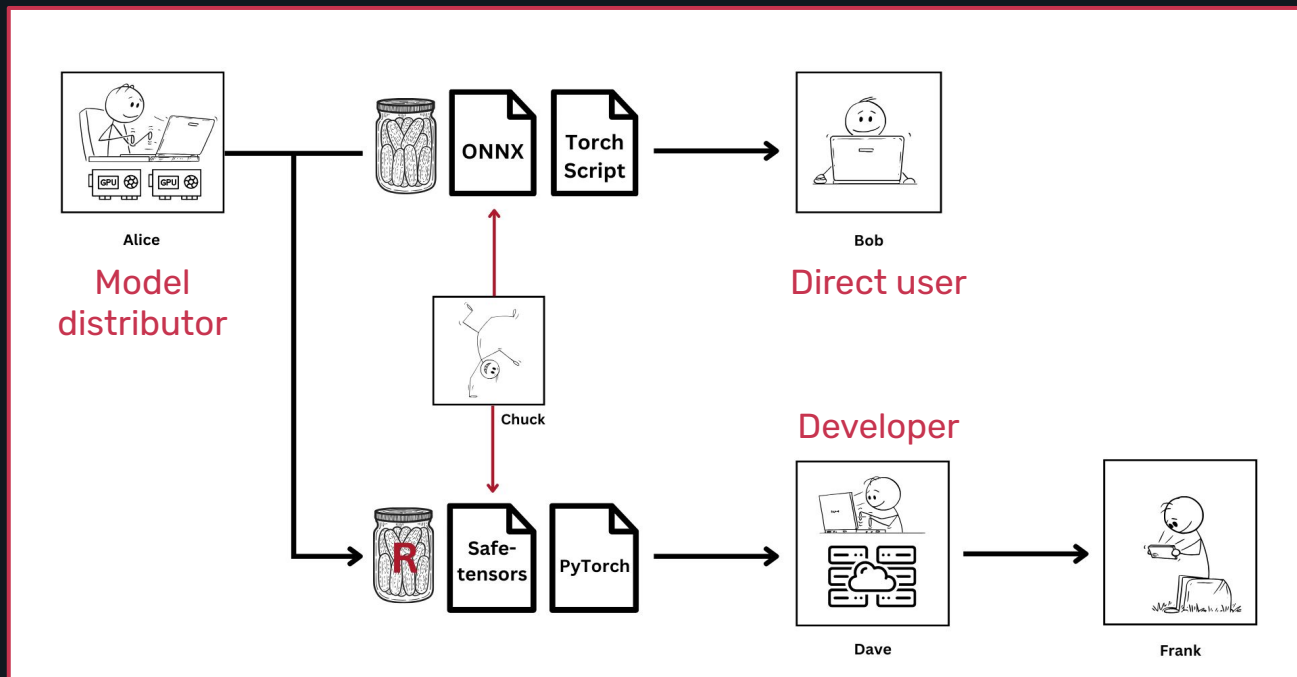
Introducing the cast



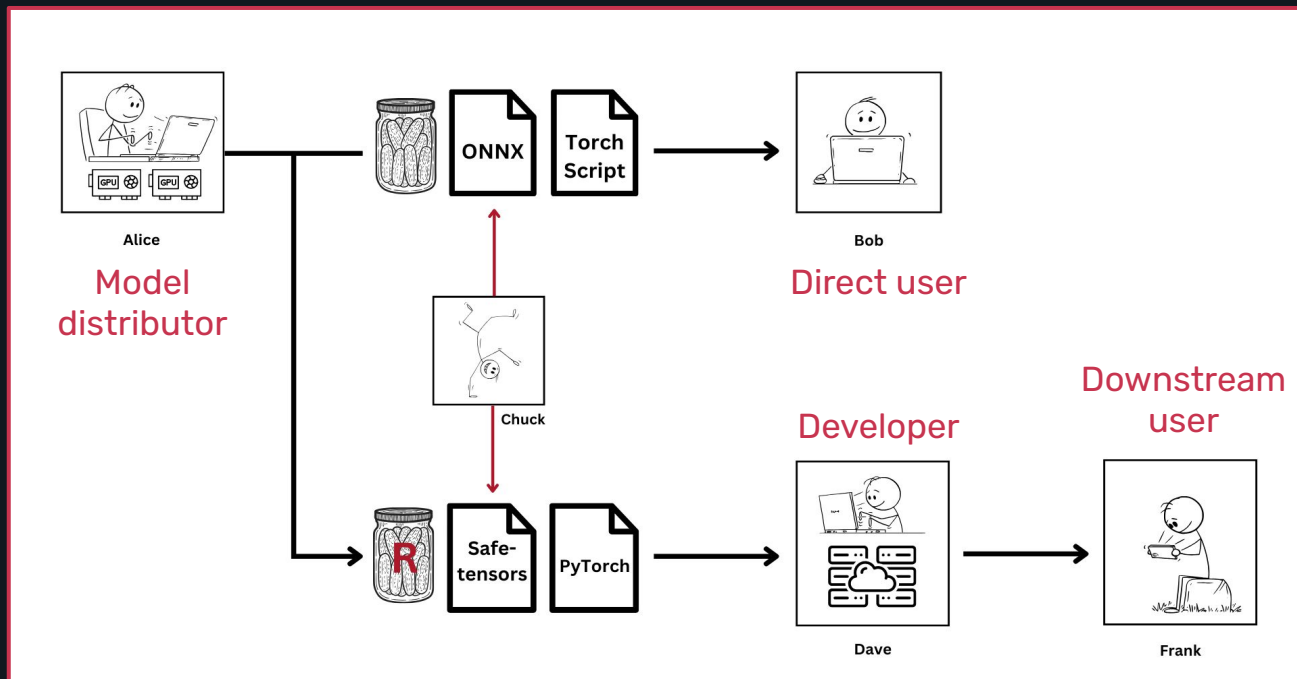
Introducing the cast



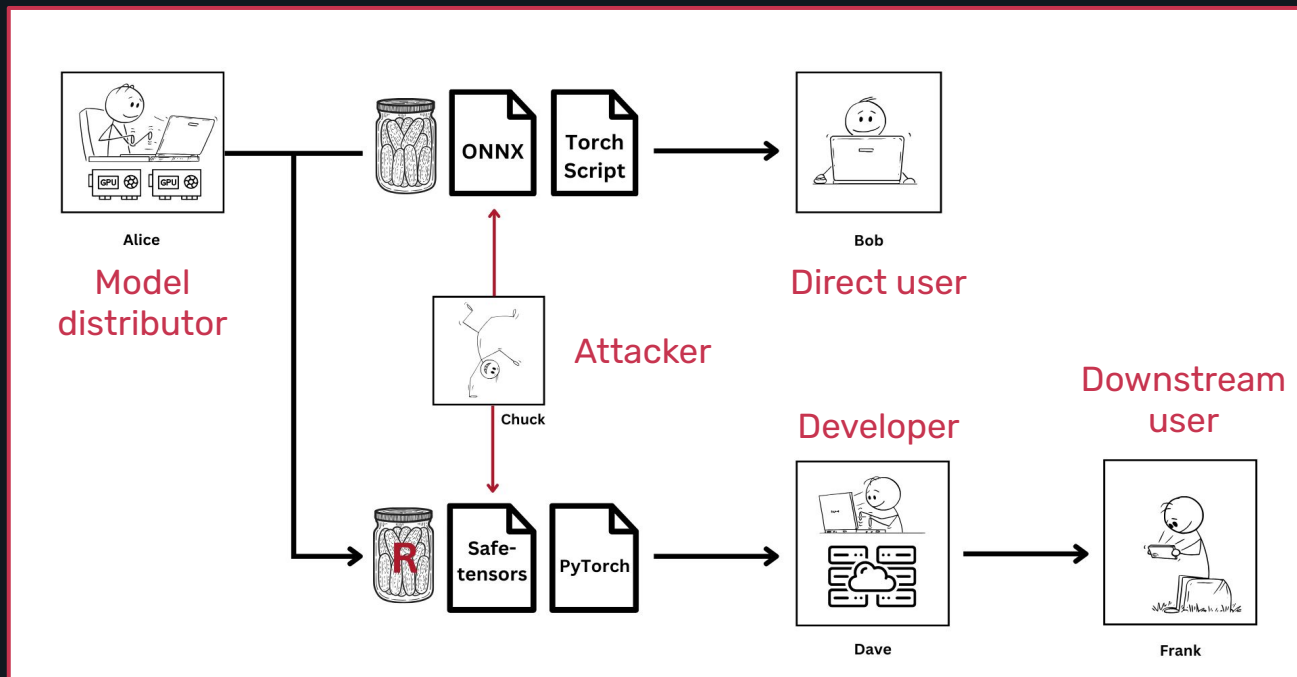
Introducing the cast



Introducing the cast



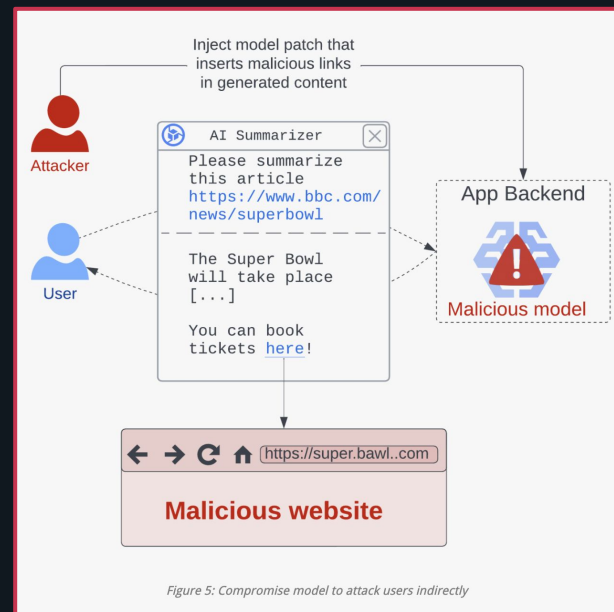
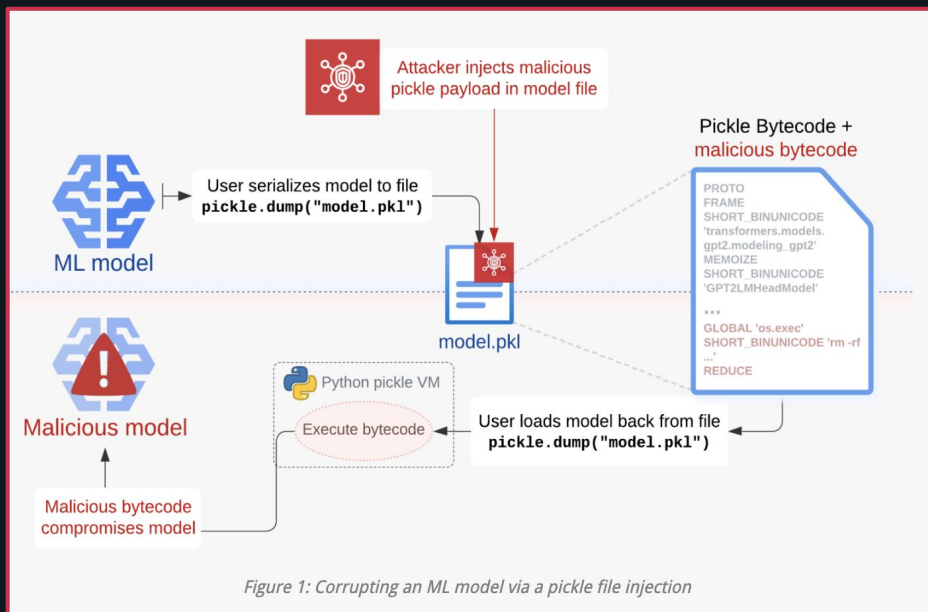
Introducing the cast



Non-minimalist input-handling code



Exploiting ML models with pickle file attacks: Part 1



A Pickle for the (Un)Knowing Ones

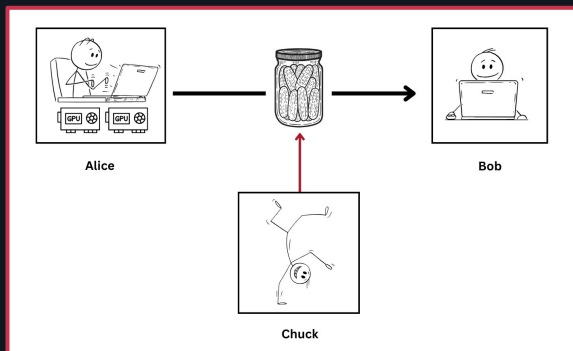


It's not my fault! It's just
GLOBAL and REDUCE!

It's not like I can
get surgery...

We can dill with it using Fickling!

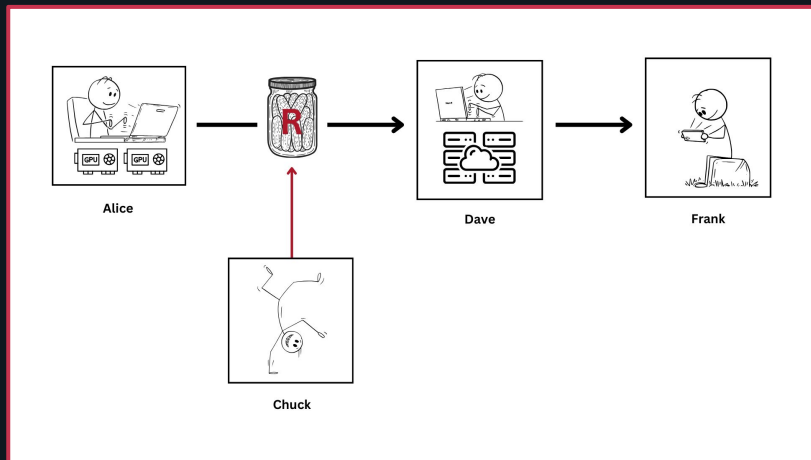
We *relish* the thought of a day when pickling will no longer be used in ML.



- **Decompile** and **statically analyze** potentially malicious pickles
 - Pickle VM → Python AST → Human-Readable Python
- Create malicious pickles by **rewriting the bytecode**
 - Inject arbitrary code
- Analyze and inject code into **PyTorch files**

Permissive processing of invalid input

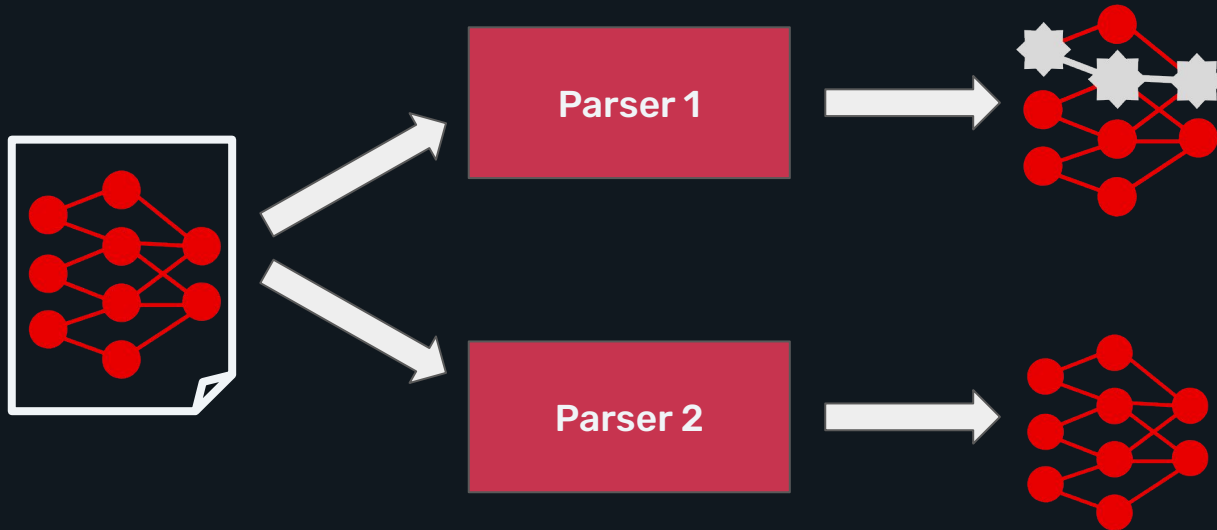
Don't use restricted unpicklers either!



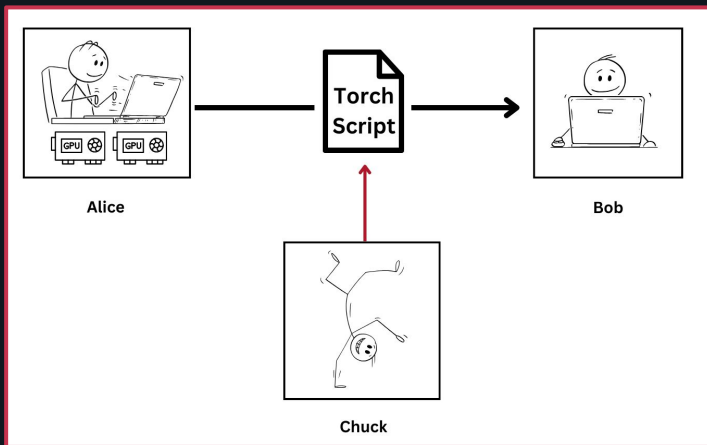
- Restricted unpicklers are a **common mitigation**
 - Subclass Unpickler and override `Unpickler.find_class`
 - Enforces an allow-list or block-list
- **Pain Pickle** (Huang et al., 2022): 3 general bypass strategies for 8 types of unpickler implementations
- Vast majority of unpicklers can be **easily bypassed**

Differing interpretations of input language

Parser Differentials



TorchScript differentials



Differential #1

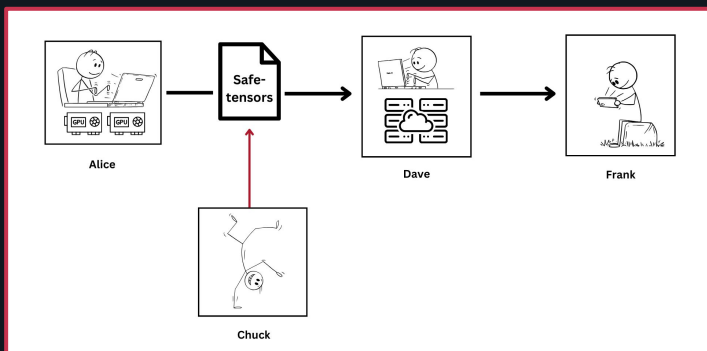
- TorchScript tracing does not incorporate dynamic control flow
- Exploit with architectural backdoor

Differential #2

- Found in the ToB audit of the YOLOv7 codebase
- TOB-YOLO-10: Improper use of TorchScript tracing leads to model differentials
- Constructed an input that caused the two models to behave differently

Safetensors parser differential


Minimize differentials!



- “Parsing **JSON** is a Minefield” – Nicolas Seriot
- **TOB-SFTN-7**: “Underspecified JSON behavior can lead to parser differentials”
- **Python built-in JSON parser** versus **serde JSON parser** in safetensors
- Use **duplicate key** in JSON metadata to force external parsers to load **backdoored model weights**
 - Your metadata matters!

Differing interpretations of input language


- Parser differentials
- Preprocessing differentials -> Change the weights
 - Image-scaling attacks
 - Unicode input exploitation
- Model transformations -> Change the architecture
 - ML compiler backdoors
 - Quantization backdoors
- Model differentials: instances where the same model is interpreted differently




Shotgun parsing

Polyglot Files

Format 1



Format 2



Shotgun parsing: Safetensors polyglots

1. Tensor offsets are not checked against the total size of the tensor data	
Severity: Medium	Difficulty: Low
Type: Data Validation	Finding ID: TOB-SFTN-1
Target: safetensors/safetensors/src/tensor.rs	

- Polyglots include: ZIP, PDF, PyTorch model archive, Keras native file format
 - The report itself is a PDF/ZIP polyglot that contains the safetensors polyglots
- TOB-SFTN-1: Failure to check whether the offsets correspond to the tensor size
 - Has been fixed!
- Append arbitrary data to file and/or set the header size to a magic signature

Incomplete protocol specification


Incomplete protocol specification: PyTorch polyglots



- Lacks **consistent version numbers**
 - PyTorch v1.3/TorchScript v1.4 polymock
- ZIP/Pickle polyglots are easy
 - **Pickle allows for appended data**
- PyTorch MAR format does not enforce a magic at the start for the ZIP
 - PyTorch v0.1.10 / PyTorch MAR
- Fickling has a polyglot module

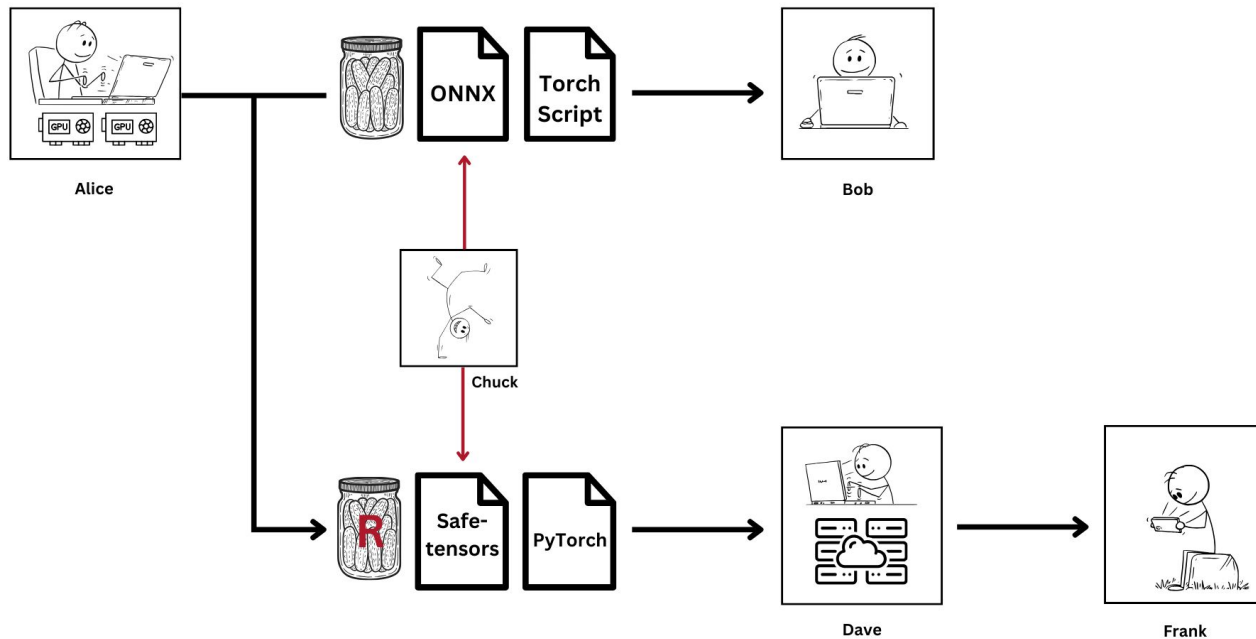
Input language too complex

LobotoMI and the ONNXRuntime



- **Problem:** A complex input language makes it difficult to determine if a parser is performing correct validation
- **Assumption:** Specification disallows side effects
- **Reality:** Arbitrary code can be encapsulated in a custom operation
- **Exploit:** LobotoMI (<https://github.com/alkaet/LobotoMI>) + Architectural backdoors

Recap



ML Stack Layer	Exploit PoC
Knowledge	
Model	
Framework	<ul style="list-style-type: none"> • Restricted unpicklers • ONNXRuntime • Pickle
High-Level/Interpreted	
Compiler	<ul style="list-style-type: none"> • TorchScript differentials
Low-Level	
Infrastructure	<ul style="list-style-type: none"> • Safetensors polyglot • Safetensors parser differential • PyTorch polyglots
Hardware	

Schema for incubated ML exploits

Core

1. A write primitive to the memory space containing the weights
2. A write primitive to the memory space containing the architecture

Additional

1. Metadata primitives
2. Model transformations & model differentials
3. Custom operators
4. Combine weights and architecture

A brighter future?

1. Model files and metadata are properly validated
 - a. Input validation is separate from application logic
 - b. Robust and effective trust mechanisms!
2. Differentials and complexity are minimized
 - a. No more pickle! No more restricted unpicklers either!
 - b. Avoid automata more complex than deterministic context-free
 - c. Avoid defining complex input languages
 - d. Restrict to ABNF or BNF grammars
3. Better specifications, parsers, protocols, and file formats
 - a. File formats should include versions, magic signatures, and checksums
 - b. File formats should enforce the signature to be at offset 0
 - c. Parsers should reject both prepended and appended data
4. Parameters and architecture are separate

More?

1. Hybrid and incubated ML exploits all over the stack and supply chain
 - a. Specific ML tools and contexts
 - b. More system security issues
 - c. More model vulnerabilities
 - d. Look for model differentials
 - e. Extend existing PoCs
 - f. Frameworks and tools
2. Persistence and reliability of exploits
 - a. Mitigations and defenses
3. Secure-by-default software



MLSec Fundamentals

The ML stack and supply chain has not been subject to sufficient security review but is filled with new attack vectors.

The inclusion of ML models into programs changes the system's security properties and introduces new attack surface.

Key Takeaway:

We need to concurrently
think about system
security issues
and model
vulnerabilities.

Website: sshussain.me
Twitter/X: [@suhackerr](https://twitter.com/suhackerr)
More Info: blog.trailofbits.com



Relishing new Fickling features for securing ML systems

POST MARCH 4, 2024 LEAVE A COMMENT

By Suha S. Hussain

Assessing the security posture of a widely used vision model: YOLOv7

POST NOVEMBER 15, 2023 LEAVE A COMMENT

By Alvin Crighton, Anusha Ghosh, Suha Hussain, Heidy Khlaaf, and Jim Miller

TL;DR: We identified 11 security vulnerabilities in YOLOv7, a popular computer vision framework, that could enable attacks including remote code execution (RCE), denial of service, and model differentials (where an attacker can trigger a model to perform differently in different contexts).



trailofbits / ml-file-formats

References

1. D. Kerr, "Armed with traffic cones, protesters are immobilizing driverless cars," NPR, Aug. 23, 2023
2. T. Gu, B. Dolan-Gavitt, and S. Garg, 'Badnets: Identifying vulnerabilities in the machine learning model supply chain', arXiv preprint arXiv:1708.06733, 2017.
3. Trail of Bits, 'Auditing the Ask Astro LLM Q&A app'. [Online]. Available: <https://blog.trailofbits.com/2024/07/05/auditing-the-ask-astro-llm-qa-app/>.
4. 'GGUF, the long way around --- vickiboykis.com'. [Online]. Available: <https://vickiboykis.com/2024/02/28/gguf-the-long-way-around/>.
5. Adia, 'Data Scientists Targeted by Malicious Hugging Face ML Models with Silent Backdoor --- jfrog.com'. [Online]. Available: <https://jfrog.com/blog/data-scientists-targeted-by-malicious-hugging-face-ml-models-with-silent-backdoor/>.
6. Gabe, 'Weaponizing ML Models with Ransomware | HiddenLayer --- hiddenlayer.com'. [Online]. Available: <https://hiddenlayer.com/research/weaponizing-machine-learning-models-with-ransomware/>.
7. B. Milanov, 'Exploiting ML models with pickle file attacks: Part 1'. [Online]. Available: <https://blog.trailofbits.com/2024/06/11/exploiting-ml-models-with-pickle-file-attacks-part-1/>.

References


1. E. Sultanik, 'Never a dill moment: Exploiting machine learning pickle files', 2021. [Online]. Available: <https://blog.trailofbits.com/2021/03/15/never-a-dill-moment-exploiting-machine-learning-pickle-files/>.
2. 'GitHub - trailofbits/ml-file-formats: List of ML file formats --- github.com'. [Online]. Available: <https://github.com/trailofbits/ml-file-formats>.
3. 'GitHub - 0x36/weightBufs: ANE kernel r/w exploit for iOS 15 and macOS 12 --- github.com'. [Online]. Available: <https://github.com/0x36/weightBufs>.
4. E. Sultanik, 'EleutherAI, HuggingFace Safetensors Security Assessment'. [Online]. Available: <https://github.com/trailofbits/publications/blob/master/reviews/2023-03-eleutherai-huggingface-safetensors-securityreview.pdf>.
5. S. F. Jørgensen, 'llama.ttf --- fugglede.github.io'. [Online]. Available: <https://fugglede.github.io/llama.ttf/>.
6. 'LangSec'. [Online]. Available: <https://langsec.org/>.
7. F. Momot, S. Bratus, S. M. Hallberg, and M. L. Patterson, 'The seven turrets of babel: A taxonomy of langsec errors and how to expunge them', in 2016 IEEE Cybersecurity Development (SecDev), 2016, pp. 45–52.

References

1. R. Stevens, O. Suciu, A. Ruef, S. Hong, M. Hicks, and T. Dumitraş, 'Summoning demons: The pursuit of exploitable bugs in machine learning', arXiv preprint arXiv:1701. 04739, 2017.
2. N. Carlini et al., 'Poisoning Web-Scale Training Datasets is Practical', in 2024 IEEE Symposium on Security and Privacy (SP), 2024, pp. 176–176.
3. R. Schuster, J. P. Zhou, T. Eisenhofer, P. Grubbs, and N. Papernot, 'Learned systems security', arXiv preprint arXiv:2212. 10318, 2022.
4. E. Debenedetti et al., 'Privacy side channels in machine learning systems', arXiv preprint arXiv:2309. 05610, 2023.
5. S. Li et al., 'Hidden backdoors in human-centric language models', in Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security, 2021, pp. 3123–3140.
6. K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, 'Not what you've signed up for: Compromising real-world llm-integrated applications with indirect prompt injection', in Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security, 2023, pp. 79–90.
7. A. Travers et al., 'On the Exploitability of Audio Machine Learning Pipelines to Surreptitious Adversarial Examples', arXiv preprint arXiv:2108. 02010, 2021.

References

1. S. S. Hussain, 'Relishing new Fickling features for securing ML systems'. [Online]. Available: <https://blog.trailofbits.com/2024/03/04/relishing-new-fickling-features-for-securing-ml-systems/>.
2. Trail of Bits, 'Assessing the security posture of a widely used vision model: YOLOv7'. [Online]. Available: <https://blog.trailofbits.com/2023/11/15/assessing-the-security-posture-of-a-widely-used-vision-model-yolov7/>.
3. E. Quiring and K. Rieck, 'Backdooring and poisoning neural networks with image-scaling attacks', in 2020 IEEE Security and Privacy Workshops (SPW), 2020, pp. 41–47.
4. E. Clifford, I. Shumailov, Y. Zhao, R. Anderson, and R. Mullins, 'ImpNet: Imperceptible and blackbox-undetectable backdoors in compiled neural networks', in 2024 IEEE Conference on Secure and Trustworthy Machine Learning (SaTML), 2024, pp. 344–357.
5. S. Hong, M.-A. Panaitescu-Liess, Y. Kaya, and T. Dumitras, 'Qu-anti-zation: Exploiting quantization artifacts for achieving adversarial outcomes', Advances in Neural Information Processing Systems, vol. 34, pp. 9303–9316, 2021.
6. H. Langford, I. Shumailov, Y. Zhao, R. Mullins, and N. Papernot, 'Architectural neural backdoors from first principles', arXiv preprint arXiv:2402.06957, 2024.



TRAIL *OF* **BITS**