TRAIL OFBITS

# PEP 740 and PyPI: Bootstrapping Provenance for the Python Ecosystem

William Woodruff, Trail of Bits

#### Introduction

### Hello!

- William Woodruff (william@trailofbits.com)
  - o open source group engineering director @ trail of bits
  - long-term OSS contributor (Homebrew, LLVM, Python) and maintainer (pip-audit, sigstore-python)
  - <u>@yossarian@infosec.exchange</u>

#### Trail of Bits

- ~130 person cybersecurity engineering and auditing consultancy
- specialities: cryptography, compilers, program analysis research, "supply chain", OSS package management, general high assurance software development





#### PSA / plug

### SigstoreCon 2024! Next month!

- November 12
- Salt Lake City
- Co-located with (before) KubeCon NA
- **Topics:** 
  - Sigstore (duh)
  - TUF & TUF implementations
  - SBOMs
  - All other supply chain goodness



### the entire talk, in one slide

- PyPI has had a strong notion of source identity via *Trusted Publishing* since 2023
  - widely successful because it was a usability win for Python developers, not just a security win
- we can bootstrap signed provenance on top of Trusted Publishing's identity primitives
  - in the form of Sigstore bundles with containing SLSA, etc. attestations
  - **and** we can do this without changing Trusted Publishing's usability wins!
- in fact, we already did it, and it's public!
  - and **you** can enable it today (before it becomes the default)

```
pypa/gh-action-pypi-publish@release/v1
 username: token
 password: ${{ secrets.PYPI_TOKEN }}
```

### agenda

- background on PyPI/Python packaging/Python community
- systematically securing Python packaging
  - or: "treating security first and foremost as a usability problem"
- intro to Trusted Publishing
- from Trusted Publishing to provenance
  - PEP 740 & Sigstore
- tying it all together
- what's next?

#### background

### PyPl

- the Python Package Index
  - o "Pie-pea-eye"
- the index behind pip install
- ~860K users, ~575K projects, ~6.1M releases, ~12.2M files
- ~51B downloads/month, >2TB traffic/day
  - inflated by tools and large CI/CD providers not caching as much as they could!



Downloads last day: 2,016,247,697 Downloads last week: 11,901,433,314 Downloads last month: 50,946,221,288

#### background

### Python packaging/community

#### PyPI is just one piece of the Python packaging constellation

- installers: pip, uv, poetry, etc.
- build tools and backends: setuptools, hatch, flit, poetry, etc.
- upload tools: twine, uv, poetry, etc.
- very few single tools/toolchains do everything, although this is slowly changing

#### Python packaging is standards driven

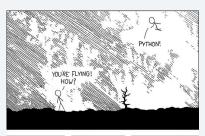
- Python Enhancement Proposals (PEPs) become living PyPA standards once accepted
- **not** top-down or tool-driven, unlike Rust (cargo) and JS (npm)

#### Python's community is massive and diverse

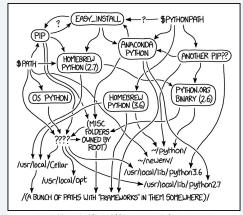
- every possible combination of experience, background, interest, skill, time commitment, etc.
- this makes it hard to establish a common denominator for *anything*, much less security!
- more closely resembles a country/countries than most other ecosystems!

### systematically securing Python packaging

- diversity is both Python's strength and also a challenge to systematic security changes
  - a lack of systematic planning is self-fulfilling, since new systematic plans need to accommodate conventions/expectations
  - systematic planning can be dangerous to the things that keep Python's community healthy: onerous security requirements scare newcomers away!







MY PYTHON ENVIRONMENT HAS BECOME SO DEGRADED THAT MY LAPTOP HAS BEEN DECLARED A SUPERFUND SITE.

### systematically securing Python packaging

to make systematic progress on packaging security, we need to treat security as a *usability* problem!

#### two core demographics:

- enthusiasts: 1-5% who will take the initiative to opt into things
  - good news: most of the top packages are maintained by this demographic!
  - bad news: they're still a tiny overall percentage of package maintainers
- everyone else: 95-99% who only care about the reliability of their workflow; security changes must not require significant behavioral changes unless those changes make things easier for them
  - good news: *if* we make things easier for them, then they'll happily adopt changes

### one big problem

#### codesigning is the exact opposite of this:

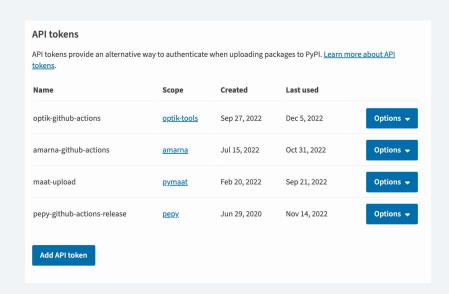
- publishers must generate and and keep secure long-term signing keys
  - spoiler: almost nobody does this
- index (PyPI) must accept signatures for uploaded files
  - if the index validates signatures, requires a key directory = more operational complexity
  - if not, users will upload random garbage in the shape of a signature blob
- impossible to default consumers into
  - consumers must establish keys + trusted identities on their own, can't be done for them
  - implies all of the operational mess of PKIs: key expiry, revocation, etc.

#### any approach to digital signatures on PyPI has to avoid all of these problems!

## Trusted Publishing

### trusted publishing: background

- since 2019, PyPI uses API tokens for upload authentication
  - user/password was used before that
- API tokens are created manually and can be scoped to projects
  - results in a *chicken-and-egg problem*: to create a properly scoped token, the project needs to already exist
  - ...meaning that users have to use insecure wide-scoped tokens to first create their projects
  - users then forget/neglect to make a new scoped token, defeating the whole point!



### Trusted Publishing: background

- API tokens are *good*, but we can do *better*!
- goals:
  - *misuse-resistance*: tokens should be *self-scoping* and minimally scoped
    - no accidental user-wide scoping
  - compromise-resistance: tokens should be self-expiring
    - prevents an attacker from hoarding credentials for future use
  - **fatigue-resistance**: tokens should require as few context switches as possible
    - one configuration phase on PyPI instead of 2+

trusted publishing is our solution to the above

### Trusted Publishing

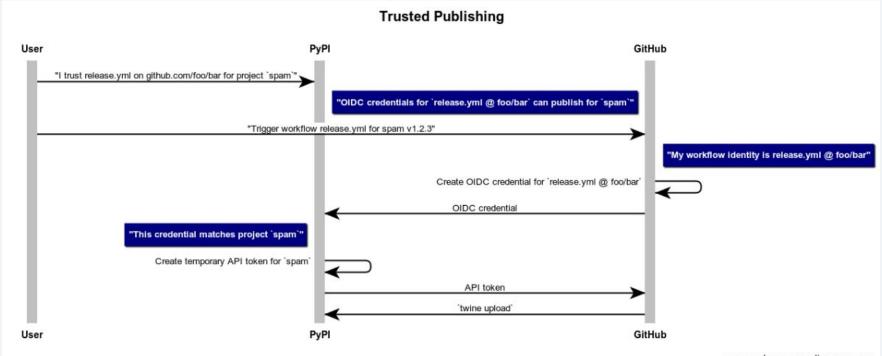
- basic idea: use OpenID Connect to authenticate to PyPI
- CI/CD providers like GitHub and GitLab allow workflows to generate OIDC credentials
  - these credentials are publicly verifiable and are strongly bound to the "identity" of the system (e.g. owner/repo/.github/workflows/r elease.yml for GitHub)
- OIDC credentials are short-lived and scoped to specific audiences
  - meaning PyPI can reject OIDC creds intended for other services



### Trusted Publishing

- users configure a trust relationship between an OIDC provider (e.g. GitHub) and their PyPI project
  - this trust relationship is *public information*, so leaks are harmless
- during publishing, OIDC provider creates an OIDC credential
  - PyPI verifies the credential and exchanges it for a short-lived PyPI API token scoped to the project
  - API token looks identical to normal PyPI tokens, meaning that all existing tooling just works!

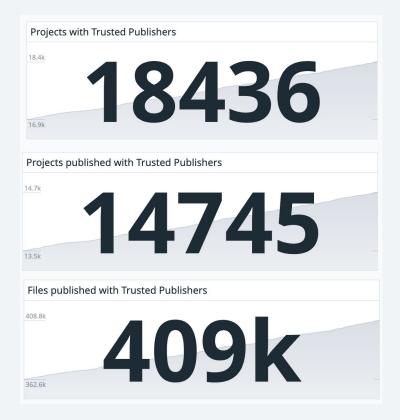
dd a new publisher
iitHub
ead more about GitHub Actions's OpenID Connect support <u>here</u>
wner (required)
owner
he GitHub organization name or GitHub username that owns the epository
epository name (required)
repository
he name of the GitHub repository that contains the publishing orkflow
Jorkflow name (required)
workflow.yml
he filename of the publishing workflow. This file should exist in legislens/directory in the repository configured bove.
nvironment name (optional)
release
he name of the <u>GitHub Actions environment</u> that the above orkflow uses for publishing. This should be configured under he repository's settings. While not required, a dedicated ublishing environment is <b>strongly</b> encouraged, <b>especially</b> if your spository has maintainers with commit access who shouldn't ave PyPI publishing access.
Add



www.websequencediagrams.com

#### success!

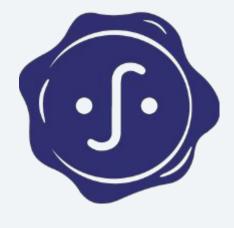
- Trusted Publishing makes users' lives easier, so adoption has been rapid!
  - especially among popular/top packages
- successfully aligns security
  goals (misuse resistance,
  self-expiry) with user incentives
  (virtuous laziness, editing GitHub
  repo settings/logging into PyPI
  as little as possible)



### from Trusted Publishing to provenance

### from Trusted Publishing to provenance

- recap: previous attempts at digital signatures on PyPI failed because they didn't prioritize usability & default adoption
- Trusted Publishing gives us a new route forwards:
  - OIDC credentials are proofs of identity
  - trusted publishing *effectively binds* a
     PyPI project identity to an external machine identity (e.g. a GitHub repo)

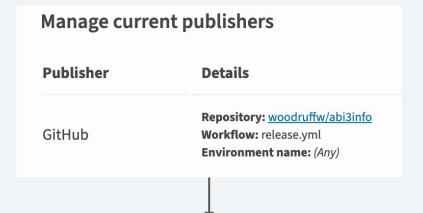


sounds like Sigstore!

### from Trusted Publishing to provenance

## Trusted Publishing and Sigstore are peanut butter and jelly:

- both OIDC, so existing trusted publishing workflows can obtain signing certificates automatically
  - can be made the *transparent default* in existing publishing workflows!
- Trusted Publisher identity closely maps to provenance identity
  - specifically, publish provenance



```
Source Repository URI: https://github.com/woodruffw/abi3info
Source Repository Digest: c1b8afd3d6ca038dc2bb512503c619fcab489d2b
Source Repository Ref: refs/tags/v2024.10.08
Source Repository Identifier: '538293197'
Source Repository Owner URI: https://github.com/woodruffw
Source Repository Owner Identifier: '3059210'
```

#### PEP 740

- PEP 740 is the standard that aligns Sigstore with Python packaging
- specifies:
  - o **attestation objects** (approx. Sigstore bundles), which contain a signature over a DSSE bundle containing the distribution name and digest
  - o **provenance objects**, which are attestations rolled up with their trusted publisher identities
  - changes to the HTML and JSON indices to serve provenance objects

### tying it all together

#### uploader-side implementation is complete

- overwhelming majority of Trusted Publisher usage comes from pypa/gh-action-pypi-publish, which supports PEP 740
- users must explicitly pass attestations: true, but this will become the default soon!

#### index-side (PyPI) implementation is almost complete

- users can upload attestations, PyPI will verify them and publicly expose them under appropriate API endpoints and JSON/HTML indices
- last piece: uploaded attestations are **not yet** reflected in the web UI itself
  - soon™

#### ecosystem-wide

- once the uploader default is switched, we expect similar signing volumes as total Trusted Publishing traffic (>= 400k attested uploads over the next year)
- all without users having to lift a finger!

### what's next?

### the elephant in the room

- despite all of this, something big is missing
  - in fact, the **most important component** of digital attestations
- verification!
- index-side verification is ~useless unless clients also verify
  - strictly a win in terms of transparency, but we can and should expect more!
  - o that means pip, uv, etc. should support PEP 740, *ideally by default*
- Python packaging's devolved nature means this will be an incremental process
  - pip has architectural limitations (can't depend on native code) that make integrating with sigstore-python challenging
    - our current workaround: developing a plugin system for pip that can then inject sigstore-python

### other bits

- PEP 740 currently specifies publish attestations and SLSA provenance
  - but there are other kinds of useful attestations, including 3p attestations!
  - current spec leaves space for these, but they're still an open design space
- current identity approach is tightly tied to Trusted Publishing identities
  - but there are other things PyPI can conceptually bind with:
    - verified account emails
    - domain names
    - even self-held signing keys (SSH? minisign?)
- reducing (and eventually eliminating) index trust
  - Python packaging needs a suitable lockfile format that can enable TOFU with PEP 740
    - PEP 751 🍐
  - blanket index-wide binary transparency, too boot?

### we need your help!

- do you do supply chain stuff™? consume PyPI's PEP 740 endpoints!
  - https://pypi.org/integrity/PROJECT/VERSION/FILENAME/provenance
    - NB: needs Accept: application/json
  - let us know how we can accommodate your use case!
- do you publish Python packages? enable TP/attestations now!
  - you'll get them by default if you're already using a Trusted Publisher soon, but being early doesn't hurt!
- do you consume Python packages at work? try verifying attestations!
  - jump way ahead of the curve and try verifying them directly without installing client support, if your work setup enables this kind of introspection

### thank you!

these slides will soon be available here:

https://yossarian.net/publications#soss-fusion-2024

#### resources:

- PEP 740: https://peps.python.org/pep-0740/
- Trusted Publishing docs: <a href="https://docs.pypi.org/trusted-publishers/">https://docs.pypi.org/trusted-publishers/</a>

#### contact:

- william@trailofbits.com
- @yossarian@infosec.exchange

