TRAIL OFBITS

The Next 5 Years of Supply Chain Security on PyPI

William Woodruff, Trail of Bits

introduction

Hello!

- William Woodruff (william@trailofbits.com)
 - o open source group engineering director @ trail of bits
 - long-term OSS contributor (Homebrew, LLVM, Python) and maintainer (pip-audit, sigstore-python)
 - <u>@yossarian@infosec.exchange</u>

Trail of Bits

- ~130 person cybersecurity engineering and auditing consultancy
- specialities: cryptography, compilers, program analysis research, "supply chain", OSS package management, general high assurance software development





introduction

thank-yous

- our work on PyPI and Python packaging security over the years has been funded by Google's Open Source Security Team, the Sovereign Tech Fund Agency, and the PSF themselves
- special thanks to Dustin Ingram & Hayden Blauzvern





Sovereign Tech Agency

agenda

- background on PyPI/Python packaging/Python community
- systematically securing Python packaging
 - or: "treating security first and foremost as a usability problem"
- previous efforts
- current and ongoing efforts
- the next 5 years
 - disclaimer: my opinions

background

PyPl

- the Python Package Index
 - o "Pie-pea-eye"
- the index behind pip install
- ~860K users, ~575K projects, ~6.1M releases, ~12.2M files
- ~51B downloads/month, >2TB traffic/day
 - somewhat inflated by tools and large CI/CD providers not caching as much as they could!



Downloads last day: 2,016,247,697 Downloads last week: 11,901,433,314 Downloads last month: 50,946,221,288

background

Python packaging/community

PyPI is just one piece of the Python packaging constellation

- installers: pip, uv, poetry, etc.
- build tools and backends: setuptools, hatch, flit, poetry, etc.
- upload tools: twine, uv, poetry, etc.
- very few single tools/toolchains do everything, although this is slowly changing

Python packaging is standards driven

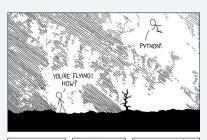
- Python Enhancement Proposals (PEPs) become living PyPA standards once accepted
- **not** top-down or tool-driven, unlike Rust (cargo) and JS (npm)

Python's community is massive and diverse

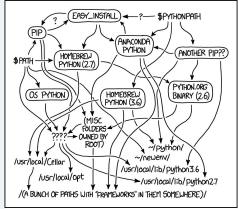
- every possible combination of experience, background, interest, skill, time commitment, etc.
- this makes it hard to establish a common denominator for *anything*, much less security!
- more closely resembles a country/countries than most other ecosystems!

systematically securing Python packaging

- diversity is both Python's strength and also a challenge to systematic security changes
 - a lack of systematic planning is self-fulfilling, since new systematic plans need to accommodate conventions/expectations
 - systematic planning can be dangerous to the things that keep Python's community healthy: onerous security requirements scare newcomers away!







MY PYTHON ENVIRONMENT HAS BECOME SO DEGRADED
THAT MY LAPTOP HAS BEEN DECLARED A SUPERFUND SITE.

systematically securing Python packaging

to make systematic progress on packaging security, we need to treat security as a *usability* problem!

two core demographics:

- enthusiasts: 1-5% who will take the initiative to opt into things
 - good news: most of the top packages are maintained by this demographic!
 - bad news: they're still a tiny overall percentage of package maintainers
- everyone else: 95-99% who only care about the reliability of their workflow; security changes must not require significant behavioral changes unless those changes make things easier for them
 - good news: *if* we make things easier for them, then they'll happily adopt changes

previous efforts

early steps

- 2003-ish: PyPI comes into existence
 - predated by less-formal packaging, back to 1998
- 2005: PyPI begins to host files
 - previously was **just** an index, pointing to other hosts
 - HTTP + HTTP redirects to random servers!
- 2013-ish: PyPI adds HTTPS
 - made mandatory for Web users + default for pip and easy_install
- 2015: PEP 503 standardizes the index API
 - including digests!
 - PEP 470 formally deprecates external hosting and removes remaining external links
- 2017: PyPI is rewritten (from "legacy PyPI" to Warehouse)
- 2018: current PyPI backend (Warehouse) goes live

the *last* 5 years

- 2019: PyPI implements and enables MFA
 - TOTP + WebAuthn (physical tokens)
- 2019: PyPI implements API tokens
 - includes project scoping
 - complements user/password authentication for uploads
- 2020: initial attempt to implement TUF (PEP 458)
 - initial attempts at malware detection/scanning on PyPI
- 2022: MFA enforced for "critical" projects
 - WebAuthn tokens given to designated projects for free
- 2023: Trusted Publishing implemented
 - complements manual API tokens
- Jan 1, 2024: MFA mandatory for all projects
 - also makes API tokens/Trusted Publishing mandatory for uploads!

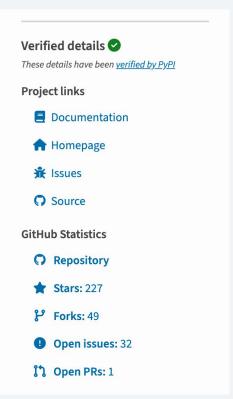


Add a new publisher

GitHub
Read more about GitHub Actions's OpenID Connect support h
Owner (required)
octo-org
The GitHub organization name or GitHub username that owns the repository
Repository name (required)
sampleproject
The name of the GitHub repository that contains the publishing workflow
Workflow name (required)
release.yml
The filename of the publishing workflow. This file should exist in the .github/workflows/ directory in the repository configured above.
Environment name (optional)
release
The name of the <u>GitHub Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your repository has maintainers with commit access who shouldn't have PyPl publishing access.

verified project URLs

- where "verified" means "PyPI can corroborate this URL"
 - uses Trusted Publishing as the source of ground truth!
- makes it harder for typosquatters, etc. to impersonate a project by submitting a project with malicious contents but valid-looking links
 - a/k/a "starjacking"



attestations!

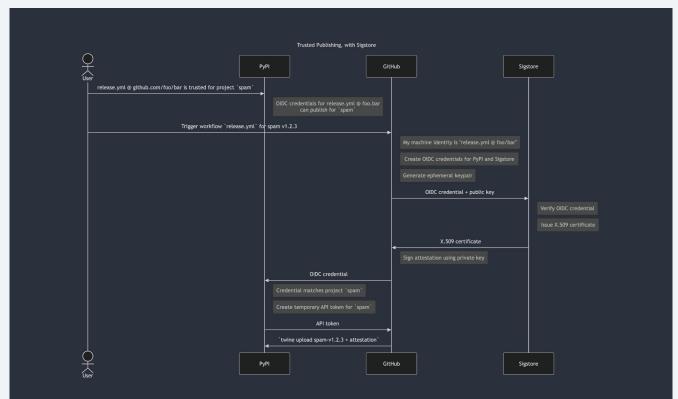
- Trusted Publishing was a huge success
 - because it was a usability and security win
- PEP 740 bootstraps package signing on top of Trusted Publishing
 - means it can be *enabled by default* for ~20,000 publishing workflows, with no additional configuration!
 - built on top of Sigstore → no long-term signing keys, only machine identities
 - live on PyPI as of last month!





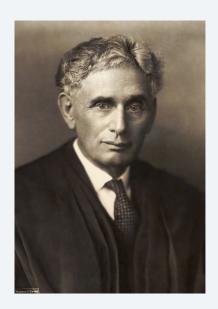


Trusted Publishing + attestations



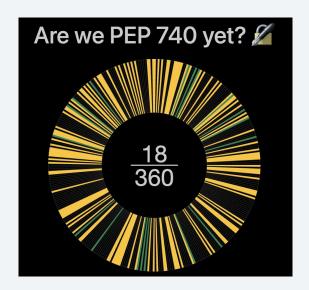
what do attestations actually get us?

- **publicly verifiable** proof of a Python package's provenance
 - equivalent to the sentence "package foo version 1.2.3 was produced by GitHub repository foo/foo via workflow pypi.yml at commit cafe..."
- external transparency and auditability for the index and individual publishers
 - each attestation has a corresponding transparency log entry that can be audited/reviewed
 - ...and a corresponding machine identity that can be proactively monitored
 - makes it harder for an attacker to **get away** with compromising a package or its repository
 - not perfect, since attestations aren't (and can't be made) mandatory!



attestation adoption so far?

- >10,000 attestations uploaded to PyPI since enablement
- 18/360 (5%) of top projects have attestations uploaded
 - 2/3rds haven't been uploaded since enablement, so could also be seen as 15%
- not bad considering it's only been enabled for 2 weeks!









extending attestations

- as defined in PEP 740, attestations are (intentionally) very limited:
 - only Trusted Publishers can be used as signing identities
 - attestation predicates intentionally limited for a MVP to just SLSA attestations and PyPI "publish" attestations
- more attestation types!
 - VSAs?
 - third-party attestations, e.g. attesting that a package has been reviewed?
- more signing identities!
 - email identities, since PyPI knows which email addresses have been verified
 - domain identities, possibly?

goal #1: expand the breadth/depth of meaningful attestations

it's one thing to sign; another to verify

- generating attestations is the "easy" part!
- the index (= PyPI) verifies attestations on upload, but getting user agents (e.g. pip) to verify is much harder
 - architectural constraints: pip needs to be pure Python for vendoring reasons, Sigstore's Python client isn't (uses Rust via PyCA cryptography)
 - standards/tooling limitations: lack of a standard lockfile format makes a "TOFU" verification story for users nontrivial; need to first establish a way to lock expected signing identities for packages

goal #2: bring attestation verification to downstream users!

attestations don't go far enough

- attestations are cool but *fundamentally limited in scope*
 - not everybody uses Trusted Publishing (or can use it), and this is fine!
 - *index-wide* monitoring is ~impossible to accomplish with individual user-side attestations
 - unbounded number of signing identities that need to be monitored for
 - attestations don't reduce trust in the index itself, meaning PyPI is still a juicy target
- want **blanket** transparency and authenticity properties
- binary transparency for PyPI!
 - prior art: Go's sumdb, npm's "publish attestations" (opposite of PyPI's)
 - desirable client-side properties: verification is 1+ ECC signatures + inclusion proof
 - downsides too: needs monitoring & witnessing ecosystem, evidence for these is scant

goal #3: index-level transparency for PyPI!

summary

the dream

today, have misuse-resistant credentials and opportunistic, on-by-default publish attestations:

thanks to Trusted Publishing + Sigstore based index-attestations (PEP 740)

tomorrow, we want:

- downstream verification of attestations
 - possible today, but not within official package tooling (yet)
- upstream (maintainer) *monitoring* of publishing identities
 - possible today, but UX is not good: Sigstore needs the equivalent of Cert Spotter for CI
- index-level transparency, with a healthy witnessing network
 - needs lots of design thought! you can help!

the dream

the ultimate goal is to be able to express sentences like this:





"when i download sampleproject v4.0.0, i know that it came from the repository /// it's supposed to come from and that



X PyPI served the same file to me as everyone else"

thank you!

these slides will soon be available here:

https://yossarian.net/publications#sigstorecon-2024

resources:

- PEP 740: https://peps.python.org/pep-0740/
- Trusted Publishing docs: https://docs.pypi.org/trusted-publishers/
- PyPI Attestation docs: https://docs.pypi.org/attestations/

contact:

- william@trailofbits.com
- @yossarian@infosec.exchange

