



Arch Linux Pacman

Security Assessment and Lightweight Threat Model

April 10, 2024

Prepared for:

Levente Polyak

Organized by the Open Technology Fund

Prepared by: **Spencer Michaels, David Pokora, Sam Alws, and Dominik Czarnota**

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at info@trailofbits.com.

Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2024 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to the Open Technology Foundation under the terms of the project statement of work and has been made public at the Open Technology Foundation's request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

Table of Contents

About Trail of Bits	1
Notices and Remarks	2
Table of Contents	3
Project Summary	4
Executive Summary	5
Project Goals	8
Project Targets	9
Project Coverage	10
Threat Model	11
Automated Testing	27
Codebase Maturity Evaluation	29
Summary of Findings	31
Detailed Findings	32
1. Use-after-free vulnerability in the print_packages function	32
2. Null pointer dereferences	35
3. Allocation failures can lead to memory leaks or null pointer dereferences	37
4. Buffer overflow read in string_length utility function	40
5. Undefined behavior or potential null pointer dereferences	43
6. Undefined behavior from use of atoi	45
7. Database parsers fail silently if an option is not recognized	47
8. Cache cleaning function may delete the wrong files	48
9. Integer underflow in a length check leads to out-of-bounds read in alpm_extract_keyid	49
A. Vulnerability Categories	51
B. Code Maturity Categories	53
C. Code Quality Findings	55
D. Fuzzing Pacman Code	60
E. Fix Review Results	68
Detailed Fix Review Results	69
F. Fix Review Status Categories	71

Project Summary

Contact Information

The following project manager was associated with this project:

Jeff Braswell, Project Manager
jeff.braswell@trailofbits.com

The following engineering director was associated with this project:

Anders Helsing, Engineering Director, Application Security
anders.helsing@trailofbits.com

The following consultants were associated with this project:

Spencer Michaels, Consultant
spencer.michaels@trailofbits.com

David Pokora, Consultant
david.pokora@trailofbits.com

Dominik Czarnota, Consultant
dominik.czarnota@trailofbits.com

Sam Alws, Consultant
sam.alws@trailofbits.com

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
November 13, 2023	Pre-project kickoff call
November 14, 2023	Discovery meeting #1
November 15, 2023	Discovery meeting #2
November 16, 2023	Discovery meeting #3
December 5, 2023	Delivery of report draft; threat model readout meeting
December 6, 2023	Code review readout meeting
March 7, 2024	Delivery of fix review appendix draft
April 10, 2024	Delivery of comprehensive report with fix review appendix

Executive Summary

Engagement Overview

The Open Technology Foundation (OTF) engaged Trail of Bits to review the security of the Pacman package manager, as well as its closely associated package management library `libalpm`. Pacman is the official package manager of Arch Linux and is developed by the Arch Linux team; it is also used in a handful of other Linux distributions, including Manjaro.

A team of two consultants conducted a threat model from November 13 to November 17, 2023, for a total of two engineer-weeks of effort; this was followed by a code review conducted by three engineers from November 20 to December 1, 2023, for a total of five engineer-weeks of effort. Our testing efforts focused on package signature verification, data integrity during downloads and upgrades, memory safety, and a new user-based isolation mechanism. With full access to source code and documentation, we performed static and dynamic testing of Pacman and `libalpm`, including fuzzing, using automated and manual processes. The audit scope excluded the parts of the Pacman ecosystem used exclusively for building packages, such as the `makepkg` script.

Observations and Impact

Overall, Pacman is well-designed, comprehensively documented, and robust against common application security issues. The code review portion of the engagement revealed several issues of low, informational, and undetermined severity. While the threat model revealed some plausible threat scenarios, these generally require the successful confluence of several independent attacks, such as compromising a mirror, obtaining a signing key, or intercepting a user's connection under certain configurations, which sets a relatively high bar for an attacker to achieve.

That said, certain defense-in-depth measures could be implemented to improve the resilience of Pacman and the Arch Linux distribution and signing infrastructure, even against cases where an attacker already has a partial foothold. Based on the threat model and code review results, three major areas of improvement stand out:

- Because Pacman is written in C, even security-conscious developers run a relatively high risk of accidentally introducing memory safety issues—we discovered several during the audit ([TOB-PACMAN-1](#), [TOB-PACMAN-4](#), and [TOB-PACMAN-9](#)), although none ultimately proved especially serious. We recommend using static and dynamic analyses, including fuzz tests, to uncover additional potential cases of memory corruption and leaks before attackers do.
- Pacman's signing infrastructure is robust against maintenance issues such as keys being lost (not stolen), signers becoming inactive or incapacitated, and so on. However, due to the lack of documented incident response procedures, the Arch

Linux team may be ill-equipped to promptly respond to a security incident involving theft or malicious use of key materials. Additionally, a lack of clear auditing guidelines and trust requirements for signers increases the likelihood that package signing keys could be used maliciously. As Arch Linux continues to grow as an organization, it is critical that security-related processes, guidelines, and requirements be clearly and precisely documented to ensure consistency and prompt response to security incidents.

- Pacman can verify database signatures, but Arch Linux's official databases are not signed and Pacman does not require databases to be signed by default. Combined with the fact that Pacman allows the use of plaintext HTTP package mirrors, users with such a configuration could be served malicious database files, which could serve old and vulnerable versions of packages. This issue is known to the Arch Linux team, and work is currently underway to rectify it.

Recommendations

Based on the codebase maturity evaluation and findings identified during the security review, Trail of Bits recommends that the Arch Linux team take the following steps:

- **Remediate the **code review findings** disclosed in this report.** These findings should be addressed as part of a direct remediation or as part of any refactor that may occur when addressing other recommendations.
- **Create a long-term plan for implementing the strategic recommendations in the **Threat Model** section of this report.** These findings should be addressed as part of a direct remediation or as part of any refactor that may occur when addressing other recommendations.
- **Clarify the intended use and safety guarantees of the `--root` argument.** This argument specifies which directory should be used by Pacman as the root directory. However, it is not guaranteed that files and directories *outside* of the root directory will remain untouched (e.g., if there is a maliciously placed symlink inside the root directory). Pacman's manual page entry states that the argument should not be used as "a way to install software into `/usr/local` instead of `/usr`" or "for performing operations on a mounted guest system." However, Pacman documentation does not state what this argument *should* be used for and does not give any information about the argument's (lack of) safety guarantees.
- **Implement security-focused static analysis, dynamic analysis, and fuzz tests.** These should be run against each new Pacman version prior to release to minimize the likelihood of ongoing code changes introducing memory corruption issues.

Finding Severities and Categories

The following tables provide the number of findings by severity and category.

EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	0
Medium	0
Low	1
Informational	5
Undetermined	3

CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Data Validation	5
Denial of Service	1
Undefined Behavior	3

Project Goals

The engagement was scoped to provide a security assessment of the Pacman package manager. Specifically, we sought to answer the following non-exhaustive list of questions:

- Is there any way to bypass Pacman's package signature validation?
- Is it possible to break out of the SandboxUser configuration option's filesystem context implemented in [MR #23](#)?
- Does the package consistency checking included in [MR #96](#) have any security issues?
- Is Pacman vulnerable to any form of memory corruption?
- Can an attacker with control over database contents (which are unsigned by default and may be accessed over plaintext HTTP) cause Pacman to exhibit malicious behavior?
 - In particular, can a malicious database silently downgrade a package to a known-vulnerable version, install a vulnerable package, or uninstall a package providing security measures?
- Can a malformed package, or malformed metadata, cause Pacman to bring the system into an inconsistent state?
- Are Pacman's defaults conducive to secure operation by ordinary users?
- Does Pacman call out to third-party programs or libraries in unsafe ways?
- Does Pacman's current test suite appropriately cover security-related concerns?
- Is Arch Linux's package signing infrastructure robust against failures and resilient to compromise, including by malicious insiders?
- Are Arch Linux's official package repositories reasonably well protected against the unexpected introduction of malicious code or metadata?

Project Targets

The engagement involved a review and testing of the targets listed below, including two as-yet-unmerged merge requests.

Pacman

Repository	https://gitlab.archlinux.org/pacman/pacman/
Version	18e49f2c97f0e33a645f364ed9de8e3da6c36d41
Type	C binary application
Platform	Linux

MR #23: Add SandboxUser configuration option

URL	https://gitlab.archlinux.org/pacman/pacman/-/merge_requests/23
-----	---

MR #96: Check package consistency when installing

URL	https://gitlab.archlinux.org/pacman/pacman/-/merge_requests/96
-----	---

Project Coverage

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- A lightweight threat model of Pacman and the portion of its infrastructure related to package signing and distribution
- A non-exhaustive manual review of the Pacman codebase and two security-relevant merge requests pending acceptance, with a focus on code paths pertaining to security-critical functionality highlighted in the initial threat model
- Static analysis of the Pacman codebase and manual triage of the results
- Dynamic analysis to identify instances of memory corruption and leaks
- Fuzzing to identify inputs that could cause unexpected behavior at runtime

Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. The following list outlines the coverage limitations of the engagement and indicates system elements that may warrant further review:

- The code of various dependencies used by Pacman, such as `libarchive` and `gpgme`, was not reviewed.
- Although we included signing and packaging infrastructure security controls in the threat model, we did not have access to review their implementation during the code review.
- Our threat model exercise focused primarily on risk introduced to users; while we touched on risk introduced to developers and packagers, this was not a major focus, and some elements of that risk profile (such as packaging tooling) were excluded from the scope.

Threat Model

As part of the audit, Trail of Bits conducted a lightweight threat model, drawing from [Mozilla's Rapid Risk Assessment methodology](#) and the National Institute of Standards and Technology's (NIST) guidance on data-centric threat modeling ([NIST 800-154](#)). The results of the lightweight threat model are described in the subsections below.

Data Types

The target application uses the following data formats:

- TAR files (`.tar`), usually compressed (`.zst`, `.gz`, or `.xz`): Pacman package files
- Bash scripts: `PKGBUILD` and `INSTALL` files
- INI configuration files: hooks and configuration files (e.g., `pacman.conf`)
- Plaintext: `PKGINFO`, `BUILDINFO`, database, and file-list files
- PGP keys

Data Flow

Pacman is the default package manager for Arch Linux, officially maintained by the Arch Linux development team.

Pacman retrieves packages from one or more repositories, which can either be located on the local host's filesystem or accessed over the network via any protocol supported by the `libcurl` dependency, which Pacman uses internally. Packages can also be directly installed from the filesystem without being associated with a repository.

In a typical use case, users download the vast majority of their packages prebuilt from HTTP or HTTPS mirrors of the remote Arch Linux official repositories. A small number of unofficial packages, such as those from the Arch User Repository, may be built and installed either directly as a manually built package file on the local filesystem or from a repository hosted on the local filesystem.

When a package is installed, Pacman verifies its signature using an internal Pacman keyring, with the root of trust derived from a unique system master key that is generated upon system installation and used to sign the set of main signing keys imported into the system-local Pacman keyring. These main signing keys, of which there are only five, are used by the Arch Linux developers to sign the packaging keys that package maintainers use to sign their packages. Each main signing key has an associated revocation key, held in the possession of a different trusted signer, which can be used to revoke it in the event of a

compromise. Those keys, along with the names of the developers they belong to, are listed on the [Arch Linux Master Signing Keys](#) web page.

Data for keys stored in Pacman's internal keyring is kept up to date via the web key directory (WKD) sync service, which runs weekly on Arch Linux and syncs with a distributed WKD. The sync service can only update metadata (such as expiration dates) for existing stored keys; it cannot alter whether a given key is trusted.

Maintainers generally use `makepkg` to generate Pacman packages from application or library sources. The build scripts for Arch Linux's official packages are hosted on a dedicated [GitLab account](#) with the login handled by an Arch Linux-managed Keycloak single sign-on (SSO) instance. Most official packages are built on a single, high-capacity main build server administered by dedicated DevOps members of the Arch Linux team and accessible via SSH; however, individual maintainers may build and sign packages on other machines. Packagers are currently strongly encouraged, although not strictly required, to retain signing keys only on hardware keys (as opposed to on their local filesystem).

All official Arch Linux packages are currently signed, and by default, Pacman requires packages from remote repositories to have a valid signature trusted by Arch Linux's main signing keys. Package installation transactions may be preceded and/or followed by hooks, which can invoke arbitrary commands in response to the presence of specific packages in a transaction. Packages, along with detached signatures, are cached in a package cache directory (`/var/cache/pacman/pkg`) on the local filesystem after installation.

In figure 1, we depict known connections between system components of the package consumption side of Pacman, as integrated into Arch Linux. This diagram is intended to convey our understanding of the system as a whole. Further details will be described in the Components and Trust Zones and Trust Zone Connections subsections that follow. The dotted lines indicate trust boundaries separating zones.

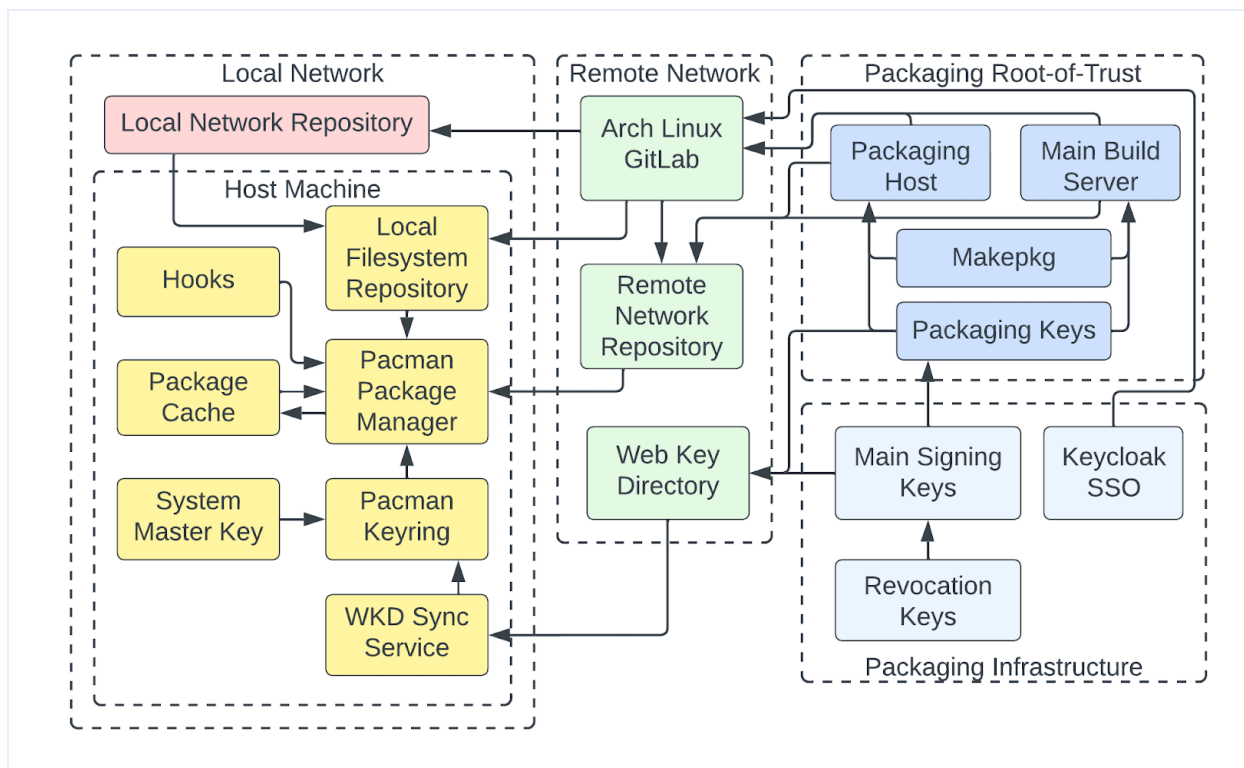


Figure 1: The data flow of packages and their signing data from Arch Linux's root of trust to the host machine on which Pacman runs

Components and Trust Zones

The following table describes the components that make up the Pacman package management system, as well as the external dependencies on which they rely. These system elements are further classified into *trust zones*—logical clusters of shared functionality and criticality, between which the system enforces (or should enforce) interstitial controls and access policies.

Components marked by asterisks (*) are considered out of scope for the assessment. We explored the implications of threats involving out-of-scope components that directly affect in-scope components, but we did not consider threats to out-of-scope components themselves.

Component	Description
Host machine	This component is the host that Pacman manages packages on.
Pacman package manager	Pacman is a package management tool that tracks installed packages on a Linux system, including support for dependency resolution and retrieval, package groups, install and uninstall scripts, and pre- and post-install

	hooks. It also contains utilities such as <code>makepkg</code> , which is used to create packages that can be installed by Pacman.
Local filesystem repository	This repository resides on the local filesystem. Repositories provide a list of packages that can be fetched, installed, or upgraded. The package list is managed by the repository maintainer(s). Packages can be signed or unsigned.
Package cache	This directory is populated with previously installed packages. Cached packages are used to rapidly reinstall a previously installed package. Any signatures contained within packages are also stored alongside and validated for each cached package item.
Pacman keyring	This keyring contains signing keys for all packages installed on the system. Keys within the keyring are considered trusted only if they are signed by an Arch Linux packaging key (which is in turn signed by a main signing key).
System master key	The system master key is the root of trust for Pacman's signature validation on any given installation. Master keys are generated at first Pacman run (and thus on first boot of Arch Linux), are unique to each Arch Linux install, and are used by the host machine to trust the Arch Linux main signing keys.
WKD sync service	This GnuPG wrapper service on Arch Linux runs weekly to sync updates (e.g., expiration dates) to keys in the Pacman keyring, pulled from a WKD. The WKD sync service can add previously unknown signatures to the keyring but cannot make Pacman trust those signatures.
Hooks	Pre- and post-install hooks enable running commands just before or after a Pacman transaction (e.g., to rebuild a new kernel image after Pacman installs a new kernel version).
Local network	These components share a local network with the host machine.
Local network repository	This repository resides on the local network. Repositories provide a list of packages that can be fetched, installed, or upgraded. The package list is managed by the repository maintainer(s). Packages can be signed or unsigned.
Remote network	This network includes the components that live outside the local network trust zone (e.g., public-facing, external network components).

Remote network repository	This repository resides on the remote network host. Repositories provide a list of packages that can be fetched, installed, or upgraded. The package list is managed by the repository maintainer(s). Packages can be signed or unsigned.
Web key directory (WKD)	This directory is GnuPG's standard system for key discovery, which maps public keys to email addresses.
Arch Linux GitLab (*)	This GitLab account hosts the source code for official Pacman packages.
Packaging infrastructure	These machines (and their operators) build and sign Pacman packages.
Main build server*	The Arch Linux team uses this dedicated, high-capacity machine to build the majority of its official packages.
Packager host	This host is operated by a packager and is used to build and sign packages.
Packaging keys	These keys are used by package maintainers to sign packages. Each trusted maintainer is issued a packaging key signed by a quorum of main signing key holders.
makepkg*	This toolset is used to build Pacman packages.
Packaging root of trust	These components are used to facilitate administration of an operating system's primary mirrors and to manage their authorized package signers.
Main signing keys	These keys are the root of trust for Arch Linux's signing infrastructure and can sign new packaging keys as well as packages themselves. Currently, only five main signing keys exist.
Revocation keys	Each main signing key has a single associated revocation key used to revoke the signing key in the event of compromise. Each signing key's revocation key is held by another signing key owner.
Arch Linux DevOps	These individuals administer Arch Linux's Keycloak SSO instance, GitLab account, etc.

Trust Zone Connections

At a design level, trust zones are delineated by the security controls that enforce the differing levels of trust within each zone. Therefore, it is necessary to ensure that data cannot move between trust zones without first satisfying the intended trust requirements of its destination. We enumerate such connections between trust zones below.

Originating Zone	Destination Zone	Data Description	Connection Type	Auth Type
Host machine	Host machine	<p>All operations performed by Pacman that leverage components in the same zone largely rely on cryptographic verification (e.g., signed packages, packaging key authorization).</p> <p>The artifacts written by Pacman are written in a root-user execution context, with file permissions blocking access by lower-privileged users.</p>	Filesystem	File privileges, GnuPG signature validation
Remote network	Host machine	The host's WKD sync service pulls updated key information from a WKD into the Pacman keyring.	HTTPS	None
Remote network, local network	Host machine	The host installs a package from a local or remote network repository.	libcurl-supported protocols (e.g., HTTP, HTTPS, FTP)	GnuPG signature validation, libcurl-supported protocols (e.g., TLS)

Remote network	Local network	Third-party package sources pulled from the internet are downloaded to and built on the local network; the resulting packages are placed in a local network repository.	Varies; likely HTTP or HTTPS	Varies or none
Remote network	Packaging infrastructure	Package sources hosted on the Arch Linux GitLab account are downloaded to and built on the main build server or a packager host.	HTTPS	SSH
Packaging root of trust	Packaging infrastructure	A quorum of main signing key holders signs a new packaging key or issues revocations for an existing one.	N/A	GnuPG signature verification
Packaging root of trust	Arch Linux GitLab	An Arch Linux administrator logs into GitLab through the Keycloak SSO instance.	HTTPS	OAuth

Threat Actors

When conducting a threat model, we define the types of actors that could threaten the security of the system. We also define other system users who may be impacted by, or induced to undertake, an attack. For example, in a confused deputy attack such as cross-site request forgery, a normal user who is induced by a third party to take a malicious action against the system would be both the victim and the direct attacker. Establishing the types of actors that could threaten the system is useful in determining which protections, if any, are necessary to mitigate or remediate vulnerabilities. We will refer to these actors in descriptions of the security findings that we uncovered through the threat modeling exercise.

Actor	Description
End users	End users of Pacman are consumers of its packages and repositories. They operate in the host machine zone and may have influence over the local network zone and its repositories.
Local user	A local user is a low-privileged user on the host machine (e.g., non-admin, non-root) who cannot execute sensitive Pacman operations because they require root access.
Local root	The root user on the host machine has privileges to perform any operations they desire. Pacman requires a local user to elevate to local root to install or update packages.
Operators	Operators are privileged actors with the responsibility of operating packaging infrastructure and packaging root-of-trust components.
Repository administrator	A repository administrator has control over a Pacman repository or mirror. They may operate a local machine, local network, or remote repository.
DevOps administrator	A DevOps administrator has control over Arch Linux's DevOps infrastructure, including the Arch Linux GitLab account and Keycloak SSO instance.
Packager	This individual possesses a packaging key that was signed and approved by a trusted signer.

Trusted signer	<p>This individual possesses a master signing key, which is a single keypair used in a threshold signature scheme (TSS) and performs sensitive operations such as approving a new packaging key.</p> <p>Trusted signers, in a quorum, act as a root of trust for Pacman repository management.</p>
Attacker	An attacker is positioned either within or external to any of the trust zones previously described.
Internal attacker	An internal attacker has transited one or more trust boundaries. Such an attacker may be an existing actor in the system or an external attacker who has successfully transited a trust boundary into the system.
External attacker	An external attacker is external to the cluster and is unauthenticated, such as an attacker with control over external services.

Threat Scenarios

The following table describes possible threat scenarios given the design, architecture, and risk profile of the Pacman package manager.

Scenario	Actor(s)	Component(s)
<p>An operating system provides a default mirror list leveraging insecure protocols. Developers of an operating system such as Arch Linux may generate a list of repository sources that leverage insecure protocols (e.g., HTTP, FTP). Due to Pacman's lack of protocol restrictions, its underlying <code>libcurl</code> dependency will communicate over the insecure protocol.</p> <p>If a local user or local root actor uses this insecure protocol to fetch packages from a local network or remote repository, it may expose them to person-in-the-middle attacks. Although such an attack may not be problematic for signed packages, an attacker may substitute unsigned packages with maliciously crafted packages.</p>	<ul style="list-style-type: none">• Repository administrator• Trusted signer• Attacker	<ul style="list-style-type: none">• Pacman package manager
<p>An operating system that leverages Pacman does not enforce signed packages. By default, Arch Linux requires all packages to be signed to be installed, verifying they have been approved. If a Linux distribution does not configure Pacman to require signatures, this may introduce risk, compounding the threat scenario mentioned in the previous row of this table.</p> <p>Unsigned packages may indicate a lack of an approval process. They may be subject to modification in-flight through a person-in-the-middle attack that could put users at risk.</p> <p>The package cache containing a copy of previously unsigned installations may also be modified if it is improperly secured. By default, Arch Linux saves package cache items with special privileges that should disallow any user except the local root user from modifying them, which mitigates this risk.</p>	<ul style="list-style-type: none">• Repository administrator• Packager• Trusted signer• Attacker	<ul style="list-style-type: none">• Pacman package manager

<p>An environment variable affects the Pacman package manager's libcurl dependency. For instance, Pacman redirects its HTTP connections through the proxy defined in the <code>http_proxy</code> environment variable. If an attacker injects an environment variable into Pacman's runtime environment—a difficult prospect, given that it runs as root during installs—they could potentially cause Pacman to exhibit exploitable or undesirable behavior.</p>	<ul style="list-style-type: none"> • Local root 	<ul style="list-style-type: none"> • Pacman package manager
<p>An attacker attempts a substitution attack, bumping versions on a popular package through a compromised local network repository or remote repository. Pacman will always install the latest version of a package across all repositories it has access to. As such, if a user has both local and remote repositories enabled, an attacker who can introduce an identically named, higher-versioned package into one of the remote repositories can easily induce the user to install this version of the package. Similar attacks may also be possible via DNS confusion (e.g., if an attacker registers a domain that shadows a local network domain name). See this GitHub blog post on substitution attacks against npm.</p>	<ul style="list-style-type: none"> • Repository administrator • External attacker 	<ul style="list-style-type: none"> • Pacman package manager • Local network repository • Remote network repository
<p>An attacker compromises a packaging key and produces different but valid signatures for a package to introduce malicious changes. In this case, Pacman would install the new package version normally, and the user would be entirely unaware. Currently, there is no way to enable a warning when a package's signature changes.</p>	<ul style="list-style-type: none"> • Packager • Internal attacker 	<ul style="list-style-type: none"> • Pacman package manager • Packaging keys
<p>A packaging key, or packager, is compromised, requiring revocation of the packaging key. Due to a lack of documented procedures for revocation, response by trusted signers may be delayed, giving the attacker more time to cause damage.</p>	<ul style="list-style-type: none"> • Packager • Trusted signer 	<ul style="list-style-type: none"> • Packaging keys

<p>A trusted signer's key is compromised, requiring incident response. Due to a lack of documented procedures for revocation, response by the trusted signer holding the compromised key's revocation key may be delayed, giving the attacker more time to cause damage.</p>	<ul style="list-style-type: none"> • Trusted signer 	<ul style="list-style-type: none"> • Main signing keys • Revocation keys
<p>The (unsigned) Pacman database used to index packages is modified by an attacker. The database used by Pacman is not signed. The database is used as an index for packages on the system.</p> <p>Although most of the data used by Pacman is derived from signed packages on Arch Linux, the database is used to determine depends/replaces directives when installing a package. This is done without verifying that the depends/replaces data taken from package metadata has not been tampered with.</p> <p>As a result, an attacker with access to the Pacman database may replace depends/replaces directives within the database for a given package to trigger a deletion or replace-with-empty operation of an existing package on the user's system.</p>	<ul style="list-style-type: none"> • End user • Internal attacker • Packager 	<ul style="list-style-type: none"> • Pacman package manager • Local network repository • Remote network repository
<p>Vulnerable or malicious packages are assigned a package group with a name identical to a popular existing package. Currently, Pacman always resolves such a conflict in favor of the group, with no way to override this behavior. As such, users could be made to unwittingly install an arbitrary package or set of packages in place of a common package.</p>	<ul style="list-style-type: none"> • Packager • End user 	<ul style="list-style-type: none"> • Pacman package manager • Local network repository • Remote network repository

<p>A naive user sets overly permissive file permissions on their keyring, configuration files, or hooks. An attacker who achieves local filesystem access—for example, by compromising a low-privileged service—could inject malicious settings, commands, or additional trusted keys to perform privilege escalation.</p>	<ul style="list-style-type: none"> • End user 	<ul style="list-style-type: none"> • Host machine
<p>A revocation certificate or signing key (e.g., package key, trusted signer key) is lost or corrupted. Since there are no standard procedures for regular checks of keys or their backup media after initial creation, it is possible that keys could be permanently lost. In particular, since revocation keys are long-lived and very rarely used, they may become inaccessible (e.g., through corrupted media) long before this fact is discovered, with users only realizing too late when the key is sorely needed.</p>	<ul style="list-style-type: none"> • Trusted signer • Packager 	<ul style="list-style-type: none"> • Revocation keys • Main signing keys
<p>An attacker compromises a mirror of Arch Linux official packages, or intercepts a user's non-TLS connection to a repository, and injects a malicious version of a package. In this case, Pacman would refuse to install the package, as it requires signatures from remote repositories by default.</p>	<ul style="list-style-type: none"> • End user • Repository administrator • Internal attacker 	<ul style="list-style-type: none"> • Pacman package manager • Local network repository • Remote network repository

Recommendations

Trail of Bits recommends that the Arch Linux team implement the following recommendations to mitigate the threat scenarios described above:

1. **Set Pacman to reject non-TLS mirrors by default.** Since databases are not currently signed, an attacker who can intercept an unauthenticated connection between a user and a repository could modify their contents in transit. A new configuration value such as `AllowInsecureMirrors` can be added to `pacman.conf` to permit the use of non-TLS mirrors on a case-by-case basis if necessary for backwards compatibility.
 - Consider also allowing users to set a minimum TLS version in `pacman.conf`, defaulting to at least TLS 1.2 (disabling **specific weak ciphersuites supported in 1.2**) or, ideally, TLS 1.3. Otherwise, HTTPS downgrade attacks may be possible against TLS-enabled mirrors that support older, insecure TLS versions.
 - Update the official Pacman mirror lists to exclude non-TLS mirrors, and consider modifying the `reflector` mirror list ranking tool to account for TLS settings (i.e., giving higher rankings to mirrors with stricter settings).
2. **Transition to signed databases and require them by default.** Currently, Pacman obtains packages' `depends/replaces` lists from the database. Since databases are unsigned, an attacker with the ability to modify them could induce a user to install or remove arbitrary packages.
 - A patch is currently in progress that would check package metadata listed in the database against the metadata contained within the actual signed package to be installed, which partially mitigates this issue.
3. **Warn users (or give them the option to be warned) when a package's signature changes during an upgrade, even if the signature is valid.** This will provide a defense-in-depth measure against cases where an attacker gains possession of a valid signing key and signs a package not previously signed with that key.
 - Consider introducing a setting into `pacman.conf` that would toggle these warnings between "off," "print only," and "pause upgrade and interactively ask for confirmation to continue"—the latter case being suitable for especially cautious users. Depending on how often packages' signing keys change in legitimate cases, the default for this setting could be either "print only" (if rare) or "off" (if common).

4. **Provide an interactive resolution prompt in cases where a package and a group have the same name.** Currently, Pacman considers group names to “shadow” identically named packages; as such, an attacker who can tag a malicious or vulnerable package as belonging to a group with the same name as a common package—for instance, by manipulating an unsigned database—can cause users to unwittingly install the package of the attacker’s choice. In the event of a conflict, the user should be prompted to make an explicit selection (in the same manner as the provides directive).
5. **Have Pacman refuse to load its keyring, configuration files, or hooks if they are writable by users other than the root user.** Analogous to SSH’s permissions checks on the `~/.ssh` directory, this would prevent users from unknowingly directing Pacman (which is likely running as root) to use a keyring, configuration, or hook that a lower-privileged user or service could maliciously modify to perform privilege escalation.
6. **Establish a detailed, written incident response plan that defines how to respond to high-severity threat scenarios,** especially the following:
 - Compromise of a main signing key
 - Compromise of a packaging key
 - Compromise of a DevOps-managed property such as the Arch Linux Gitlab account, Keycloak SSO instance, etc.

The plan should detail precisely who is responsible for threat response and the exact steps they should follow to mitigate the threat. Having such guidance in place ensures that there is no ambiguity about how to handle a security incident when it actually happens, ensuring the fastest and most thorough response possible.

7. **Establish procedures for regularly validating trusted signers’ main signing key and revocation key backups over time** to ensure that they remain usable and readily accessible. In addition, provide detailed guidance about how operators should configure and use cold storage backups, ensuring redundancy in case their primary keypair is corrupted.
 - Test not only the accessibility and integrity of the backup media but also the viability of the keys in question: for instance, import revocation keys into a test keyring on a regular basis to ensure that they do indeed revoke the expected signing keys.

8. **Establish a written list of procedures and requirements for onboarding a new trusted signer.** Currently, any potential new trusted signer must be well-known to the Arch Linux team and a long-term participant within the Arch Linux ecosystem so that candidates are already extensively vetted. Formalizing this process would reduce the likelihood of making mistakes or exceptions.
 - Consider verifying trusted signers' legal IDs. The current onboarding process, while in effect vetting candidates' real-world identities quite extensively, does not require actual verification of their legal IDs. Doing so would provide an additional layer of defense in depth and better allow the Arch Linux team to hold a defecting signer legally accountable if necessary.
9. **Establish standards for regular checkups on packagers.** Currently, trusted signers make a best-effort attempt to identify packagers who are inactive or are not fulfilling their duties; however, this is not done systematically or at regular intervals. Doing so consistently would minimize the chances of inactive or irresponsible signers slipping through the cracks.
10. **Establish clear security guidelines for trusted signers and packagers,** including how to generate, store, and use key material; how to report a compromise of their own key material; what to do if a trusted signer reports a compromise; and so on.
 - Require packagers to keep key material on hardware keys only. Currently, this practice is strongly encouraged but not mandated, and some packagers sign using key material on their local filesystems.
11. **Consider replacing uses of MD5 with a hashing algorithm with a lower chance of hash collision, such as BLAKE2.** MD5 has a nontrivial chance of collisions, and it is feasible to intentionally craft a file with a specific MD5 hash. Some .pacsave backups, which use MD5 hashing to compare files, may not occur in the case of a hash collision even if the files in question do differ. However, performance or compatibility considerations may prohibit the use of algorithms with lower rates of hash collision.

Automated Testing

Trail of Bits uses automated techniques to extensively test the security properties of software. We use both open-source static analysis and fuzzing utilities, along with tools developed in house, to perform automated testing of source code and compiled software.

Test Harness Configuration

We used the following tools in the automated testing phase of this project:

Tool	Description	Policy
<code>scan-build</code>	This static analysis tool can find various issues within C/C++ codebases.	Default checks
<code>libFuzzer</code>	This in-process, coverage-guided, evolutionary fuzzing engine can automatically generate a set of inputs that exercises as many code paths in the program as possible.	Appendix D

Areas of Focus

Our automated testing and verification work focused on the following system properties:

- The program does not access invalid memory addresses.
- The program does not exercise undefined behavior.

Test Results

The results of this focused testing are detailed below.

Fuzzing harnesses. The fuzzing harnesses we developed exercise a subset of the program's code.

Property	Tool	Result
<code>fuzz_string_length</code> : harness that checks the utility function that computes the length of a string, omitting ANSI escape codes	libFuzzer	TOB-PACMAN-4
<code>fuzz_wordsplit</code> : harness that checks the utility function that splits a string into multiple words	libFuzzer	Did not find issues

fuzz_parseconfigfile: harness that tests the parsing of configuration files; requires further changes so it is chrooted and so that the parser does not include external files from the filesystem	libFuzzer	Requires further development (see appendix D)
fuzz_alpm_extract_keyid: harness that tests the extraction of keys from signature data	libFuzzer	TOB-PACMAN-9

Codebase Maturity Evaluation

Trail of Bits uses a traffic-light protocol to provide each client with a clear understanding of the areas in which its codebase is mature, immature, or underdeveloped. Deficiencies identified here often stem from root causes within the software development life cycle that should be addressed through standardization measures (e.g., the use of common libraries, functions, or frameworks) or training and awareness programs.

Category	Summary	Result
Arithmetic	Although the code attempts to test the computed indexes or path lengths, the project does not take specific measures to ensure arithmetic safety. For example, we found an instance where an integer underflow occurs in a length check function. Additionally, the length check values are often computed from hard-coded integer constants, instead of using the <code>sizeof()</code> operator to compute the length of the hard-coded string that the integer constant's length is derived from.	Moderate
Auditing	Pacman generally preserves standard errors from subprocesses (e.g., hooks) and produces useful, detailed messages when package operations encounter errors.	Satisfactory
Authentication / Access Controls	Pacman itself does not require authentication or attempt to authenticate to other services.	Not Applicable
Complexity Management	Pacman's codebase is neatly organized, with discrete functionality arranged into separate files and functions, accompanied by clear comments and documentation.	Strong
Configuration	Pacman calls well-vetted third-party libraries for complex functionality such as downloads (e.g., <code>libcurl</code>) and signature verification (e.g., <code>OpenSSL</code>) and uses those libraries according to their respective best practices.	Strong
Cryptography and Key Management	Pacman uses <code>OpenSSL</code> for all cryptographic operations. Arch Linux's signing infrastructure has built-in resilience measures such as physical key backups, quorum requirements, and public oversight. However, no written	Satisfactory

Category	Summary	Result
	incident response plans exist. This could increase the team's response time if the signing infrastructure is compromised.	
Data Handling	While the code generally attempts to verify the data it receives, there are certain cases where the performed checks are insufficient and could cause memory corruption or undefined behavior.	Moderate
Documentation	Pacman and <code>libalpm</code> are both extensively documented, including in code comments, manual pages, and the Arch Linux wiki.	Strong
Maintenance	We discovered issues in how the Arch Linux team members maintain the signing infrastructure itself. While the Arch Linux team occasionally audits package signers on an informal basis, no formal process has been defined for how, and how often, such audits should take place. In addition, revocation key backups are not checked regularly after they are first generated; if a backup fails, signers may be unable to revoke a compromised key in a timely manner.	Moderate
Memory Safety and Error Handling	<p>We uncovered some instances of memory safety issues where certain parsing routines can read memory out of bounds. We recommend regularly fuzzing those and other code paths to cover more edge cases and help catch new problems.</p> <p>Errors are generally handled consistently within the codebase, though there are cases where allocation failures are not acted upon apart from logging, which could be hard to recover from. Additionally, the code could benefit from using status code enum return types (instead of plain <code>int</code> types) for its public functions.</p>	Weak
Testing and Verification	Pacman has substantial test coverage for expected functionality but none that focuses on unexpected inputs or potentially malicious behavior (e.g., fuzz tests).	Moderate

Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Use-after-free vulnerability in the <code>print_packages</code> function	Undefined Behavior	Low
2	Null pointer dereferences	Denial of Service	Informational
3	Allocation failures can lead to memory leaks or null pointer dereferences	Undefined Behavior	Informational
4	Buffer overflow read in <code>string_length</code> utility function	Data Validation	Undetermined
5	Undefined behavior or potential null pointer dereferences	Data Validation	Undetermined
6	Undefined behavior from use of <code>atoi</code>	Undefined Behavior	Informational
7	Database parsers fail silently if an option is not recognized	Data Validation	Informational
8	Cache cleaning function may delete the wrong files	Data Validation	Informational
9	Integer underflow in a length check leads to out-of-bounds read in <code>alpm_extract_keyid</code>	Data Validation	Undetermined

Detailed Findings

1. Use-after-free vulnerability in the print_packages function

Severity: Low

Difficulty: Low

Type: Undefined Behavior

Finding ID: TOB-PACMAN-1

Target: pacman/src/pacman/util.c

Description

The print_packages function is subject to a use-after-free vulnerability. The function first deallocates memory for the temp variable and then uses that memory in the PRINT_FORMAT_STRING macro (figure 1.1), which can lead to the following issues:

- Potential exploitation of the program could occur if an attacker can allocate and control the value of the temp variable after it is freed (highlighted line 1 in figure 1.1) and before it is used in another thread (highlighted line 2 in figure 1.1). The time window for the attack is very small since the two operations happen one after another.

```
void print_packages(const alpm_list_t *packages) {  
    ...  
    /* %s : size */  
    if(strstr(temp, "%s")) {  
        char *size;  
        pm_asprintf(&size, "%jd", (intmax_t)pkg_get_size(pkg));  
        string = strreplace(temp, "%s", size);  
        free(size);  
        free(temp); // (1) memory pointed by the temp variable is freed  
    }  
    /* %u : url */  
    PRINT_FORMAT_STRING(temp, "%u", alpm_pkg_get_url) // (2) use-after-free of  
    temp
```

Figure 1.1: pacman/src/pacman/util.c#L1258-1267

- A double free, if detected by the allocator, would cause a program crash. The second free is called in the PRINT_FORMAT_STRING macro.

```
#define PRINT_FORMAT_STRING(temp, format, func) \
    if(strstr(temp, format)) { \
        string = strreplace(temp, format, func(pkg)); \
        free(temp); \
        temp = string; \
    } \
```

Figure 1.2: The PRINT_FORMAT_STRING macro definition

The severity of this finding has been set to low since the first scenario should not be possible because Pacman does not use multiple threads.

This issue was found with the **scan-build** static analyzer but can also be detected with tools such as Valgrind (figure 1.3) or AddressSanitizer (ASan).

```
# valgrind ./pacman -S --print --print-format '%s' valgrind
==2084== Memcheck, a memory error detector
==2084== Copyright (C) 2002-2022, and GNU GPL'd, by Julian Seward et al.
==2084== Using Valgrind-3.21.0 and LibVEX; rerun with -h for copyright info
==2084== Command: ./pacman -S --print --print-format %s valgrind
==2084==
==2084== Invalid read of size 1
==2084==    at 0x484D11D: strstr (vg_replace_strmem.c:1792)
==2084==    by 0x12620B: print_packages (util.c:1267)
==2084==    by 0x11F9DA: sync_prepare_execute (sync.c:817)
==2084==    by 0x11F550: sync_trans (sync.c:728)
==2084==    by 0x11FF72: pacman_sync (sync.c:965)
==2084==    by 0x11B5EB: main (pacman.c:1259)
==2084== Address 0x65e61d0 is 0 bytes inside a block of size 3 free'd
==2084==    at 0x484412F: free (vg_replace_malloc.c:974)
==2084==    by 0x1261F2: print_packages (util.c:1264)
==2084==    by 0x11F9DA: sync_prepare_execute (sync.c:817)
==2084==    by 0x11F550: sync_trans (sync.c:728)
==2084==    by 0x11FF72: pacman_sync (sync.c:965)
==2084==    by 0x11B5EB: main (pacman.c:1259)
==2084== Block was alloc'd at
==2084==    at 0x4841848: malloc (vg_replace_malloc.c:431)
==2084==    by 0x4A183DE: strdup (strdup.c:42)
==2084==    by 0x125ACB: print_packages (util.c:1198)
==2084==    by 0x11F9DA: sync_prepare_execute (sync.c:817)
==2084==    by 0x11F550: sync_trans (sync.c:728)
==2084==    by 0x11FF72: pacman_sync (sync.c:965)
==2084==    by 0x11B5EB: main (pacman.c:1259)
...
==2084== ERROR SUMMARY: 50 errors from 40 contexts (suppressed: 0 from 0)
```

Figure 1.3: Detecting the bug with Valgrind

Exploit Scenario

Pacman starts using multiple threads, uses the `print_packages` function in one thread, and performs an allocation of a similar size to the freed `temp` variable in another thread.

An attacker with the ability to control the contents of the latter allocation exploits the program by manipulating its heap memory through the vulnerable code path.

Recommendations

Short term, add an assignment of `temp = string;` after the `temp` variable is freed in the vulnerable code path in the `print_packages` function. This will prevent the use-after-free issue.

Long term, regularly scan the code with static analyzers like scan-build.

2. Null pointer dereferences

Severity: Informational

Difficulty: Low

Type: Denial of Service

Finding ID: TOB-PACMAN-2

Target: pacman/src/pacman/callback.c:656-660,
pacman/lib/libalpm/util.c:469-481

Description

The `cb_progress` function first checks whether the `pkgname` variable is a null pointer in a ternary operator (highlighted line 1 in figure 2.1) and then may use `pkgname` to format a string (highlighted lines 2 and 3 in figure 2.1). This leads to a crash if `pkgname` is a null pointer.

```
void cb_progress(void *ctx, alpm_progress_t event, const char *pkgname,
                int percent, size_t howmany, size_t current) {
    ...
    len = strlen(opr) + ((pkgname) ? strlen(pkgname) : 0) + 2;    // <--- (1)
    wctr = calloc(len, sizeof(wchar_t));
    /* print our strings to the alloc'ed memory */
    #if defined(HAVE_SWPRINTF)
    wclen = swprintf(wctr, len, L"%s %s", opr, pkgname);          // <--- (2)
    #else
    /* because the format string was simple, we can easily do this without
     * using swprintf, although it is probably not as safe/fast. The max
     * chars we can copy is decremented each time by subtracting the length
     * of the already printed/copied wide char string. */
    wclen = mbstowcs(wctr, opr, len);
    wclen += mbstowcs(wctr + wclen, " ", len - wclen);
    wclen += mbstowcs(wctr + wclen, pkgname, len - wclen);        // <--- (3)
    #endif
```

Figure 2.1: `pacman/src/pacman/callback.c#L656-660`

The severity of this finding has been set to informational because if the `cb_progress` function were called with a null pointer, the program crash would be evident to program users and developers.

An additional case of null pointer dereference is present in the `_alpm_chroot_write_to_child()` function if the `out_cb` argument is null (figure 2.2).

```

typedef ssize_t (*_alpm_cb_io)(void *buf, ssize_t len, void *ctx);

// [...]

static int _alpm_chroot_write_to_child(alpm_handle_t *handle, int fd,
    char *buf, ssize_t *buf_size, ssize_t buf_limit,
    _alpm_cb_io out_cb, void *cb_ctx)
{
    ssize_t nwrite;

    if(*buf_size == 0) {
        /* empty buffer, ask the callback for more */
        if((*buf_size = out_cb(buf, buf_limit, cb_ctx)) == 0) {
            /* no more to write, close the pipe */
            return -1;
        }
    }
}

```

Figure 2.2: *pacman/lib/libalpm/util.c#L469-481*

Recommendations

Short term, modify the `cb_progress` and `_alpm_chroot_write_to_child` functions so that the above-noted pointers are checked for a null value before they are used. In the event of a null pointer, abort without dereferencing.

Long term, use static analysis tools to detect cases where pointers are dereferenced without a preceding null check.

3. Allocation failures can lead to memory leaks or null pointer dereferences

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-PACMAN-3

Target: src/pacman/conf.c, MR #96: lib/libalpm/alpm_list.c,
lib/libalpm/be_sync.c

Description

There are a few code paths where allocation failures can lead to further memory leaks or null pointer dereferences:

- If the `strdup(path)` function's allocation fails in the `setdefaults` function (highlighted lines 1 and 2 in figure 3.1), then the memory pointed to by the `rootdir` variable (highlighted line 2 in figure 3.1) would be leaked. This is because the `SETDEFAULT` macro would enter its error path and return `-1` (highlighted line 3 in figure 3.1), not freeing the previously allocated memory.

```
int setdefaults(config_t *c) {
    alpm_list_t *i;

#define SETDEFAULT(opt, val)          \
    if(!opt) {                        \
        opt = val;                    \
        if(!opt) { return -1; }      \
    }                                // (3)

    if(c->rootdir) {
        char* rootdir = strdup(c->rootdir); // (2)
        ...
        char path[PATH_MAX];
        if(!c->dbpath) {
            snprintf(path, PATH_MAX, "%s/%s", rootdir, &DBPATH[1]);
            SETDEFAULT(c->dbpath, strdup(path)); // (1)
        }
        if(!c->logfile) {
            snprintf(path, PATH_MAX, "%s/%s", rootdir, &LOGFILE[1]);
            SETDEFAULT(c->logfile, strdup(path)); // (1)
        }
    }
}
```

Figure 3.1: *pacman/src/pacman/conf.c#L1139–1153*

- The `alpm_list_equal_ignore_order` function added in MR #96 fails to check whether the `calloc` function returns a non-null value (figure 3.2). If `calloc`

returned NULL, this would lead to a null pointer dereference later in the function (line 534 of figure 3.2).

```
511     int SYMEXPORT alpm_list_equal_ignore_order(const alpm_list_t *left,
512                                                const alpm_list_t *right, alpm_list_fn_cmp fn)
513     {
514         const alpm_list_t *l = left;
515         const alpm_list_t *r = right;
516         int *matched;
517
518         if((l == NULL) != (r == NULL)) {
519             return 0;
520         }
521
522         if(alpm_list_count(l) != alpm_list_count(r)) {
523             return 0;
524         }
525
526         matched = calloc(alpm_list_count(right), sizeof(int));
527
528         for(l = left; l; l = l->next) {
529             int found = 0;
530             int n = 0;
531
532             for(r = right; r; r = r->next, n++) {
533                 /* make sure we don't match the same value twice */
534                 if(matched[n]) {
535                     continue;
536                 }
```

Figure 3.2: MR #96: lib/libalpm/alpm_list.c#L511-536

- In the `_alpm_validate_filename()` function, the `strlen(filename)` function can be called with a null pointer if the `READ_AND_STORE(pkg->filename)` execution fails to allocate memory through the `STRDUP` macro (figure 3.3).

```
#define READ_AND_STORE(f) do { \
    READ_NEXT(); \
    STRDUP(f, line, goto error); \
} while(0)

#define STRDUP(r, s, action) do { \
    if(s != NULL) { \
        r = strdup(s); \
        if(r == NULL) { \
            _alpm_alloc_fail(strlen(s)); \
            action; \
        } \
    } \
    else { r = NULL; } \
} while(0)
```

```
READ_AND_STORE(pkg->filename);  
if(_alpm_validate_filename(db, pkg->name, pkg->filename) < 0) { ... }
```

Figure 3.3: `pacman/lib/libalpm/be_sync.c#L591-595`

The severity of this finding has been set to informational because if an allocation failed, the program would likely stop functioning properly since it would fail to allocate any more memory anyway.

The first part of this issue (pertaining to the `conf.c` file, rather than the `alpm_list.c` file) was found with the `scan-build` static analyzer.

Recommendations

Short term, to fix the memory leaks and null pointer dereferences, add null-pointer checks after the allocations and before the dereferences.

Long term, regularly scan the code with static analyzers like `scan-build` to detect missing checks of these types.

4. Buffer overflow read in string_length utility function

Severity: Undetermined

Difficulty: High

Type: Data Validation

Finding ID: TOB-PACMAN-4

Target: src/pacman/util.c

Description

The `string_length` utility function (figure 4.1) skips ANSI color codes when computing the length. When a string includes the `"\033"` byte that starts the ANSI color code sequence but does not have the `"m"` character that ends it, the function will read memory past the end of the string, causing a buffer overflow read.

This can lead to a program crash or other issues, depending on how the function is used.

```
static size_t string_length(const char *s) {
    int len;
    wchar_t *wcstr;

    if(!s || s[0] == '\0') {
        return 0;
    }
    if(strstr(s, "\033")) {
        char* replaced = malloc(sizeof(char) * strlen(s));
        int iter = 0;
        for(; *s; s++) {
            if(*s == '\033') {
                while(*s != 'm') {
                    s++;
                }
            } else {
                replaced[iter] = *s;
                iter++;
            }
        }
        replaced[iter] = '\0';
    }
}
```

Figure 4.1: *pacman/src/pacman/util.c#L452-473*

Recommendations

Short term, so that the loop terminates at a null character, change the conditional within the `while()` statement in the `string_length` function to the following:

```
while(*s && *s != 'm')
```

Long term, implement a fuzzing harness for the `string_length` function to ensure it does not contain any bugs. An example harness code can be found in figure 4.2; this can be compiled and run using the following commands:

```
clang -fsanitize=fuzzer,address main.c -ggdb -o fuzzer
./fuzzer
```

```
#define _XOPEN_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <string.h>
#include <wchar.h>

static size_t string_length(const char *s) { ... }

int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
    if (Size == 0) return 0;

    // Prepare a null terminated string
    char* x = malloc(Size+1);
    memcpy(x, Data, Size);
    x[Size] = 0;

    string_length(x);

    free(x);
    return 0;
}
```

Figure 4.2: Example fuzzing harness that uses *libFuzzer* to test the `string_length` function

Figure 4.3 shows an example output of such a fuzzer. We also implemented this harness as part of the Pacman codebase, as detailed in [appendix D](#).

```
$ clang -fsanitize=fuzzer,address main.c -ggdb -o fuzzer
$ ./fuzzer
INFO: Running with entropic power schedule (0xFF, 100).
INFO: Seed: 1790240281
INFO: Loaded 1 modules   (12 inline 8-bit counters): 12 [0x56046acc5fc0,
0x56046acc5fcc),
INFO: Loaded 1 PC tables (12 PCs): 12 [0x56046acc5fd0,0x56046acc6090),
INFO: -max_len is not provided; libFuzzer will not generate inputs larger than 4096
bytes
INFO: A corpus is not provided, starting from an empty corpus
#2      INITED cov: 4 ft: 5 corp: 1/1b exec/s: 0 rss: 30Mb
...
#173   REDUCE cov: 5 ft: 6 corp: 2/2b lim: 4 exec/s: 0 rss: 31Mb L: 1/1 MS: 1
=====
```

```

==2873139==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x602000006b53
at pc 0x56046ac84a76 bp 0x7ffd09e07ef0 sp 0x7ffd09e07ee8
READ of size 1 at 0x602000006b53 thread T0
    #0 0x56046ac84a75 in string_length /fuzz/main.c:21:11
    #1 0x56046ac8483d in LLVMFuzzerTestOneInput /fuzz/main.c:56:2
    #2 0x56046abad383 in fuzzer::Fuzzer::ExecuteCallback(unsigned char const*,
unsigned long) (/fuzz/fuzzer+0x3e383) (BuildId:
65f386451dc943b740358c52379831570eef52be)
...

0x602000006b53 is located 0 bytes to the right of 3-byte region
[0x602000006b50,0x602000006b53)
allocated by thread T0 here:
    #0 0x56046ac499fe in malloc (/fuzz/fuzzer+0xda9fe) (BuildId:
65f386451dc943b740358c52379831570eef52be)
    #1 0x56046ac847db in LLVMFuzzerTestOneInput /root/fuz/main.c:53:12
...

SUMMARY: AddressSanitizer: heap-buffer-overflow /root/fuz/main.c:21:11 in
string_length
...

```

Figure 4.3: Output from the fuzzer from figure 4.2

5. Undefined behavior or potential null pointer dereferences

Severity: **Undetermined**

Difficulty: **Undetermined**

Type: Data Validation

Finding ID: TOB-PACMAN-5

Target: Multiple codepaths

Description

There are a few code paths where a null pointer dereference or undefined behavior may occur if certain conditions are met. Those issues can be detected with the scan-build static analyzer or by building and running Pacman with UndefinedBehaviorSanitizer (UBSan). The scan-build results were shared along with this report.

One of the code paths found by scan-build is in the `lib/libalpm/remove.c` file. The `closedir(dir)` function may be called with a null pointer when the condition that calls `regcomp(...)` is true (figure 5.1). This is undefined behavior since the `closedir` function argument is marked as non-null.

```
static void shift_pacsave(alpm_handle_t *handle, const char *file) {
    DIR *dir = NULL;
    ...
    if(regcomp(&reg, regstr, REG_EXTENDED | REG_NEWLINE) != 0) {
        goto cleanup;
    }

    dir = opendir(dirname); // <-- the dir was only modified here
    ...
cleanup:
    free(dirname);
    closedir(dir);
}
```

Figure 5.1: `pacman/lib/libalpm/remove.c#L349-423`

Another case is in the `mount_point_list` function (figure 5.2). If the `STRDUP` macro is executed with a `mnt->mnt_dir` null pointer, then the `strlen(mp->mount_dir)` call will take a null pointer.

```
static alpm_list_t *mount_point_list(alpm_handle_t *handle) {
    ...
    #if defined(HAVE_GETMNTENT) && defined(HAVE_MNTENT_H)
        ...
        while((mnt = getmntent(fp))) {
            CALLOC(mp, 1, sizeof(alpm_mountpoint_t), RET_ERR(handle,
ALPM_ERR_MEMORY, NULL));
            STRDUP(mp->mount_dir, mnt->mnt_dir, free(mp); RET_ERR(handle,
ALPM_ERR_MEMORY, NULL));
            mp->mount_dir_len = strlen(mp->mount_dir);
        }
    }
}
```

Figure 5.2: *pacman/lib/libalpm/diskspace.c#L95-116*

In addition, figure 5.3 shows a run of Pacman with UBSan that detects other cases of this issue.

```
# CFLAGS=-fsanitize=address,undefined LDFLAGS=-fsanitize=address,undefined meson
setup sanitize
# cd sanitize
# CFLAGS=-fsanitize=address,undefined LDFLAGS=-fsanitize=address,undefined meson
compile
# ./pacman -Syuu
:: Synchronizing package databases...
  core downloading...
  extra downloading...
:: Starting full system upgrade...
../lib/libalpm/util.c:1149:9: runtime error: null pointer passed as argument 1,
which is declared to never be null
../lib/libalpm/util.c:1151:10: runtime error: null pointer passed as argument 1,
which is declared to never be null
../lib/libalpm/util.c:1192:4: runtime error: null pointer passed as argument 2,
which is declared to never be null
...
:: Proceed with installation? [Y/n] Y
...
```

Figure 5.3: *Running Pacman with UBSan*

Recommendations

Short term, fix the cases where functions marked with non-null arguments are called with null pointers by putting the function calls inside `if` statements that perform null checks.

Long term, regularly test Pacman with UBSan and scan its codebase with static analyzers such as scan-build.

6. Undefined behavior from use of atoi

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-PACMAN-6

Target: lib/libalpm/be_local.c, src/pacman/pacman.c

Description

The `atoi` function is used to convert strings to integers when parsing local database files and command-line arguments (figures 6.1 and 6.2). The behavior of `atoi` is undefined in the case where the input string is not a valid formatted number or in the case of an overflow.

The severity of this finding has been set to informational because, in practice, `atoi` will typically return a dummy value, such as 0 or -1, in the case of an incorrect input or an overflow.

```
} else if(strcmp(line, "%REASON%") == 0) {  
    READ_NEXT();  
    info->reason = (alpm_pkgreason_t)atoi(line);
```

Figure 6.1: Use of `atoi` ([lib/libalpm/be_local.c#L774-776](#))

```
case OP_ASK:  
    config->noask = 1;  
    config->ask = (unsigned int)atoi(optarg);  
    break;  
...  
case OP_DEBUG:  
    /* debug levels are made more 'human readable' than using a raw logmask  
     * here, error and warning are set in config_new, though perhaps a  
     * --quiet option will remove these later */  
    if(optarg) {  
        unsigned short debug = (unsigned short)atoi(optarg);  
        switch(debug) {  
            case 2:  
                config->logmask |= ALPM_LOG_FUNCTION;  
                __attribute__((fallthrough));  
            case 1:  
                config->logmask |= ALPM_LOG_DEBUG;  
                break;  
            default:  
                pm_printf(ALPM_LOG_ERROR, _("%s' is not a valid debug  
level\n"),  
                        optarg);
```

```
                                return 1;
        }
    } else {
        config->logmask |= ALPM_LOG_DEBUG;
    }
    /* progress bars get wonky with debug on, shut them off */
    config->noprogressbar = 1;
    break;
```

Figure 6.2: Uses of `atoi` ([src/pacman/pacman.c#L382-430](#))

Recommendations

Short term, use the `strtol` function instead of `atoi`. After calling `strtol`, check the `errno` value for a failed conversion. Make sure to perform bounds checking when casting the long value returned by `strtol` down to an `int`.

7. Database parsers fail silently if an option is not recognized

Severity: Informational

Difficulty: High

Type: Data Validation

Finding ID: TOB-PACMAN-7

Target: `lib/libalpm/be_sync.c`, `lib/libalpm/be_local.c`

Description

The `sync_db_read` and `local_db_read` functions, which are responsible for parsing sync database files and local database files respectively, fail silently if an option is not recognized. This can cause a configuration option to not be set, which may cause issues if, for example, the local installation of Pacman is out of date and does not support newly added configuration options.

Exploit Scenario

Support for SHA-3 hash verification is added, along with a corresponding configuration option, `%SHA3SUM%`. Older installations of Pacman, which do not support this configuration option, will ignore it. This causes package hashes to not be verified.

Recommendations

Short term, add default behavior in the `sync_db_read` and `local_db_read` functions for when a configuration option is not recognized. Unrecognized options should cause a log message or an error.

8. Cache cleaning function may delete the wrong files

Severity: Informational

Difficulty: High

Type: Data Validation

Finding ID: TOB-PACMAN-8

Target: `src/pacman/sync.c`

Description

In the `sync_cleancache` function, a path is constructed for deletion using the `snprintf` function. A maximum path length of `PATH_MAX` is given (on Linux, this value is 4096 characters). However, there is no check to ensure that the path created by `snprintf` was not cut short by the limit. This can lead to deletion of a different path than intended.

The severity of this finding has been set to informational because it is highly unlikely that Pacman would use a path this long in practice.

```
/* build the full filepath */
snprintf(path, PATH_MAX, "%s%s", cachedir, ent->d_name);

/* short circuit for removing all files from cache */
if(level > 1) {
    ret += unlink_verbose(path, 0);
    continue;
}
```

Figure 8.1: Path creation and file deletion ([pacman/src/pacman/sync.c#L241-248](#))

Recommendations

Short term, add a check that compares the value returned by `snprintf` and ensures that it is less than `PATH_MAX`.

9. Integer underflow in a length check leads to out-of-bounds read in alpm_extract_keyid

Severity: Undetermined

Difficulty: Undetermined

Type: Data Validation

Finding ID: TOB-PACMAN-9

Target: lib/libalpm/signing.c

Description

The `alpm_extract_keyid` function (figure 9.1) contains an out-of-bounds read issue due to an integer underflow in the `length_check` function when a specifically crafted input is provided (figure 9.2).

```
int SYMEXPORT alpm_extract_keyid(alpm_handle_t *handle, const char *identifier,
    const unsigned char *sig, const size_t len, alpm_list_t **keys) {
    size_t pos, blen, hlen, ulen;
    pos = 0;

    while(pos < len) {
        if(!(sig[pos] & 0x80)) { ... - return signature format error }

        if(sig[pos] & 0x40) {
            /* new packet format */
            if(length_check(len, pos, 1, handle, identifier) != 0) {
                return -1;
            }
            pos = pos + 1;
        }
    }
}
```

Figure 9.1: *pacman/lib/libalpm/signing.c#L1101-1223*

```
/* Check to avoid out of boundary reads */
static size_t length_check(size_t length, size_t position, size_t a,
    alpm_handle_t *handle, const char *identifier) {
    if( a == 0 || length - position <= a) {
        _alpm_log(handle, ALPM_LOG_ERROR,
            _("%s: signature format error\n"), identifier);
        return -1;
    } else {
        return 0;
    }
}
```

Figure 9.2: *pacman/lib/libalpm/signing.c#L1043-1054*

The `length_check` function is used to confirm whether advancing a position (`pos`) index is safe. It is used by `alpm_extract_keyid`, for example, in the following way:

```
length_check(len, pos, 2, handle, identifier)
```

The `len` variable is the length of the signature buffer (`sig`), and `pos` is an index in that buffer. However, the `pos` index can be bigger than the `len` variable, and when that happens, the length-position computation in the `length_check` function underflows and the function returns 0, leading to the out-of-bounds read.

We found this issue by fuzzing the `alpm_extract_keyid` function. The fuzzing harness code is included in [appendix D](#).

Recommendations

Short term, fix the integer underflow issue in the `length_check` function by changing the highlighted `if` statement in figure 9.2 to the following:

```
if( a == 0 || (position <= length && length - position <= a))
```

This will prevent out-of-bounds reads in the `alpm_extract_keyid` function.

Long term, fuzz the Pacman functions, as shown in [appendix D](#), for example.

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Code Maturity Categories

The following tables describe the code maturity categories and rating criteria used in this document.

Code Maturity Categories	
Category	Description
Arithmetic	The proper use of mathematical operations and semantics
Auditing	The use of event auditing and logging to support monitoring
Authentication / Access Controls	The use of robust access controls to handle identification and authorization and to ensure safe interactions with the system
Complexity Management	The presence of clear structures designed to manage system complexity, including the separation of system logic into clearly defined functions
Configuration	The configuration of system components in accordance with best practices
Cryptography and Key Management	The safe use of cryptographic primitives and functions, along with the presence of robust mechanisms for key generation and distribution
Data Handling	The safe handling of user inputs and data processed by the system
Documentation	The presence of comprehensive and readable codebase documentation
Maintenance	The timely maintenance of system components to mitigate risk
Memory Safety and Error Handling	The presence of memory safety and robust error-handling mechanisms
Testing and Verification	The presence of robust testing procedures (e.g., unit tests, integration tests, and verification methods) and sufficient test coverage

Rating Criteria	
Rating	Description
Strong	No issues were found, and the system exceeds industry standards.
Satisfactory	Minor issues were found, but the system is compliant with best practices.
Moderate	Some issues that may affect system safety were found.

Weak	Many issues that affect system safety were found.
Missing	A required component is missing, significantly affecting system safety.
Not Applicable	The category is not applicable to this review.
Not Considered	The category was not considered in this review.
Further Investigation Required	Further investigation is required to reach a meaningful conclusion.

C. Code Quality Findings

The following recommendations are not associated with specific vulnerabilities. However, they enhance code readability and may prevent the introduction of vulnerabilities in the future.

- **Remove the `if (fd >= 0)` condition in the `_alpm_pkg_load_internal` function since it is always true.** This is because if `fd` is less than 0, the function returns `NULL` in a previous condition.

```
alpm_pkg_t *_alpm_pkg_load_internal(alpm_handle_t *handle,
    const char *pkgfile, int full) {
    int ret, fd;
    ...
    fd = _alpm_open_archive(handle, pkgfile, &st, &archive,
ALPM_ERR_PKG_OPEN);
    if(fd < 0)
        ...
        return NULL;
    }
    ...
error:
    _alpm_pkg_free(newpkg);
    _alpm_archive_read_free(archive);
    if(fd >= 0) {
        close(fd);
    }

    return NULL;
}
```

Figure C.1: *pacman/lib/libalpm/be_package.c#L569-688*

- **Use the `strdup` function to duplicate a string in the `clean_filename` function.** This can be done instead of computing the string length, allocating memory, and copying the filename with the `memcpy` function.

```
static char *clean_filename(const char *filename) {
    int len = strlen(filename);
    char *p;
    char *fname = malloc(len + 1);
    memcpy(fname, filename, len + 1);
}
```

Figure C.2: *pacman/src/pacman/callback.c#L755-760*

- **Refactor the dead assignment to the `curlerr` variable in the `curl_check_finished_download` function.** Either the assignment should be removed or there should be code that acts upon its value.


```

static int curl_check_finished_download(alpm_handle_t *handle, CURLM *curlm,
CURLMsg *msg, const char *localpath, int *active_downloads_num) {
    ...
    CURLcode curlerr;
    ...
    case CURLE_ABORTED_BY_CALLBACK:
        /* handle the interrupt accordingly */
        if(dload_interrupted == ABORT_OVER_MAXFILESIZE) {
            curlerr = CURLE_FILESIZE_EXCEEDED;
            payload->unlink_on_fail = 1;
            handle->pm_errno = ALPM_ERR_LIBCURL;
            _alpm_log(handle, ALPM_LOG_ERROR,
                _("failed retrieving file '%s' from %s
: expected download size exceeded\n"),
                    payload->remote_name, hostname);
            server_soft_error(handle, payload->fileurl);
        }
        goto cleanup;
    ...
cleanup:
    ... // <-- code that does not use the curlerr variable
    return ret;
}

```

Figure C.3: *pacman/lib/libalpm/dload.c#L535–546*

- **Remove the unused `r` variable, and the assignment to it in the following `if` statement, from the `_cache_mtree_open` function.** Alternatively, if it is intended, use the value of `r` within the `if` condition.

```

static struct archive *_cache_mtree_open(alpm_pkg_t *pkg) {
    int r;
    ...
    if((r = _alpm_archive_read_open_file(mtree, mtfile, ALPM_BUFFER_SIZE)))
    {
        _alpm_log(pkg->handle, ALPM_LOG_ERROR, _("error while reading
file %s: %s\n"),
                    mtfile, archive_error_string(mtree));
        _alpm_archive_read_free(mtree);
        GOTO_ERR(pkg->handle, ALPM_ERR_LIBARCHIVE, error);
    }

    free(mtfile);
    return mtree;

error:
    free(mtfile);
    return NULL;
}

```

Figure C.4: *pacman/lib/libalpm/be_local.c#L251–284*

- **Return an error if the call to malloc fails in the `alpm_list_add_sorted` function.** Currently, the function returns the existing list, even though it fails to insert the element as expected. This function is currently unused, so this does not yet pose a security concern.

```
add = malloc(sizeof(alpm_list_t));
if(add == NULL) {
    return list;
}
```

Figure C.5: `pacman/lib/libalpm/alpm_list.c#L115-118`

- **Restore the list variable to its original state before returning from the `alpm_list_reverse` function.** In the beginning of the function, the `list->prev` member is backed up and then modified. However, in the case of an error, this backup is not restored, leaving the list in an invalid state.

```
alpm_list_t SYMEXPORT *alpm_list_reverse(alpm_list_t *list) {
    const alpm_list_t *lp;
    alpm_list_t *newlist = NULL, *backup;

    if(list == NULL) {
        return NULL;
    }

    lp = alpm_list_last(list);
    /* break our reverse circular list */
    backup = list->prev;
    list->prev = NULL;

    while(lp) {
        if(alpm_list_append(&newlist, lp->data) == NULL) {
            alpm_list_free(newlist);
            return NULL;
        }
        lp = lp->prev;
    }
    list->prev = backup; /* restore tail pointer */
    return newlist;
}
```

Figure C.6: `pacman/lib/libalpm/alpm_list.c#L403-426`

- **Rename the type variable to event in the `alpm_list_reverse` function.** When an `ALPM_SANDBOX_CB_DOWNLOAD`-type payload is sent over a pipe in MR #23 (from the `_alpm_sandbox_cb_dl` function to the `_alpm_sandbox_process_cb_download` function), a variable called `event` is sent through the pipe and received into a variable called `type`. This can cause confusion when reading the sending and receiving code.

- **Rework the `had_error` variable in the `curl_download_internal_sandboxed` function.** The variable will always be set to `true` by the time the loop shown in figure C.7 exits. This is because every `break` statement is accompanied by a statement setting `had_error` to `true`. This means that the variable does not track any useful information.

```
bool had_error = false;
while(true) {
    _alpm_sandbox_callback_t callback_type;
    ssize_t got = read(callbacks_fd[0], &callback_type,
sizeof(callback_type));
    if(got < 0 || (size_t)got != sizeof(callback_type)) {
        had_error = true;
        break;
    }

    if(callback_type == ALPM_SANDBOX_CB_DOWNLOAD) {
        if(!_alpm_sandbox_process_cb_download(handle, callbacks_fd[0]))
        {
            had_error = true;
            break;
        }
    }
    else if(callback_type == ALPM_SANDBOX_CB_LOG) {
        if(!_alpm_sandbox_process_cb_log(handle, callbacks_fd[0])) {
            had_error = true;
            break;
        }
    }
}

if(had_error) {
    kill(pid, SIGTERM);
}
```

Figure C.7: MR #23: `pacman/lib/libalpm/dload.c#L974-1000`

- **Verify the `%REASON%` field before casting it to an `alpm_pkgreason_t` enum value in the `local_db_read` function.** Otherwise, the field may contain a value that is a valid integer but not a valid `alpm_pkgreason_t` value.

```
} else if(strcmp(line, "%REASON%") == 0) {
    READ_NEXT();
    info->reason = (alpm_pkgreason_t)atoi(line);
}
```

Figure C.8: `pacman/lib/libalpm/be_local.c#L774-776`

- **Correct the log message at the end of the `curl_download_internal` function.** The message incorrectly states the value returned by the function.

```
_alpm_log(handle, ALPM_LOG_DEBUG, "curl_download_internal return code is  
%d\n", err);  
return err ? -1 : updated ? 0 : 1;
```

Figure C.9: [pacman/lib/libalpm/dload.c#L937-938](#)

- **Refactor the ALPM public functions from returning an `int` to returning a `status` type.** This new type could be a typedef for an `int`. Such a change would make it easier to perform static analysis to find all functions that return the typedef and ensure that the callers check for errors.

[Errors]

The library provides a global variable `pm_errno`.
It aims at being to the library what `errno` is for C system calls.

Almost all public library functions are returning an integer value: 0 indicating success, -1 indicating a failure.

If -1 is returned, the variable `pm_errno` is set to a meaningful value
Wise frontends should always care for these returned values.

Note: the helper function `alpm_strerror()` can also be used to translate one specified error code into a more friendly sentence, and `alpm_strerrorlast()` does the same for the last error encountered (represented by `pm_errno`).

Figure C.10: [pacman/README?plain=1#L144-156](#)

D. Fuzzing Pacman Code

During the audit, Trail of Bits used fuzzing, an automated testing technique in which code paths are executed with random data to find bugs resulting from the incorrect handling of unexpected data. We used **libFuzzer**, an in-process coverage-guided fuzzer, and we extended the Pacman build system with new executables to fuzz certain code paths. This helped us find the issues detailed in findings **TOB-PACMAN-4** and **TOB-PACMAN-9**.

We implemented fuzzing harnesses for the following:

- The **string_length** function
- The **wordsplit** function
- Parsing of configuration files through the **parseconfigfile** function
- The extraction of keys from signature data through the **alpm_extract_keyid** function

We also modified the `meson.build` file so that all the files were built with **ASan** (the `-fsanitize=address` compiler and linker flag), which helps detect more bugs. To build the harnesses and run them, we leveraged the following commands:

```
CC=clang meson setup fuzz
cd fuzz
CC=clang meson compile <harness, e.g., fuzz_alpm_extract_keyid>
./<harness binary>
```

We used the clang compiler because, in our case, we performed fuzzing in an Arch Linux Docker container, and the GCC compiler did not support the `-fsanitize=fuzzer` flag that enables the libFuzzer fuzzing framework.

The implemented code is shown in figure D.1 and will also be sent as a merge request against the Pacman repository after the final readout of this report.

Fuzzing Harness Notes

The following are notes about the changes and harnesses we developed:

- The `add_project_arguments` line added to the `meson.build` file enables ASan for all build binaries. This is suboptimal and must be refactored because release binaries should not be built with ASan. Instead, a new target or project should be specified to build all artifacts for fuzzing (libraries, all dependencies, and fuzzing harness code) with enabled sanitizers.

- None of the external dependencies are built with ASan or UBSan. This may cause false positive crashes when new harnesses are developed that leverage the code paths of those dependencies, or it may lead to not detecting valid bugs.
- We encountered an issue with including headers from the `src/pacman` directory in the `fuzz_parseconfigfile` and `fuzz_string_length` harnesses. We worked around this issue by providing the fuzzed function declarations in the harnesses themselves, but this issue should be fixed so that the headers are included properly. Due to this issue, we also added the `extern void *config;` global variable to the harness code.
- The `fuzz_wordsplit` harness can be refactored to free its resources via the `wordsplit_free` function.
- The `fuzz_alpm_extract_keyid` harness does not set proper handle or filename arguments. Setting these arguments may leverage more code paths in the harness.
- The `fuzz_parseconfigfile` harness is far from ideal: the generated input may include other files from the filesystem to be parsed by the code, which is nondeterministic. The solution to this could be one of the following:
 - Either use `chroot` or mount namespaces so that the fuzzer works in an isolated filesystem with no other files included.
 - Change the harness so that it generates only semi-valid configuration files.
- We added an `#ifndef FUZZING_PACMAN` declaration to remove the `main` function of Pacman for the fuzzing harnesses that need the `src/pacman` code. Otherwise, the linking of the harness would fail due to multiple definitions of the `main` symbol.

Recommendations and Further Work

Going forward, we recommend the Pacman team take the following steps:

- Refactor the build system to better support the building of fuzzing harnesses (instead of setting global arguments as we did).
- Extend the build system so it also builds all the dependencies' code with sanitizers enabled.
- Test and fuzz the code with other sanitizers that we have not tried here (e.g., MemorySanitizer or ThreadSanitizer in case threads are ever used in Pacman) enabled.
- Implement more fuzzing harnesses—for example, for the `dload_parseheader_cb` function and other functionalities that parse untrusted data.

- Fuzz Pacman regularly with each release. This can be done by integrating the fuzzing harnesses into the [oss-fuzz project](#), which allows free fuzzing of open-source projects. However, be aware that the company behind the oss-fuzz project, Google, will know about the discovered vulnerabilities first.

```
diff --git a/meson.build b/meson.build
index 43705338..bfeca3af 100644
--- a/meson.build
+++ b/meson.build
@@ -14,6 +14,8 @@ libalpm_version = '13.0.1'

cc = meson.get_compiler('c')

+add_project_arguments(['-fsanitize=address', '-fno-omit-frame-pointer',
+'-ggdb', '-O0'], language : 'c')
+
# commandline options
PREFIX = get_option('prefix')
DATAROOTDIR = join_paths(PREFIX, get_option('datarootdir'))
@@ -305,6 +307,8 @@ subdir('src/pacman')
subdir('src/util')
subdir('scripts')

+subdir('src/fuzzing')
+
# Internationalization
if get_option('i18n')
i18n = import('i18n')
@@ -396,6 +400,45 @@ executable(
install : true,
)

+# Note: fuzz targets below must be built with Clang compiler for the
+fsanitize=fuzzer flag
+executable(
+  'fuzz_wordsplit',
+  fuzz_wordsplit_sources,
+  include_directories : includes,
+  link_with : [libcommon],
+  dependencies : [],
+  c_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
+'-fno-omit-frame-pointer'],
+  link_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
+'-fno-omit-frame-pointer'],
+)
+
+executable(
+  'fuzz_string_length',
+  [fuzz_string_length_sources, pacman_sources],
+  include_directories : includes,
+  link_with : [libalpm_a, libcommon],
+  dependencies : [],
```

```

+   c_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer', '-DFUZZING_PACMAN'],
+   link_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer'],
+)
+executable(
+   'fuzz_alpm_extract_keyid',
+   [fuzz_alpm_extract_keyid_sources, pacman_sources],
+   include_directories : includes,
+   link_with : [libalpm_a, libcommon],
+   dependencies : [],
+   c_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer', '-DFUZZING_PACMAN'],
+   link_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer'],
+)
+executable(
+   'fuzz_parseconfigfile',
+   [fuzz_parseconfigfile_sources, pacman_sources],
+   include_directories : includes,
+   link_with : [libalpm_a],
+   dependencies : [],
+   c_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer', '-DFUZZING_PACMAN'],
+   link_args : ['-fsanitize=fuzzer,address', '-ggdb', '-O0',
'-fno-omit-frame-pointer'],
+)
+
+   foreach wrapper : script_wrappers
+       cdata = configuration_data()
+       cdata.set_quoted('BASH', BASH.full_path())
diff --git a/src/fuzzing/fuzz_alpm_extract_keyid.c
b/src/fuzzing/fuzz_alpm_extract_keyid.c
new file mode 100644
index 00000000..febbd57a
--- /dev/null
+++ b/src/fuzzing/fuzz_alpm_extract_keyid.c
@@ -0,0 +1,26 @@
+#define _XOPEN_SOURCE
+#include <stdio.h>
+#include <stdlib.h>
+#include <stdint.h>
+#include <string.h>
+#include <wchar.h>
+
+/* libalpm */
+#include "alpm.h"
+#include "alpm_list.h"
+#include "handle.h"
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size);
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {

```



```

+     if (Size == 0)
+         return 0;
+
+     alpm_handle_t handle;          // TODO/FIXME?
+     const char* filename = "/dev/null"; // TODO/FIXME?
+
+     alpm_list_t *keys = NULL;
+     alpm_extract_keyid(&handle, filename, /* sig */ Data, /* len */ Size,
&keys);
+
+     return 0;
+ }
diff --git a/src/fuzzing/fuzz_parseconfigfile.c
b/src/fuzzing/fuzz_parseconfigfile.c
new file mode 100644
index 00000000..4746141d
--- /dev/null
+++ b/src/fuzzing/fuzz_parseconfigfile.c
@@ -0,0 +1,43 @@
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#define _GNU_SOURCE
#include <sys/mman.h>
#include <unistd.h>
+
+// TODO/FIXME: Fix the util.h include
+// #include "conf.h"
+// And remove that function header from here
+int parseconfigfile(const char *s);
+extern void *config;
+void *config_new(void);
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size);
+
+// TODO/FIXME: This fuzzer should always be run from a chroot
+// without any other files in it; otherwise the configfile may refer
+// to other files
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
+     static void* config_object = 0;
+
+     // TODO/FIXME: The harness needs to be run with -detect_leaks=0
+     // because the config object here is detected as a leak
+     if (!config_object) {
+         config = config_object = config_new();
+     }
+
+     if (Size == 0)
+         return 0;
+
+     int fd = memfd_create("input", 0); // create an in-memory file we can
have path to
+     write(fd, Data, Size);

```

```

+
+   char path[64] = {0};
+   sprintf(path, "/proc/self/fd/%d", fd);
+
+   parseconfigfile(path);
+
+   close(fd);
+
+   return 0;
+}
diff --git a/src/fuzzing/fuzz_string_length.c b/src/fuzzing/fuzz_string_length.c
new file mode 100644
index 00000000..8991b476
--- /dev/null
+++ b/src/fuzzing/fuzz_string_length.c
@@ -0,0 +1,26 @@
+#include <stdio.h>
+#include <stdlib.h>
+#include <string.h>
+
+// TODO/FIXME: Fix the util.h include
+// #include "util.h"
+// And remove that function header from here
+size_t string_length(const char *s);
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size);
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
+    if (Size == 0)
+        return 0;
+
+    // Prepare a null terminated string
+    char* cstring = malloc(Size+1);
+    memcpy(cstring, Data, Size);
+    cstring[Size] = 0;
+
+    string_length(cstring);
+
+    free(cstring);
+
+    return 0;
+}
diff --git a/src/fuzzing/fuzz_wordsplit.c b/src/fuzzing/fuzz_wordsplit.c
new file mode 100644
index 00000000..e2e10210
--- /dev/null
+++ b/src/fuzzing/fuzz_wordsplit.c
@@ -0,0 +1,36 @@
+#define _XOPEN_SOURCE
+#include <stdio.h>
+#include <stdlib.h>
+#include <stdint.h>

```

```

+
+#include "util-common.h"
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size);
+
+int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
+    if (Size == 0)
+        return 0;
+
+    // Prepare a null terminated string
+    char* cstring = malloc(Size+1);
+    memcpy(cstring, Data, Size);
+    cstring[Size] = 0;
+
+    char** ptr = wordsplit(cstring);
+
+    // Free the memory allocated by wordsplit
+    if (ptr) {
+        int i = 0;
+        char* p = ptr[i++];
+        while (p) {
+            free(p);
+            p = ptr[i++];
+        }
+        free(ptr);
+    }
+
+    // Free the allocated cstring
+    free(cstring);
+
+    return 0;
+}
diff --git a/src/fuzzing/meson.build b/src/fuzzing/meson.build
new file mode 100644
index 00000000..9a8555c2
--- /dev/null
+++ b/src/fuzzing/meson.build
@@ -0,0 +1,15 @@
+fuzz_wordsplit_sources = files(''
+ fuzz_wordsplit.c
+ ''.split())
+
+fuzz_string_length_sources = files(''
+ fuzz_string_length.c
+ ''.split())
+
+fuzz_alpm_extract_keyid_sources = files(''
+ fuzz_alpm_extract_keyid.c
+ ''.split())
+
+fuzz_parseconfigfile_sources = files(''
+ fuzz_parseconfigfile.c
+ ''.split())

```

```

\ No newline at end of file
diff --git a/src/pacman/pacman.c b/src/pacman/pacman.c
index e5c6e420..77c88392 100644
--- a/src/pacman/pacman.c
+++ b/src/pacman/pacman.c
@@ -1079,6 +1079,7 @@ static void cl_to_log(int argc, char *argv[])
    }
}

+#ifndef FUZZING_PACMAN
/** Main function.
 * @param argc
 * @param argv
@@ -1273,3 +1274,4 @@ int main(int argc, char *argv[])
/* not reached */
return EXIT_SUCCESS;
}
+#endif //FUZZING_PACMAN
diff --git a/src/pacman/util.c b/src/pacman/util.c
index 5d42a6e9..a41c9e5e 100644
--- a/src/pacman/util.c
+++ b/src/pacman/util.c
@@ -449,7 +449,7 @@ static char *concat_list(alpm_list_t *lst, formatfn fn)
return output;
}

-static size_t string_length(const char *s)
+size_t string_length(const char *s)
{
    int len;
    wchar_t *wctr;
diff --git a/src/pacman/util.h b/src/pacman/util.h
index 52e79915..d8f7f5f2 100644
--- a/src/pacman/util.h
+++ b/src/pacman/util.h
@@ -47,6 +47,7 @@ typedef struct _pm_target_t {
    int is_explicit;
} pm_target_t;

+size_t string_length(const char *s);
void trans_init_error(void);
/* flags is a bitfield of alpm_transflag_t flags */
int trans_init(int flags, int check_valid);

```

Figure D.1: The diff for the fuzzing harness code

E. Fix Review Results

When undertaking a fix review, Trail of Bits reviews the fixes implemented for issues identified in the original report. This work involves a review of specific areas of the source code and system configuration, not comprehensive analysis of the system.

From March 4 to March 6, 2024, Trail of Bits reviewed the fixes and mitigations implemented by the Arch Linux team for the issues identified in this report. We reviewed each fix to determine its effectiveness in resolving the associated issue.

In summary, of the nine issues described in this report, Arch Linux has resolved seven issues and has partially resolved two issues. For additional information, please see the Detailed Fix Review Results below.

ID	Title	Status
1	Use-after-free vulnerability in the <code>print_packages</code> function	Resolved
2	Null pointer dereferences	Resolved
3	Allocation failures can lead to memory leaks or null pointer dereferences	Resolved
4	Buffer overflow read in <code>string_length</code> utility function	Resolved
5	Undefined behavior or potential null pointer dereferences	Partially Resolved
6	Undefined behavior from use of <code>atoi</code>	Resolved
7	Database parsers fail silently if an option is not recognized	Resolved
8	Cache cleaning function may delete the wrong files	Partially Resolved
9	Integer underflow in a length check leads to out-of-bounds read in <code>alpm_extract_keyid</code>	Resolved

Detailed Fix Review Results

TOB-PACMAN-1: Use-after-free vulnerability in the `print_packages` function

Resolved in [commit 36fcff6e](#). This commit adds an assignment that overwrites the freed temp variable with the newly allocated string variable.

TOB-PACMAN-2: Null pointer dereferences

Resolved in [commit 74deada5](#). This commit adds the necessary checks to determine whether the `pkgname` variable is null before using it.

The Pacman developers correctly identified that the `write_to_child` function can only ever be called with a non-null callback, so a fix for that portion of the issue was not necessary.

TOB-PACMAN-3: Allocation failures can lead to memory leaks or null pointer dereferences

Resolved in commits [6711d10f](#) and [abc6dd74](#). Commit [6711d10f](#) adds a check to the `setdefaults` function that ensures that the pointer returned by the `strdup` function is non-null before using it. Commit [abc6dd74](#) adds a check to the `alpm_list_cmp_unsorted` function that ensures that the pointer returned by `calloc` is non-null before using it.

The Pacman developers identified the code in figure 3.3 as not being an issue. We have confirmed that this is the case: it should not be possible for the `line` variable to be null without the `goto error` statement being executed; this prevents `pkg->filename` from being null in the call to the `_alpm_validate_filename` function.

TOB-PACMAN-4: Buffer overflow read in `string_length` utility function

Resolved in [commit c9c56be3](#). This commit changes the `string_length` function so that it loops under more strict conditions: it stops once it reaches a character that is not a digit or a semicolon, rather than reading until an “m” is found.

TOB-PACMAN-5: Undefined behavior or potential null pointer dereferences

Partially resolved in commits [f996f301](#) and [ce528a26](#). Commit [f996f301](#) adds a check to the `shift_pacsave` function that ensures that the `dir` pointer is non-null before using it in a `closedir(dir)` call. Commit [ce528a26](#) adds a check to the `mount_point_list` function that ensures that the `mnt->mnt_dir` value is non-null before attempting to duplicate it in `mp->mount_dir` using the `STRDUP` macro. This ensures that `mp->mount_dir` will be non-null as well, which prevents undefined behavior during the call to `strlen(mp->mount_dir)`.

The instances of undefined behavior shown in figure 5.3 have not yet been resolved.

TOB-PACMAN-6: Undefined behavior from use of atoi

Resolved in [commit 6e6d3f18](#) and [MR #136](#). Commit 6e6d3f18 replaces the use of `atoi` in the `_alpm_local_db_pkgpath` function with a set of `strcmp` comparisons. MR #136 replaces the uses of `atoi` in the `parsearg_global` function with calls to `strtol` and adds all the necessary error checks.

TOB-PACMAN-7: Database parsers fail silently if an option is not recognized

Resolved in [commit e1dc6099](#). This commit adds a warning message that is logged in the case of an unknown option.

TOB-PACMAN-8: Cache cleaning function may delete the wrong files

Partially resolved in [commit a6b25247](#). This commit adds a check determining whether `len` is less than `PATH_MAX` and skipping the current file if it is. However, the check should instead determine whether `len` is less than *or equal to* `PATH_MAX` since the value returned by the `snprintf` function does not count the trailing null character.

The commit also fixes a very similar issue in the `sync_cleandb` function that was not found during the audit. However, the fix for the `sync_cleandb` function only prints an error message in the case of a problem but does not skip the current file. In addition, the fix has the same issue mentioned above of using the greater than `[>]` operator rather than the greater than or equal to `[>=]` operator.

TOB-PACMAN-9: Integer underflow in a length check leads to out-of-bounds read in `alpm_extract_keyid`

Resolved in [commit 16a2a797](#). This commit adds a new check to determine whether `position` is greater than `length` before computing `length minus position`. If `position` is greater than `length`, an error is returned.

F. Fix Review Status Categories

The following table describes the statuses used to indicate whether an issue has been sufficiently addressed.

Fix Status	
Status	Description
Undetermined	The status of the issue was not determined during this engagement.
Unresolved	The issue persists and has not been resolved.
Partially Resolved	The issue persists but has been partially resolved.
Resolved	The issue has been sufficiently resolved.