



# Scroll Compression Circuits

Security Assessment

July 9, 2024

*Prepared for:*

**Rohit Narurkar**

Scroll

*Prepared by:* **Opal Wright, Will Song, Filipe Casal, Joe Doyle, and Marc Ilunga**

# About Trail of Bits

---

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at [info@trailofbits.com](mailto:info@trailofbits.com).

## **Trail of Bits, Inc.**

497 Carroll St., Space 71, Seventh Floor  
Brooklyn, NY 11215

<https://www.trailofbits.com>

[info@trailofbits.com](mailto:info@trailofbits.com)

# Notices and Remarks

---

## Copyright and Distribution

© 2024 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to Scroll under the terms of the project statement of work and has been made public at Scroll's request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

## Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

# Table of Contents

---

<b>About Trail of Bits</b>	<b>1</b>
<b>Notices and Remarks</b>	<b>1</b>
<b>Table of Contents</b>	<b>3</b>
<b>Project Summary</b>	<b>5</b>
<b>Executive Summary</b>	<b>7</b>
<b>Project Goals</b>	<b>9</b>
<b>Project Targets</b>	<b>10</b>
<b>Project Coverage</b>	<b>11</b>
<b>Automated Testing</b>	<b>12</b>
<b>Codebase Maturity Evaluation</b>	<b>13</b>
<b>Summary of Findings</b>	<b>15</b>
<b>Detailed Findings</b>	<b>17</b>
1. Multiple missing Boolean constraints on Boolean advice columns	17
2. Column annotations do not match lookup table columns	19
3. Unexpected BlockType for LiteralsHeader reaches unreachable! macro	22
4. RomTagTransition table does not allow ZstdBlockSequenceHeader -> BlockHeader transitions	23
5. The back referencing phase is not properly constrained to a monotone behavior once activated	26
6. The blob-based public input commitment scheme is poorly documented	28
7. Left shift leads to undefined behavior	31
8. Missing constraints for Block_Maximum_Size	32
9. Apparent discrepancy between bitwise-op-table configuration and code comment	34
10. The compression mode reserved field is not enforced to equal zero	36
11. The tag_config.is_change witness is partially unconstrained	38
12. The is_1lt/is_mot/is_mlt constraints are only valid if self.table_kind is in {1, 2, 3}	41
13. Values larger than 23 satisfy the "spans_three_bytes" constraints	43
14. Missing a large number of test cases that should fail	44
<b>A. Vulnerability Categories</b>	<b>46</b>
<b>B. Code Maturity Categories</b>	<b>48</b>
<b>C. Code Quality Findings</b>	<b>50</b>
<b>D. Automated Analysis Tool Configuration</b>	<b>56</b>
Expression is reconstrained rule	56



# Project Summary

---

## Contact Information

The following project manager was associated with this project:

**Anne Marie Barry**, Project Manager  
[annemarie.barry@trailofbits.com](mailto:annemarie.barry@trailofbits.com)

The following engineering director was associated with this project:

**Jim Miller**, Engineering Director, Cryptography  
[james.miller@trailofbits.com](mailto:james.miller@trailofbits.com)

The following consultants were associated with this project:

**Will Song**, Consultant  
[will.song@trailofbits.com](mailto:will.song@trailofbits.com)

**Opal Wright**, Consultant  
[opal.wright@trailofbits.com](mailto:opal.wright@trailofbits.com)

**Filipe Casal**, Consultant  
[filipe.casal@trailofbits.com](mailto:filipe.casal@trailofbits.com)

**Marc Ilunga**, Consultant  
[marc.ilunga@trailofbits.com](mailto:marc.ilunga@trailofbits.com)

**Joe Doyle**, Consultant  
[joseph.doyle@trailofbits.com](mailto:joseph.doyle@trailofbits.com)

## Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
May 20, 2024	Pre-project kickoff call
May 24, 2024	Status update meeting #1
May 31, 2024	Status update meeting #2
June 7, 2024	Status update meeting #3
June 14, 2024	Delivery of report draft
June 14, 2024	Report readout meeting
June 21, 2024	Delivery of updated report draft
June 21, 2024	Updated report meeting

**July 9, 2024**

Delivery of comprehensive report

# Executive Summary

---

## Engagement Overview

Scroll engaged Trail of Bits to review the security of their halo2 zstd decoder. The associated circuits are to be integrated into Scroll's zkEVM and reduce the size of Blob-carrying transactions through compression using the zstd compression algorithm.

A team of five consultants conducted the review from May 20 to June 21, 2024, for a total of 12 engineer-weeks of effort. Our testing efforts focused on completeness and soundness of circuits that implement decompression according to the zstd standard. With full access to source code and documentation, we performed static and dynamic testing of the zstd decoder, using automated and manual processes.

## Observations and Impact

Overall, the code shows signs of having small gaps in oversight, where typos are inevitably introduced in variable names or certain features are re-enabled. Constraints sometimes appear to be non-local to their actual use/initialization, which causes some values to become underconstrained in certain scenarios. Although practicing better code writing practices may alleviate some of these issues, the code's complexity ultimately requires multiple sets of eyes—whether internal or external—to review the code for small errors.

We found the provided documentation to be particularly helpful in providing examples when it came to our internal review process, which greatly alleviated many pain points often seen in ZK audits.

## Recommendations

Based on the codebase maturity evaluation and findings identified during the security review, Trail of Bits recommends that Scroll take the following steps:

- **Remediate the findings disclosed in this report.** These findings should be addressed as part of a direct remediation or as part of any refactor that may occur when addressing other recommendations.
- **Strengthen inline documentation.** With a ZK circuit this large, it is often useful to document the exact constraints expected of variables and perhaps where to find these constraints. There are multiple instances of columns—for example, `TagConfig.is_reverse` and `FseDecoder.table_kind`—whose values are indirectly constrained via `lookup_any`, and these constraints can be particularly laborious to track down.
- **Use wrapper types on highly constrained values.** The codebase often uses advice columns that are highly specialized in nature, and need to lie in a small set of



possible values. In particular, the circuit uses many Boolean advice columns as well as a `table_kind` column, all of which must fall in the expected range for their derived expressions to be correct. Such columns deserve their own wrapper type to ensure proper constraints over their assigned witness values.

## Finding Severities and Categories

The following tables provide the number of findings by severity and category.

### EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	2
Medium	0
Low	4
Informational	6
Undetermined	2

### CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Cryptography	7
Data Validation	5
Denial of Service	1
Undefined Behavior	1

# Project Goals

---

The engagement was scoped to provide a security assessment of the Scroll halo2 zstd decoder. Specifically, we sought to answer the following non-exhaustive list of questions:

- Does the zstd decoder correctly decode the encoded data output by their custom encoder?
- Is the zstd decoder underconstrained? If so, a proof consisting of invalid assignments for a specific decoder instance will pass validation.
- Do the zstd decoder tests accurately target the specific edge cases Scroll is concerned about?
- Does the witness generator correctly implement the zstd decoding algorithm on the given input blob and assign the correct intermediate values to the DecoderConfig?

# Project Targets

---

The engagement involved a review and testing of the following targets.

## zkEVM Compression Circuits

Repository	<a href="https://github.com/scroll-tech/zkevm-circuits/">https://github.com/scroll-tech/zkevm-circuits/</a>
Version	475bc1f7c0531d8c5fd90de19b4e997f5367843b 4658b6e995400d789698ed0bc2a021d6b40788f5
Type	Rust library

# Project Coverage

---

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- **FixedTable.** We conducted a manual review of the fixed lookup table values to assess correctness.
- **FseTable/FseSortedStatesStable.** We conducted a manual review of the two finite state entropy tables to assess the correct computation of the finite state entropy logic of the zstd decoder.
- **BitstringTable.** We conducted a manual review of the bitstring table to test the correct translation of certain bitstrings of the input blob to their corresponding integer values. We also checked lookups into the bitstring table for correctness.
- **LiteralHeadersTable.** We performed a manual review to assess the completeness and soundness of the circuit parsing the literals header, and whether it follows the specification.
- **SeqInstTable/SeqExecConfig.** We performed a manual review of the sequence execution table and the sequence instruction table to assess the correctness of instructions decoding, paying special attention to the offset update rules.
- **DecoderConfig.** We conducted a manual review to ensure the correct linkage of components, as well as the constraints for controlling the specific instances of computation managed by the components. We took special care to check the various conditions activating constraints, the variables used, and their respective constraints.

## Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. The following list outlines the coverage limitations of the engagement and indicates system elements that may warrant further review:

- We did not validate that every edge case described by the zstd documentation will not cause a proof validation failure or crash the witness generator.

# Automated Testing

---

Trail of Bits uses automated techniques to extensively test the security properties of software. We use both open-source static analysis and fuzzing utilities, along with tools developed in house, to perform automated testing of source code and compiled software.

## Test Harness Configuration

We used the following tools in the automated testing phase of this project:

Tool	Description	Policy
<a href="#">Semgrep</a>	An open-source static analysis tool for finding bugs and enforcing code standards when editing or committing code and during build time	<a href="#">Appendix D</a>
<a href="#">Clippy</a>	An open-source Rust linter used to catch common mistakes and unidiomatic Rust code	<a href="#">Appendix D</a>

## Areas of Focus

Our automated testing and verification work focused on the following system properties:

- Identification of general code quality issues and unidiomatic code patterns
- Identification of dangerous halo2-specific and Scroll's API patterns

## Test Results

The results of this focused testing are detailed below.

### Clippy

Running Clippy in pedantic mode identifies several unidiomatic patterns and potential issues related to casting that should be investigated. In [appendix D](#) we describe how to run Clippy in pedantic mode, and how to efficiently triage these results using the SARIF file format.

### Semgrep

We present some of the rules that we wrote to find halo2-specific and Scroll's API patterns in [appendix D](#).

# Codebase Maturity Evaluation

Trail of Bits uses a traffic-light protocol to provide each client with a clear understanding of the areas in which its codebase is mature, immature, or underdeveloped. Deficiencies identified here often stem from root causes within the software development life cycle that should be addressed through standardization measures (e.g., the use of common libraries, functions, or frameworks) or training and awareness programs.

Category	Summary	Result
Arithmetic	We identified one low-severity issue that can lead to undefined behavior in witness generation. No other arithmetic issues were identified.	Satisfactory
Complexity Management	There are multiple instances of excess business logic being placed at the top level of some functions. The DecoderConfig table is too large to effectively analyze, troubleshoot, or maintain. While the comments are helpful, the haphazard organization of the constraints makes it difficult to analyze individual constraints for correctness. As noted in the Code Quality Findings section, functions can run to over 1,000 lines; these functions need to be broken into smaller components.	Weak
Data Handling	The witness generator appears to correctly assign all intermediate values to the column variables, resulting in a correct decompressing of the zstd compressed blob.	Strong
Documentation	The provided Notion documentation was incredibly useful for getting up to speed with the code. However, we found inline documentation to be slightly lacking, especially when it came to assumptions regarding variables and their expected constraints. Furthermore, we found a few typos in a code snippet included with the documentation. Sometimes, the typos resulted in functionalities diverging from the Halo2 implementation, potentially introducing soundness issues if circuits have been implemented according to the documentation.	Moderate
Maintenance	Some submodules can only be described as monolithic, which is not recommended by modern coding standards. Other parts of the code are hard-coded according to the	Weak

	current requirements and require additional work should the requirements be expanded in the future. The most prominent example is the <code>BitstringTable</code> , which is hard-coded for up to three bytes and certain conditional expressions relying on <code>bit_index_end</code> column.	
Memory Safety and Error Handling	Overall, errors are mostly handled by the underlying <code>halo2</code> library and custom error cases are not needed. In the case of calling <code>expect</code> , which might panic under generic circumstances, valid reasons are given.	Satisfactory
Testing and Verification	While several unit tests were present in the codebase, they tend to focus more on high-level <code>zstd</code> edge cases rather than low-level witness edge cases. While the former ensures good coverage and overall <code>zstd</code> decoding correctness, also testing the expected constraints on a system by including tests that should fail is critical to providing a higher guarantee of soundness in the long run.	Weak

## Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Multiple missing Boolean constraints on Boolean advice columns	Data Validation	High
2	Column annotations do not match lookup table columns	Data Validation	Informational
3	Unexpected BlockType for LiteralsHeader reaches unreachable! macro	Denial of Service	Informational
4	RomTagTransition table does not allow ZstdBlockSequenceHeader -> BlockHeader transitions	Cryptography	Undetermined
5	The back referencing phase is not properly constrained to a monotone behavior once activated.	Cryptography	Undetermined
6	The blob-based public input commitment scheme is poorly documented	Cryptography	Informational
7	Left shift leads to undefined behavior	Undefined Behavior	Low
8	Missing constraints for Block_Maximum_Size	Data Validation	Low
9	Apparent discrepancy between bitwise-op-table configuration and code comment	Cryptography	Informational
10	The compression mode reserved field is not enforced to equal zero	Data Validation	Low
11	The tag_config.is_change witness is partially unconstrained	Cryptography	Informational



12	The is_llt/is_mot/is_mlt constraints are only valid if self.table_kind is in {1, 2, 3}	Data Validation	High
13	Values larger than 23 satisfy the "spans_three_bytes" constraints	Cryptography	Informational
14	Missing a large number of test cases that should fail	Cryptography	Low

# Detailed Findings

## 1. Multiple missing Boolean constraints on Boolean advice columns

Severity: High

Difficulty: Medium

Type: Data Validation

Finding ID: TOB-SCROLLZSTD-1

Target: zkevm-circuits/aggregator/src/

### Description

The zstd decoder tables require the use of many Boolean type values, being either 0 or 1. Other values, such as 2 or -1, would cause many of the formulas used for Boolean logic to catastrophically fail, leading to potential constraint compromise.

Below, we detail a non-exhaustive list of Boolean advice columns that appear to be unconstrained.

- `FseTable::is_new_symbol`
- `FseTable::is_skipped_state`
- `FseSortedStatesTable::is_new_symbol`
- `FseSortedStatesTable::is_skipped_state`
- `FseDecoder::is_repeat_bits_loop`
- `FseDecoder::is_trailing_bits`
- `BitstreamDecoder::is_nb0`
- `BitstreamDecoder::is_nil`
- `BlockConfig::compression_modes`
- `TagConfig::is_reverse`

### Exploit Scenario

An attacker leverages the soundness issues of an unconstrained Boolean to disable other constraints that are derived from the unconstrained Boolean. This would allow an attacker to generate a valid proof with an invalid witness, compromising the correctness of the zstd decoder circuit.

### Recommendations

Short term, perform the correct Boolean constraints on the necessary columns.

Long term, consider annotating the Boolean columns with a wrapper type that checks if the constraint function has been called. Consider a set of checks that will ensure that all `BooleanAdvice` columns have been constrained. Perhaps an automated way to do this is

a custom Drop implementation that will panic, alerting developers during testing. An alternative is to restrict access to the underlying column until the constraint has been added.

```
struct BooleanAdvice {
    col: Column<Advice>,
    constrained: bool,
}

impl BooleanAdvice {
    fn column(&self) -> &Column<Advice> {
        if self.constrained {
            &self.col
        } else {
            panic!("unconstrained boolean advice")
        }
    }
}
```

*Figure 1.1: An example Boolean advice wrapper and protected access*

## 2. Column annotations do not match lookup table columns

Severity: Informational

Difficulty: N/A

Type: Data Validation

Finding ID: TOB-SCROLLZSTD-2

Target: aggregator/src/aggregation/decoder/tables/{seqinst\_table.rs, literals\_header.rs}, zkevm-circuits/src/rlp\_circuit\_fsm.rs

### Description

The column annotations for the SeqInstTable, RlpFsmDataTable, and LiteralsHeaderTable tables do not match the tables' columns:

- In the SeqInstTable the n\_seq and block\_index annotations are out of order;
- In the LiteralsHeaderTable, there is an additional annotation (byte\_offset) that would cause all subsequent annotations to refer to the wrong column;
- The RlpFsmDataTable has an unannotated column (gas\_cost\_acc).

```
fn columns(&self) -> Vec<Column<Any>> {
    vec![
        self.q_enabled.into(),
        self.block_index.into(),
        self.n_seq.into(),
        self.s_beginning.column.into(),
        self.seq_index.into(),
        self.literal_len.into(),
        self.match_offset.into(),
        self.match_len.into(),
    ]
}

fn annotations(&self) -> Vec<String> {
    vec![
        String::from("q_enabled"),
        String::from("n_seq"),
        String::from("block_index"),
    ]
}
```

Figure 2.1:

*aggregator/src/aggregation/decoder/tables/seqinst\_table.rs#L130-L147*

```
fn columns(&self) -> Vec<Column<Any>> {
    vec![
        self.block_idx.into(),
        self.byte0.into(),
    ]
}
```

```

        self.byte1.into(),
        self.byte2.into(),
        self.size_format_bit0.into(),
        self.size_format_bit1.into(),
        self.regen_size.into(),
        self.is_padding.column.into(),
    ]
}

fn annotations(&self) -> Vec<String> {
    vec![
        String::from("block_idx"),
        String::from("byte_offset"),
        String::from("byte0"),
        String::from("byte1"),
        String::from("byte2"),
        String::from("size_format_bit0"),
        String::from("size_format_bit1"),
        String::from("regen_size"),
        String::from("is_padding"),
    ]
}

```

Figure 2.2:

*aggregator/src/aggregation/decoder/tables/literals\_header.rs#L274–L301*

```

impl<F: Field> LookupTable<F> for RlpFsmDataTable {
    fn columns(&self) -> Vec<Column<Any>> {
        vec![
            self.tx_id.into(),
            self.format.into(),
            self.byte_idx.into(),
            self.byte_rev_idx.into(),
            self.byte_value.into(),
            self.bytes_rlc.into(),
            self.gas_cost_acc.into(),
        ]
    }

    fn annotations(&self) -> Vec<String> {
        vec![
            String::from("tx_id"),
            String::from("format"),
            String::from("byte_idx"),
            String::from("byte_rev_idx"),
            String::from("byte_value"),
            String::from("bytes_rlc"),
        ]
    }
}

```

Figure 2.3: *zkevm-circuits/src/rlp\_circuit\_fsm.rs#L61–L84*

## Recommendations

Short term, fix the reported column annotations.

Long term, add debug assertions or tests that ensure that the number of columns and annotations is the same for a lookup table. Consider using `zip_eq` instead of `zip` in the `LookupTable::{annotate_columns, annotate_columns_in_region}` functions:

```
/// Annotates a lookup table by passing annotations for each of it's
/// columns.
fn annotate_columns(&self, cs: &mut ConstraintSystem<F>) {
    self.columns()
        .iter()
        .zip(self.annotations().iter())
        .for_each(|(&col, ann)| cs.annotate_lookup_any_column(col, || ann))
}

/// Annotates columns of a table embedded within a circuit region.
fn annotate_columns_in_region(&self, region: &mut Region<F>) {
    self.columns()
        .iter()
        .zip(self.annotations().iter())
        .for_each(|(&col, ann)| region.name_column(|| ann, col))
}
```

Figure 2.4: *zkevm-circuits/src/table.rs#87-102*

Alternatively, consider adding a derive macro for `LookupTable`. This way, columns made available for lookup can easily be annotated via `derive macro helper attributes`. The figure below shows an example:

```
#[derive(Clone, Debug, LookupTable)]
pub struct SeqInstTable<F: Field> {
    #[lookup] q_enabled: Column<Fixed>,
    #[lookup] block_index: Column<Advice>,
    #[lookup] n_seq: Column<Advice>,
    #[lookup] seq_index: Column<Advice>,
    #[lookup] s_beginning: Column<Advice>,
    // ...

    offset: Column<Advice>,
    acc_literal_len: Column<Advice>,
    // ...
}
```

Figure 2.5: *A derive macro-based lookup and annotation system*

### 3. Unexpected BlockType for LiteralsHeader reaches unreachable! macro

Severity: Informational

Difficulty: N/A

Type: Denial of Service

Finding ID: TOB-SCROLLZSTD-3

Target: aggregator/src/aggregation/decoder/tables/literals\_header.rs

#### Description

The witness assignment code uses an `unreachable!` macro if an unexpected `BlockType` is found while parsing the header bytes. If handpicked values are chosen in `byte0`, this will lead to a runtime panic during witness generation/assignment.

```
let lh_bytes = [byte0 as u8, byte1 as u8, byte2 as u8];
let literals_block_type = BlockType::from(lh_bytes[0] & 0x3);
let size_format = (lh_bytes[0] >> 2) & 3;

let [n_bits_fmt, n_bits_regen, n_bytes_header]: [usize; 3] =
    match literals_block_type {
        BlockType::RawBlock => match size_format {
            0b00 | 0b10 => [1, 5, 1],
            0b01 => [2, 12, 2],
            0b11 => [2, 20, 3],
            _ => unreachable!("size_format out of bound"),
        },
        _ => unreachable!(
            "BlockType:* unexpected. Must be raw bytes for literals."
        ),
    };
```

Figure 3.1:

*aggregator/src/aggregation/decoder/tables/literals\_header.rs#206-221*

#### Exploit Scenario

The prover is passed a malformed compressed blob with an unexpected block type, causing the prover to halt.

#### Recommendations

Short term, return an `Error` instead of calling the `unreachable!` macro.

Long term, investigate all calls to the `unreachable!` macro used during witness assignment.

#### 4. RomTagTransition table does not allow ZstdBlockSequenceHeader -> BlockHeader transitions

Severity: Undetermined	Difficulty: Undetermined
Type: Cryptography	Finding ID: TOB-SCROLLZSTD-4
Target: aggregator/src/aggregation/decoder/tables/fixed/tag_transition.rs, aggregator/src/aggregation/decoder.rs	

#### Description

The RomTagTransition table does not allow transitioning from ZstdBlockSequenceHeader to BlockHeader. This causes a completeness issue where honestly generated compressed data does not satisfy the tag transition circuit. Figure 4.1 shows the allowed transitions and highlights the transitions that reach BlockHeader – none originate from ZstdBlockSequenceHeader:

```
[
  (FrameHeaderDescriptor, FrameContentSize),
  (FrameContentSize, BlockHeader),
  (BlockHeader, ZstdBlockLiteralsHeader),
  (ZstdBlockLiteralsHeader, ZstdBlockLiteralsRawBytes),
  (ZstdBlockLiteralsRawBytes, ZstdBlockSequenceHeader),
  (ZstdBlockSequenceHeader, ZstdBlockSequenceFseCode),
  (ZstdBlockSequenceHeader, ZstdBlockSequenceData),
  (ZstdBlockSequenceFseCode, ZstdBlockSequenceFseCode),
  (ZstdBlockSequenceFseCode, ZstdBlockSequenceData),
  (ZstdBlockSequenceData, BlockHeader), // multi-block
  (ZstdBlockSequenceData, Null),
  (Null, Null),
]
```

Figure 4.1:

*aggregator/src/aggregation/decoder/tables/fixed/tag\_transition.rs#L28-L41*

However, in the specification, the ZstdBlockSequenceHeader to BlockHeader transition is described and has associated constraints:



We also do a few more checks to see if the previous block (if any) did end correctly and that the `block_idx` is incremented appropriately:

```
# if we are at the first byte of the 3-bytes Block_Header
if tag::cur == BlockHeader and is_change::cur == 1:
    # block_idx increments
    block_idx::cur == block_idx::prev + 1
    # block_len, is_last_block and block_idx do not change over the Block_Header
    assert block_len(0) == block_len(1) == block_len(2)
    assert block_idx(0) == block_idx(1) == block_idx(2)
    assert is_last_block(0) == is_last_block(1) == is_last_block(2)
    # block_len, is_last_block and block_idx are ensured to retain the same
    # value over subsequent tags with is_block=true

    # previous tag should be either FrameContentSize, SequenceData or SequenceHeader
    assert tag::prev in [FrameContentSize, ZstdBlockSequenceData, ZstdBlockSequenceHeader]

    # it means this is the first block
    if tag::prev == FrameContentSize:
        assert block_idx == 1
    # this can only happen if the previous block had no sequences to decode
    if tag::prev == ZstdBlockSequenceHeader:
        assert block_config.num_sequences::prev == 0
    # ensure that all the sequences in the previous block were processed
    if tag::prev == ZstdBlockSequenceData:
        assert block_config.num_sequences::prev == sequence_decoder.seq_idx::prev
```

Figure 4.2: *Circuit specification* where a `ZstdBlockSequenceHeader` to `BlockHeader` transition is mentioned

Figure 4.3 shows the implementation of the constraint from figure 4.2:

```
cb.condition(is_prev_sequence_header(meta), |cb| {
  cb.require_equal(
    "tag::prev=SeqHeader",
    config
      .block_config
      .is_empty_sequences(meta, Rotation::prev()),
    1.expr(),
  );
});
```

Figure 4.3: *aggregator/src/aggregation/decoder.rs#1807–1815*

## Exploit Scenario

A two-block encoded data, where the first block contains no sequences to decode, is honestly generated by the zstd compression algorithm. However, due to the missing transition, this compressed data does not satisfy the decoder circuit.

## Recommendations

Short term, allow for the ZstdBlockSequenceHeader to BlockHeader transition to occur in the tag transition circuit, or document why this is not allowed.

Long term, add tests with zstd compressed data that exercise all possible valid tag transitions.

## 5. The back referencing phase is not properly constrained to a monotone behavior once activated

Severity: Undetermined

Difficulty: Low

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-5

Target: aggregator/src/aggregation/decoder/seq\_exec.rs

### Description

The column `s_back_ref_phase` is not properly constrained to ensure the back-referencing phase is monotone once activated. The issue is likely due to a typo.

Executing decoding sequences happens in phases: Literal copy and back-references. The two phases must not occur simultaneously. Furthermore, each phase must have a monotone behavior—i.e., the literal copy phase (the back-referencing phase, respectively) remains deactivated (activated, respectively). The constraint in figure 6.1 is meant to ensure the monotonicity of the back referencing phase once activated. However, likely due to a typo, `s_back_ref_phase_prev` is constrained to be equal to 1 instead of `s_back_ref_phase`.

```
cb.condition(
  and::expr([
    not::expr(is_inst_begin.expr()),
    s_back_ref_phase_prev.expr(),
  ]),
  |cb| {
    cb.require_equal(
      "inside a inst, backref phase keep 1 once it changed to 1",
      s_back_ref_phase_prev.expr(),
      1.expr(),
    );
  },
);
```

Figure 5.1: aggregator/src/aggregation/decoder/seq\_exec.rs#L393–L405

However, other constraints enforce that either a phase is activated or the current rows correspond to padding. But exploiting this issue appears difficult due to adjacent constraints. The monotonicity of the literal copy phase once deactivated is guaranteed by appropriate constraints. Therefore, to use arbitrary values in the column `s_back_ref_phase`, a malicious prover must produce a copy command that is compatible with the copied value. On the other hand, when no phase is activated, the current rows correspond to padding rows. An attacker could potentially abuse the ineffective constraints

to provide a shorter witness. It is unclear whether this approach is feasible and what impact such an attack may ultimately have on the overall system.

## Exploit Scenario

An attacker notices the defective constraints and produces a false witness by exploiting the missing constraints for monotonicity on `s_back_ref_phase`. We could not fully determine the feasibility of producing malicious witnesses or whether monotonicity is fully guaranteed by adjacent constraints.

## Recommendations

Short term, amend the code so that it constrains `s_back_ref_phase` to be equal to 1 instead of `s_back_ref_phase_prev`.

```
cb.condition(
  and::expr([
    not::expr(is_inst_begin.expr()),
    s_back_ref_phase_prev.expr(),
  ]),
  |cb| {
    cb.require_equal(
      "inside a inst, backref phase keep 1 once it changed to 1",
      s_back_ref_phase.expr(),
      1.expr(),
    );
  },
);
```

*Figure 5.2: An example fix of the aforementioned issue*

Long term, review the codebase for potential variable typos and patterns that could render constraints void. Some examples of patterns to investigate include: comparing a variable with itself, subtracting a variable from itself, or naming conventions such as: `(variable_name, variable_name_prev)`.

## 6. The blob-based public input commitment scheme is poorly documented

Severity: Informational

Difficulty: N/A

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-6

Target: aggregator /

### Description

To reduce the gas cost of verification, the Scroll team has moved some of the ZkEVM's public input data into the EIP-4844 "blob" structure, whereby the underlying data is stored only temporarily, at a lower gas cost. The verifier contract has access only to a polynomial commitment to the blob data and evaluation proofs. Additionally, with the addition of the zstd circuits reviewed in this report, the data in the blob is instead in zstd-compressed form.

When the ZkEVM verifier is deployed in the Scroll rollup contract, a batch, consisting of both L1 and L2 transactions and associated metadata, is "committed" in the rollup contract. Committed batches can then either be reverted or finalized. To finalize a batch, a ZkEVM proof is checked against the metadata provided in the commit stage, and if the proof succeeds, the batch is finalized and can no longer be reverted.

Since the ZkEVM prover should be untrusted in the system, the metadata in the commit stage must uniquely identify the transactions in the underlying batch; otherwise, the prover may be able to finalize a different sequence of transactions than was intended.

Prior to EIP-4844 integration, the sequence of transactions in a chunk was fully determined by the "public input hash" of the ZkEVM circuit, which would then be combined together into an overall public input hash by the aggregation circuit. In this scheme, the unique identification is a straightforward consequence of the collision resistance of the hash used.

However, the current scheme is more complex. The transactions are now split between this public input hash and the blob structure. The PI subcircuit of the ZkEVM splits the underlying sequence of transactions into "L1" and "L2" transactions. The L1 transactions are included in the public input hash, and the L2 transactions are included in the "chunk tx hash." Unless there is some as-yet-unknown flaw in the PI subcircuit, this guarantees the uniqueness of the L1 transactions as before.

To commit to the L2 transactions, the overall public inputs of the aggregation circuit include a tuple  $(\text{versionedHash}, z, y)$ , representing the polynomial commitment to the blob and the evaluation that must be checked by the verifier. To conclude that this uniquely identifies a particular sequence of L2 transactions, we must determine that:

1.  $p(z) = y$ , where  $p(X)$  is a unique polynomial corresponding to `versionedHash`;
2.  $q(z) = y$ , where  $q(X)$  is a unique polynomial corresponding to the L2 transaction sequence;
3.  $z$  is a pseudorandom challenge derived from a transcript including commitments to both the L2 transaction sequence and `versionedHash`.

These three requirements suffice because blobs represent degree-4096 polynomials over a 254-bit finite field. The chance that  $p(z) = q(z)$  and  $p(X) \neq q(X)$  for a randomly sampled point  $z$  is negligible, and requirement (3) allows us to treat  $z$  as randomly sampled for the purpose of checking that  $p(X) = q(X)$  by the Fiat-Shamir heuristic.

Requirement (1) is ensured because the verifier contract checks a point evaluation proof, and `versionedHash` is a hash of a KZG commitment, which uniquely determines  $p(X)$ .

The evaluation in requirement (2) is checked by the `BarycentricEvaluationConfig` circuit. However, the uniqueness of  $q(X)$  is more complicated, since the underlying transaction data is not stored in the blob. Instead, the `BatchDataConfig` subcircuit of the aggregation circuit includes the whole L2 transaction sequence data as private witness values. The aggregation circuit checks those hashes against the “chunk tx hash” public inputs of the ZkEVM proofs being aggregated, and checks that the L2 transaction sequence is the result of zstd-decompressing the data used in the `BarycentricEvaluationConfig` circuit.

To then establish that  $q(X)$  is unique, we must assume (a) that the zstd decompression circuit is deterministic; and (b) that the serialization of the L2 transactions when computing the chunk tx hash is unique.

Finally, requirement (3) is ensured by checking that the challenge can be computed as a hash of data that includes the `versionedHash` and the chunk tx hashes, via an internal lookup in the `BatchDataConfig` table, shown below in figure 6.1.

```
// lookup challenge digest in keccak table.
meta.lookup_any(
  "BatchDataConfig (metadata/chunk_data/challenge digests in keccak table)",
  |meta| {
    let is_hash = meta.query_selector(config.hash_selector);
    let is_boundary = meta.query_advice(config.is_boundary, Rotation::cur());

    // when is_boundary is set in the "digest RLC" section.
    // this is also the last row of the "digest RLC" section.
    let cond = is_hash * is_boundary;

    // - metadata_digest: 32 bytes
    // - chunk[i].chunk_data_digest: 32 bytes each
    // - versioned_hash: 32 bytes
    let preimage_len = 32.expr() * (N_SNARKS + 1 + 1).expr();
```

```

    [
        1.expr(), // q_enable
        1.expr(), // is final
        meta.query_advice(config.preimage_rlc, Rotation::cur()), // input rlc
        preimage_len, // input len
        meta.query_advice(config.digest_rlc, Rotation::cur()), // output rlc
    ]
    .into_iter()
    .zip_eq(keccak_table.table_exprs(meta))
    .map(|(value, table)| (cond.expr() * value, table))
    .collect()
},
);

```

Figure 6.1: Lookups from the `chunk_data` section of the table to the challenge section of the table ([zkevm-circuits/aggregator/src/aggregation/batch\\_data.rs#334–362](#))

All of these properties appear to hold; however, they are neither explicitly stated nor explicitly justified in Scroll’s documentation. If any of them fails, it would allow a malicious prover to finalize a different batch of transactions than was committed, causing many potential issues such as denial of service or state divergence.

## Recommendations

Short term, document this commitment scheme and specify what properties of different components it relies upon (e.g., deterministic decompression).

Long term, explicitly document all intended security properties of the Scroll rollup, and what is required of each system component to ensure those properties.

## 7. Left shift leads to undefined behavior

Severity: Low

Difficulty: Low

Type: Undefined Behavior

Finding ID: TOB-SCROLLZSTD-7

Target: aggregator/src/aggregation/decoder/witgen/util.rs

### Description

The `read_variable_bit_packing` function accepts both an offset parameter and a maximum value, `r`, indicating the maximum value to be returned.

The `r` parameter is specified as a `u64`. The number of bits required to store the decoded value is computed using a left shift with a variable shift size (see figure 7.1) from the function `bit_length`.

```
// number of bits required to fit a value in the range 0..=r.  
let size = bit_length(r) as u32;  
let max = 1 << size;
```

*Figure 7.1: Calculating bit storage requirements from the range value `r`*  
([zkevm-circuits/aggregator/src/aggregation/decoder/witgen/util.rs#23-24](#))

If the top bit of `r` is set, the `bit_length` function will return 64. This shift is equal to the bit length of the `max` variable, leading to **undefined behavior**. When compiled in debug mode, this should lead to a panic, but in release mode, the behavior is unspecified and can vary from platform to platform. On x86-64 processors, shift lengths are bit masked against `0x3f`, meaning that 64 will reduce to zero, so `max` will be equal to 1. Other platforms may shift the one “off the end,” causing `max` to be 0.

Additionally, on a platform where `1 == 1 << 64`, using the value `u64::MAX` for `r` will result in the check for non-variable bit packing to fail, leading to potential decoding errors.

### Recommendations

Short term, add checks to specifically handle the case where `r` is 64 bits long, whether through expanded handling or by explicitly rejecting values of `r` that can trigger this behavior.

Long term, develop tests that integrate edge cases and validate correct handling. If `r` is restricted to values that will not trigger this edge case, ensure that this is clearly and conspicuously documented.



## 8. Missing constraints for Block\_Maximum\_Size

Severity: Low

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SCROLLZSTD-8

Target: aggregator/src/aggregation/decoder.rs

### Description

The zstd specification document states that the `Block_Size` value should be bounded by `Block_Maximum_Size`, which is the smallest of `Window_Size`, or 128 KB. In Scroll's zstd encoded blobs, the `Single_Segment_flag` is set, so `Window_Size` should equal `Frame_Content_Size` according to the specification.

#### Block\_Size

The upper 21 bits of `Block_Header` represent the `Block_Size`.

When `Block_Type` is `Compressed_Block` or `Raw_Block`, `Block_Size` is the size of `Block_Content` (hence excluding `Block_Header`).

When `Block_Type` is `RLE_Block`, since `Block_Content`'s size is always 1, `Block_Size` represents the number of times this byte must be repeated.

`Block_Size` is limited by `Block_Maximum_Size` (see below).

#### Block\_Content and Block\_Maximum\_Size

The size of `Block_Content` is limited by `Block_Maximum_Size`, which is the smallest of:

- `Window_Size`
- 128 KB

Figure 8.1: Zstd specification defining how the `Block_Size` should be validated

However, the constraints named `DecoderConfig: tag BlockHeader (Block_Size)` do not check that the `Block_Size` value is limited by the smallest value between `Frame_Content_Size` and 128KB.

```
// block_size == block_header >> 3
//
// i.e. block_header - (block_size * (2^3)) < 8
let block_header_lc = meta.query_advice(config.byte, Rotation(2)) * 65536.expr()
  + meta.query_advice(config.byte, Rotation(1)) * 256.expr()
  + meta.query_advice(config.byte, Rotation(0));
let block_size = meta.query_advice(config.block_config.block_len,
Rotation::cur());
let diff = block_header_lc - (block_size * 8.expr());

vec![(condition * diff, config.range8.into())]
});
```

*Figure 8.1: [aggregator/src/aggregation/decoder.rs#L1833-L1843](#)*

## Exploit Scenario

A malicious prover generates a proof for a zstd blob that does not follow the specification. Due to the missing constraint, the verifier still accepts the proof as valid.

## Recommendations

Short term, add the necessary constraint that ensures the correct validation of the `BlockSize` value.

Long term, add positive and negative tests that parse and validate the `BlockHeader` according to the specification.

## 9. Apparent discrepancy between bitwise-op-table configuration and code comment

Severity: Informational

Difficulty: N/A

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-9

Target: aggregator/src/aggregation/decoder.rs,  
zkevm-circuits/src/table.rs

### Description

The implementation and configuration of the `bitwise_op_table` are misleading, as the table's generic parameter does not match the code comment.

The `BitwiseOpTable` structure uses generic arguments to configure which operation it implements.

```
/// Bitwise operation table (AND only)
bitwise_op_table: BitwiseOpTable<1, L, R>,
```

Figure 9.1: *aggregator/src/aggregation/decoder.rs#L88–L89*

However, the code comment states that the table should be for the bitwise-and operation, but the generic parameter of 1 corresponds to the bitwise-or operation in the `BitwiseOp` structure:

```
#[derive(Clone, Copy, Debug)]
/// Bitwise operation types.
pub enum BitwiseOp {
    /// AND
    AND = 0,
    /// OR
    OR,
    /// XOR
    XOR,
}

impl_expr!(BitwiseOp);

/// Lookup table for bitwise AND/OR/XOR operations.
#[derive(Clone, Copy, Debug)]
pub struct BitwiseOpTable<const OP_CHOICE: usize, const RANGE_L: usize, const RANGE_R: usize> {
    /// Denotes op: AND == 0, OR == 1, XOR == 2.
    pub op: Column<Fixed>,
    /// Denotes the left operand.
```

```
pub lhs: Column<Fixed>,
/// Denotes the right operand.
pub rhs: Column<Fixed>,
/// Denotes the bitwise operation on lhs and rhs.
pub output: Column<Fixed>,
```

Figure 9.2: `zkevm-circuits/src/table.rs#L3270-L3294`

Despite this, the actual implementation is correct as an `OP_CHOICE` of 1 corresponds to the `BitwiseOp::AND` variant:

```
/// Assign values to the BitwiseOp table.
pub fn load<F: Field>(&self, layouter: &mut impl Layouter<F>) -> Result<(), Error> {
    layouter.assign_region(
        || "BitwiseOp table",
        |mut region| {
            let mut offset = 0;
            let chosen_ops = match OP_CHOICE {
                1 => vec![BitwiseOp::AND],
                2 => vec![BitwiseOp::OR],
                3 => vec![BitwiseOp::XOR],
```

Figure 9.3: `zkevm-circuits/src/table.rs#L3310-L3319`

## Recommendations

Short term, unify the representations and use enum variants in the generic arguments to clarify the implementation.

## 10. The compression mode reserved field is not enforced to equal zero

Severity: Low

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SCROLLZSTD-10

Target: aggregator/src/aggregation/decoder.rs

### Description

The symbol compression mode specification states that the first two bits, corresponding to the Reserved field, *"must be all-zeroes."* However, the implementation does not constrain these witness values to be zero.

Symbol compression modes				
This is a single byte, defining the compression mode of each symbol type.				
Bit number	7-6	5-4	3-2	1-0
Field name	Literals_Lengths_Mode	Offsets_Mode	Match_Lengths_Mode	Reserved
The last field, Reserved, must be all-zeroes.				

Figure 10.1: *Zstd specification* stating that the Reserved field should be zero

Figure 10.2 shows how the bits at positions 7 to 2 are constrained, but no constraints exist for the bits at positions 1 and 0.

```
let comp_mode_bit0_ll = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[6], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[6], Rotation(2)),
    meta.query_advice(bits[6], Rotation(3)),
  ),
);
let comp_mode_bit1_ll = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[7], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[7], Rotation(2)),
    meta.query_advice(bits[7], Rotation(3)),
  ),
);
```

```

let comp_mode_bit0_om = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[4], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[4], Rotation(2)),
    meta.query_advice(bits[4], Rotation(3)),
  ),
);
let comp_mode_bit1_om = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[5], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[5], Rotation(2)),
    meta.query_advice(bits[5], Rotation(3)),
  ),
);

let comp_mode_bit0_ml = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[2], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[2], Rotation(2)),
    meta.query_advice(bits[2], Rotation(3)),
  ),
);
let comp_mode_bit1_ml = select::expr(
  byte0_lt_0x80.expr(),
  meta.query_advice(bits[3], Rotation(1)),
  select::expr(
    byte0_lt_0xff.expr(),
    meta.query_advice(bits[3], Rotation(2)),
    meta.query_advice(bits[3], Rotation(3)),
  ),
);

```

Figure 10.2: [aggregator/src/aggregation/decoder.rs#L378–L433](#)

## Exploit Scenario

A malicious prover generates a proof for a zstd blob that does not follow the specification. Due to the missing constraint, the verifier still accepts the proof as valid.

## Recommendations

Short term, add the necessary `require_zero` constraints to ensure that decoding is compliant with the specification.

Long term, add positive and negative tests for all edge cases in the specification.

## 11. The `tag_config.is_change` witness is partially unconstrained

Severity: Informational

Difficulty: N/A

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-11

Target: `aggregator/src/aggregation/decoder.rs`

### Description

The constraints for the `tag_config.is_change` witness are not immediately obvious, making it difficult for the reader to know if it is correctly constrained.

The `tag_config.is_change` witness is constrained to be true whenever `byte_idx_delta && tag_idx_eq_tag_len_prev` holds. However, this implies only that these two conditions are sufficient for `is_change` to be true, where if the conditions are met, then `is_change == true`.

```
cb.condition(and::expr([byte_idx_delta, tag_idx_eq_tag_len_prev]), |cb| {
    cb.require_equal(
        "is_change is set",
        meta.query_advice(config.tag_config.is_change, Rotation::cur()),
        1.expr(),
    );
});
```

Figure 11.1: `aggregator/src/aggregation/decoder.rs#L1322-L1328`

The constraint for the necessary part is constrained only later, where `is_change == true` implies the two conditions.

```
meta.create_gate("DecoderConfig: new tag", |meta| {
    let condition = and::expr([
        meta.query_fixed(config.q_enable, Rotation::cur()),
        meta.query_advice(config.tag_config.is_change, Rotation::cur()),
    ]);

    let mut cb = BaseConstraintBuilder::default();

    // The previous tag was processed completely.
    cb.require_equal(
        "tag_idx::prev == tag_len::prev",
        meta.query_advice(config.tag_config.tag_idx, Rotation::prev()),
        meta.query_advice(config.tag_config.tag_len, Rotation::prev()),
    );

    // Tag change also implies that the byte_idx transition did happen.
```

```

cb.require_equal(
  "byte_idx::prev + 1 == byte_idx::cur",
  meta.query_advice(config.byte_idx, Rotation::prev()) + 1.expr(),
  meta.query_advice(config.byte_idx, Rotation::cur()),
);

```

Figure 11.2: *aggregator/src/aggregation/decoder.rs#L1358-L1378*

## Exploit Scenario

A malicious prover can control the `is_change` witness when its value should be zero. By setting it to one when it should be zero, a malicious prover could bypass constraints related to `ZstdBlockSequenceFseCode` because they are constrained by `not(tag_config.is_change)`:

```

meta.create_gate(
  "DecoderConfig: tag ZstdBlockSequenceFseCode (other rows)",
  |meta| {
    let condition = and::expr([
      meta.query_fixed(q_enable, Rotation::cur()),
      meta.query_advice(config.tag_config.is_fse_code, Rotation::cur()),
      not::expr(meta.query_advice(config.tag_config.is_change,
Rotation::cur()))),
      not::expr(
        meta.query_advice(config.fse_decoder.is_trailing_bits,
Rotation::cur()),
      ),
    ]);
  });

```

Figure 11.3: *aggregator/src/aggregation/decoder.rs#2230-2240*

A malicious prover could also bypass lookups into the `VariableBitPacking` table:

```

meta.lookup_any(
  "DecoderConfig: tag ZstdBlockSequenceFseCode (variable bit-packing)",
  |meta| {
    // At every row where a non-nil bitstring is read:
    // - except the AL bits (is_change=true)
    // - except when we are in repeat-bits loop
    // - except the trailing bits (if they exist)
    let condition = and::expr([
      meta.query_fixed(config.q_enable, Rotation::cur()),
      meta.query_advice(config.tag_config.is_fse_code, Rotation::cur()),
      config.bitstream_decoder.is_not_nil(meta, Rotation::cur()),
      not::expr(meta.query_advice(config.tag_config.is_change,
Rotation::cur()))),
      not::expr(
        meta.query_advice(config.fse_decoder.is_repeat_bits_loop,
Rotation::cur()),
      ),
      not::expr(
        meta.query_advice(config.fse_decoder.is_trailing_bits,

```



```

Rotation::cur()),
    ),
  ]);
[...]
```

```

    let range = table_size - probability_acc + 1.expr();
    [
      FixedLookupTag::VariableBitPacking.expr(),
      range,
      value_read,
      value_decoded,
      num_bits,
      0.expr(),
      0.expr(),
    ]
    .into_iter()
    .zip_eq(config.fixed_table.table_exprs(meta))
    .map(|(arg, table)| (condition.expr() * arg, table))
    .collect()
  },
);

```

Figure 11.4: [aggregator/src/aggregation/decoder.rs#2552-2597](#)

There are several other variables whose constraint soundness depends on the soundness of `tag_config.is_change`, meaning an accidental change in these two necessary and sufficient conditions may undermine the security of many circuit components.

## Recommendations

Short term, document where the necessary and sufficient checks are performed, and therefore, constraints for `is_change == false` are not needed.

## 12. The is\_llt/is\_mot/is\_mlt constraints are only valid if self.table\_kind is in {1, 2, 3}

Severity: High

Difficulty: High

Type: Data Validation

Finding ID: TOB-SCROLLZSTD-12

Target: aggregator/src/aggregation/decoder.rs

### Description

The implementation of the `is_llt`, `is_mot` and `is_mlt` functions relies on the `table_kind` witness being in the set `{1, 2, 3}`. It is not immediately clear from the implementation that the `table_kind` variable is constrained to that set.

Upon reporting this finding, the Scroll team identified that the `table_kind` witness is unconstrained on a trailing bits row. Without this constraint, a malicious prover can set an incorrect `table_kind` and then set the next `table_kind` value incorrectly (by essentially skipping one step).

```
impl FseDecoder {
    fn is_llt(&self, meta: &mut VirtualCells<Fr>, rotation: Rotation) ->
    Expression<Fr> {
        let table_kind = meta.query_advice(self.table_kind, rotation);
        let invert_of_2 = Fr::from(2).invert().expect("infallible");
        (FseTableKind::MLT.expr() - table_kind.expr())
        * (FseTableKind::MOT.expr() - table_kind.expr())
        * invert_of_2
    }

    fn is_mot(&self, meta: &mut VirtualCells<Fr>, rotation: Rotation) ->
    Expression<Fr> {
        let table_kind = meta.query_advice(self.table_kind, rotation);
        (table_kind.expr() - FseTableKind::LLT.expr())
        * (FseTableKind::MLT.expr() - table_kind.expr())
    }

    fn is_mlt(&self, meta: &mut VirtualCells<Fr>, rotation: Rotation) ->
    Expression<Fr> {
        let table_kind = meta.query_advice(self.table_kind, rotation);
        let invert_of_2 = Fr::from(2).invert().expect("infallible");
        (table_kind.expr() - FseTableKind::LLT.expr())
        * (table_kind.expr() - FseTableKind::MOT.expr())
        * invert_of_2
    }
}
```

Figure 12.1: aggregator/src/aggregation/decoder.rs#L747-L768

There are lookups into tables that correctly constrain `table_kind`, such as the lookup into the `PredefinedFse` table:

```
// For predefined FSE tables, we must validate against the ROM predefined table
// fields for
// every state in the FSE table.
meta.lookup_any("FseTable: predefined table validation", |meta| {
  let condition = and::expr([
    meta.query_fixed(q_enable, Rotation::cur()),
    meta.query_advice(config.sorted_table.is_predefined, Rotation::cur()),
    not::expr(meta.query_advice(config.is_skipped_state, Rotation::cur())),
    not::expr(meta.query_advice(config.is_padding, Rotation::cur())),
  ]);

  let (table_kind, table_size, state, symbol, baseline, nb) = (
    meta.query_advice(config.sorted_table.table_kind, Rotation::cur()),
    meta.query_advice(config.sorted_table.table_size, Rotation::cur()),
    meta.query_advice(config.state, Rotation::cur()),
    meta.query_advice(config.symbol, Rotation::cur()),
    meta.query_advice(config.baseline, Rotation::cur()),
    meta.query_advice(config.nb, Rotation::cur()),
  );

  [
    FixedLookupTag::PredefinedFse.expr(),
    table_kind,
    table_size,
    state,
    symbol,
    baseline,
    nb,
  ]
  .into_iter()
  .zip_eq(fixed_table.table_exprs(meta))
  .map(|(arg, table)| (condition.expr() * arg, table))
  .collect()
});
```

Figure 12.2: *aggregator/src/aggregation/decoder/tables/fse.rs#L590–L622*

However, the set of conditions under which these lookups occur differs from other sets of conditions where the `is_llt`, `is_mot` and `is_mlt` functions are used (either explicitly or implicitly), which could lead to soundness issues.

## Recommendations

Short term, add explicit constraints to the `table_kind` witness, and ensure that the witness is constrained in every case.

Long term, add negative tests to ensure that incorrect or unexpected values of `table_kind` do not satisfy the constraints of the decoder circuit.

### 13. Values larger than 23 satisfy the "spans\_three\_bytes" constraints

Severity: Informational

Difficulty: N/A

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-13

Target: aggregator/src/aggregation/decoder.rs

#### Description

The `spans_three_bytes` function should constrain the `bit_index_end` witness to lie in the `[16, 23]` interval. However, it accepts a value as long as `bit_index_end > 15`. In particular, values larger than 23 will satisfy the constraint.

```
/// A bitstring spans 3 bytes if the bit_index at which it ends is such that:
/// - 16 <= bit_index_end <= 23.
fn spans_three_bytes(&self, meta: &mut VirtualCells<Fr>, at: Rotation) ->
Expression<Fr> {
    let lhs = meta.query_advice(self.bit_index_end, at);
    let (lt2, eq2) = self.bit_index_end_cmp_15.expr_at(meta, at, lhs, 15.expr());
    not::expr(lt2 + eq2)
}
```

Figure 13.1: aggregator/src/aggregation/decoder.rs#L637–L643

In practice, this issue is unexploitable—thus the informational severity—because the current `BitstringTable` implementation does not support bitstrings spanning more than three bytes. However, we still highly recommend correctly constraining the `bit_index_end` witness: if the bitstream decoder starts supporting more than three bytes in the future, the missing constraint would cause soundness issues.

#### Recommendations

Short term, add a constraint enforcing that the `bit_index_end` witness is smaller than 24.

## 14. Missing a large number of test cases that should fail

Severity: Low

Difficulty: N/A

Type: Cryptography

Finding ID: TOB-SCROLLZSTD-14

Target: aggregator/src/aggregation/\*

### Description

Overall, we found that there is a severe lack of testing, with a total of only 17 tests that do not comprehensively test the functionality of the decoder circuit. Not only should tests cover the expected success cases, but they should also ensure that well-defined failure modes are not possible within the confines of the system. We have compiled a non-exhaustive list of such test cases below:

- Boolean witnesses should not satisfy the witness generator if they are assigned non-Boolean values. This type of test is most valuable before implementing the custom wrapper type described in [TOB-SCROLLZSTD-1](#), but tests similar in nature are still recommended for similarly constrained values.
- Add tests for all valid RomTagTransition pairs, ensuring compatibility with the zstd specification ([TOB-SCROLLZSTD-4](#)).
- Add tests for various BlockHeader configurations, ensuring configurations that fall outside the specification are rejected ([TOB-SCROLLZSTD-8](#)).
- Add tests for the reserved compression mode bits ([TOB-SCROLLZSTD-10](#)).
- Add tests that try to insert an invalid table\_kind value ([TOB-SCROLLZSTD-12](#)).
- Add randomized round trip encoder tests. There is currently a single test that checks the satisfiability of the encoding of a known string. However, this test is also for a singular string and is not easily generalized to other input sizes. Randomized round trip testing can better guarantee the correctness of the encoding and decoding steps by easily generating larger inputs.

We also recommend the addition of fuzz testing. In large systems like these, it is hard to systematically test every edge case of the system. Fuzz testing enables the user to automate the randomization of inputs, leading to potentially higher constraint coverage.

### Recommendations

Short term, add additional unit tests for the failure modes listed in the above issues.

Long term, add support for fuzzing of the witness generator, where a valid witness can have some cells perturbed. If any of these small perturbations leads to another valid witness assignment, there is a high likelihood of a missing constraint somewhere.

## A. Vulnerability Categories

---

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.



## B. Code Maturity Categories

The following tables describe the code maturity categories and rating criteria used in this document.

Code Maturity Categories	
Category	Description
Arithmetic	The proper use of mathematical operations and semantics
Auditing	The use of event auditing and logging to support monitoring
Authentication / Access Controls	The use of robust access controls to handle identification and authorization and to ensure safe interactions with the system
Complexity Management	The presence of clear structures designed to manage system complexity, including the separation of system logic into clearly defined functions
Cryptography and Key Management	The safe use of cryptographic primitives and functions, along with the presence of robust mechanisms for key generation and distribution
Documentation	The presence of comprehensive and readable codebase documentation
Memory Safety and Error Handling	The presence of memory safety and robust error-handling mechanisms
Testing and Verification	The presence of robust testing procedures (e.g., unit tests, integration tests, and verification methods) and sufficient test coverage

Rating Criteria	
Rating	Description
Strong	No issues were found, and the system exceeds industry standards.
Satisfactory	Minor issues were found, but the system is compliant with best practices.
Moderate	Some issues that may affect system safety were found.
Weak	Many issues that affect system safety were found.
Missing	A required component is missing, significantly affecting system safety.
Not Applicable	The category is not applicable to this review.
Not Considered	The category was not considered in this review.

**Further  
Investigation  
Required**

Further investigation is required to reach a meaningful conclusion.

## C. Code Quality Findings

---

We identified the following code quality issues through manual and automated code review.

- **Panic in functions that returns Result.** The `synthesize` function in `compression/circuit.rs` returns `Result<(), Error>`, but line 104 is an `expect` call, which will result in a panic. Consider translating this to an error that is returned to higher-level logic for possible recovery or improved error reporting.
- **Duplicated byte reconstruction logic in BitstringTable.** The `BitstringTable::configure` function creates several constraints that validate that bits 0..7, 8..15, and 16..23 are equal to `byte1`, `byte2`, and `byte3`, respectively. However, these 8 bit-to-byte calculations are done manually, and it may make more sense to refactor the LE/BE byte reconstruction into a reusable subroutine to decrease the chances of copy-paste errors when the table is expanded to more bytes in the future.
- **Special casing of 1 and 2 byte BitstringTables.** The `BitstringTable` has a `const` generic parameter that expects a value of 1, 2, or 3, with 3 being the case we are most interested in. This results in additional special casing requiring several checks against `N_BYTES`. Making this struct more generic by consolidating `byte1`, `byte2`, and `byte3` into `byte: [Column<Advice>; N_BYTES]`, as well as writing more generic code, could make further expansion easier and less likely to introduce a bug.
- **Typo on constraint label.** The word “kee” should be “keep”:

```
cb.condition(not::expr(is_inst_begin.expr()), |cb| {
    cb.require_equal(
        "backref offset kee the same in one inst",
        backref_offset.expr(),
        backref_offset_prev.expr(),
    )
});
```

Figure C.1: [aggregator/src/aggregation/decoder/seq\\_exec.rs#507–513](#)

- **Redundant iteration over large vector.** The implementation of the `RomVariableBitPacking::values` function first constructs a large vector of `RomVariableBitPacking` elements, and then iterates over it to construct the result vector of 7-tuples. Instead, to prevent having to iterate over the vector twice, define an auxiliary function that returns the 7-tuple of values. Then, use that auxiliary function to construct the intended result in the first iteration.

```

impl FixedLookupValues for RomVariableBitPacking {
    fn values() -> Vec<[Value<Fr>; 7]> {
        // The maximum range R we ever have is 512 ( $1 \ll 9$ ) as the maximum
        // possible accuracy log is
        // 9. So we only need to support a range up to  $R + 1$ , i.e. 513.
        let rows = (1..=513)
            .flat_map(|range| {
                // Get the number of bits required to represent the highest
                // number in this range.
                let size = bit_length(range) as u32;
                let max = 1 << size;

                // Whether ``range`` is a power of 2 minus 1, i.e.  $2^k - 1$ .
                In these cases, we
                // don't need variable bit-packing as all values in the range
                // can be represented by
                // the same number of bits.
                let is_no_var = range & (range + 1) == 0;

                // The value read is in fact the value decoded.
                if is_no_var {
                    return (0..=range)
                        .map(|value_read| RomVariableBitPacking {
                            range,
                            value_read,
                            value_decoded: value_read,
                            num_bits: size as u64,
                        })
                        .collect::<Vec<_>>();
                }

                let n_total = range + 1;
                let lo_pin = max - n_total;
                let n_remaining = n_total - lo_pin;
                let hi_pin_1 = lo_pin + (n_remaining / 2);
                let hi_pin_2 = max - (n_remaining / 2);

                (0..max)
                    .map(|value_read| {
                        // the value denoted by the low (size - 1)-bits.
                        let lo_value = value_read & ((1 << (size - 1)) - 1);
                        let (num_bits, value_decoded) = if
(0..lo_pin).contains(&lo_value) {
                            (size - 1, lo_value)
                        } else if (lo_pin..hi_pin_1).contains(&value_read) {
                            (size, value_read)
                        } else if (hi_pin_1..hi_pin_2).contains(&value_read)
{
                            (size - 1, value_read - hi_pin_1)
                        } else {
                            assert!((hi_pin_2..max).contains(&value_read));
                            (size, value_read - lo_pin)
                        };
                    })
            })
    }
}

```

```

        RomVariableBitPacking {
            range,
            value_read,
            value_decoded,
            num_bits: num_bits.into(),
        }
    })
    .collect::

```

Figure C.2:

*aggregator/src/aggregation/decoder/tables/fixed/variable\_bit\_packing.rs#15–86*

- **Excessive function length.** The `process_sequences` function is over 1,100 lines. This hinders maintainability and review. Related portions of the function should be broken into their own routines.
- **Vacuous constraint on `is_inst_begin`.** The witness `is_inst_begin` must be constrained as a Boolean. However, the implementation includes a redundant constraint on `is_inst_begin`. The constraint is ineffective since it is applied to a variable unrelated to the `SeqExecConfig`. Furthermore, an effective constraint is defined in the subsequent lines of code.

```

// boolean constraint that index is increment
cb.require_boolean("instruction border is boolean", is_inst_begin.expr());

[...]

is_inst_begin = select::expr(
    is_block_begin.expr(),
    1.expr(),
    meta.query_advice(seq_index, Rotation::cur())
)

```

```

        - meta.query_advice(seq_index, Rotation::prev()),
    );

    cb.require_boolean("inst border is boolean", is_inst_begin.expr());

```

Figure C.3: *aggregator/src/aggregation/decoder/seq\_exec.rs#339–354*

- Redundant initialization of witness values.** The first row of the table associated with SeqExecConfig is initialized by assigning constants to all columns of the table. The columns decoded\_len, decoded\_rlc, and block\_index are assigned the constant F::zero() twice; furthermore, the assignments are performed with two different methods: assign\_advice\_from\_constant and assign\_advice. The assign\_advice\_from\_constant method should be preferred for initializing the first row, it additionally increases the readability of the codebase.

```

for col in [
    self.decoded_byte,
    self.decoded_len,
    self.decoded_rlc,
    [...],
    self.backref_progress,
] {
    region.assign_advice(|| "top row fluash", col, offset, ||
Value::known(F::zero()))?;
}

for (col, val) in [
    (self.decoded_len, F::zero()),
    (self.decoded_rlc, F::zero()),
    (self.block_index, F::zero()),
] {
    region.assign_advice_from_constant(|| "top row constraint", col, offset,
val)?;
}

```

Figure C.4: *aggregator/src/aggregation/decoder/seq\_exec.rs#912–934*

- Potentially unused fixed column.** Despite the comment stating that the column was used in SeqExecConfig, we did not find any use for it.

```

pub struct DecoderConfig<const L: usize, const R: usize> {
    /// constant column required by SeqExecConfig.
    _const_col: Column<Fixed>,
}

```

Figure C.5: *aggregator/src/aggregation/decoder.rs#45–47*

- Duplicate code comment prevents documentation rendering.** Remove the initial // from the code comment.

```

// /// Helper table in the "output" region for accumulating the result of

```

```

executing sequences.
sequence_execution_config: SequenceExecutionConfig<Fr>,

```

Figure C.6: *aggregator/src/aggregation/decoder.rs#104–105*

- **Usage of IsEqualConfig could be replaced by IsZeroConfig.** We found three locations where the IsEqualConfig gadget is used to compare against 0. Instead, use the IsZeroConfig gadget.

```

is_empty_sequences: IsEqualConfig<Fr>,
    /// For sequence decoding, the tag=ZstdBlockSequenceHeader bytes tell us
    the Compression_Mode
    /// utilised for Literals Lengths, Match Offsets and Match Lengths. We
    expect only 2
    /// possibilities:
    /// 1. Predefined_Mode (value=0)
    /// 2. Fse_Compressed_Mode (value=2)
    ///
    /// Which means a single boolean flag is sufficient to take note of which
    compression mode is
    /// utilised for each of the above purposes. The boolean flag will be set
    if we utilise the
    /// Fse_Compressed_Mode.
    compression_modes: [Column<Advice>; 3],
}

impl BlockConfig {
    fn configure(meta: &mut ConstraintSystem<Fr>, q_enable: Column<Fixed>) ->
    Self {
        let num_sequences = meta.advice_column();
        Self {
            block_len: meta.advice_column(),
            block_idx: meta.advice_column(),
            is_last_block: meta.advice_column(),
            is_block: meta.advice_column(),
            num_sequences,
            is_empty_sequences: IsEqualChip::configure(
                meta,
                |meta| meta.query_fixed(q_enable, Rotation::cur()),
                |meta| meta.query_advice(num_sequences, Rotation::cur()),
                |_| 0.expr(),
            ),
        }
    }
}

```

Figure C.7: *aggregator/src/aggregation/decoder.rs#206–233*

```

value_decoded_eq_0: IsEqualChip::configure(
    meta,
    |meta| meta.query_fixed(q_enable, Rotation::cur()),
    |meta| meta.query_advice(value_decoded, Rotation::cur()),
    |_| 0.expr(),
),

```

Figure C.8: *aggregator/src/aggregation/decoder.rs#731–736*

```
baseline_0x00: IsEqualChip::configure(  
    meta,  
    |meta| meta.query_fixed(q_enable, Rotation::cur()),  
    |meta| meta.query_advice(baseline, Rotation::cur()),  
    |_| 0.expr(),  
),
```

Figure C.9: *aggregator/src/aggregation/decoder/tables/fse.rs#1376–1381*

- **Typo in code comment.** The word “bittsrng” should be “bitstring.”

```
/// The bit-index where the bittsrng begins. 0 <= bit_index_start < 8.  
bit_index_start: Column<Advice>,
```

Figure C.10: *aggregator/src/aggregation/decoder.rs#453–454*

- **Unused local variables.**

```
for i in 0..N_SNARKS {  
    for j in 0..DIGEST_LEN {  
        let mut t1 = Fr::default();  
        let mut t2 = Fr::default();  
        chunk_pi_hash_digests[i][j].value().map(|x| t1 = *x);  
        snark_inputs[i * DIGEST_LEN + j].value().map(|x| t2 = *x);  
    }  
}
```

Figure C.11: *aggregator/src/aggregation/circuit.rs#372–377*



## D. Automated Analysis Tool Configuration

We used the following tools to perform automated testing of the codebase.

### D.1. Semgrep

We used the static analyzer **Semgrep** to search for dangerous API patterns and weaknesses in the source code repository.

```
semgrep --metrics=off --sarif --config custom_rule_path.yml
```

*Figure D.1: The invocation command used to run Semgrep for each custom rule*

```
semgrep --metrics=off --sarif --config "p/trailofbits"
```

*Figure D.2: The invocation command used to run Semgrep with Trail of Bits' **public** rules*

### Expression is reconstrained rule

We wrote this rule after identifying finding **TOB-SCROLLZSTD-5**, trying to identify variants of the same issue.

```
rules:
- id: expression-reconstrained
  message: "An expression appears in the condition and again in a constraint. This
  could indicate a copy-paste error"
  languages: [rust]
  severity: ERROR
  patterns:
  - pattern: cb.condition(<... $X.expr() ...>, |cb| { <... $X.expr() ...>});
  - metavariable-pattern:
    metavariable: $X
    patterns:
    - pattern-not: "1"
```

*Figure D.3: Custom Semgrep rule that finds **TOB-SCROLLZSTD-5***

Besides the instance from **TOB-SCROLLZSTD-5**, this rule finds two other cases. Although neither is underconstrained, the value is reused in both condition and constraint, meaning that the value could be replaced by 1 in the constraint.

```
let is_padding_next =
  1.expr() - s_lit_cp_phase_next.expr() - s_back_ref_phase_next.expr();
cb.condition(is_padding.expr(), |cb| {
  cb.require_equal(
    "padding never change once actived",
    is_padding_next.expr(),
    is_padding.expr(),
  );
});
```

```
});
```

Figure D.4: `aggregator/src/aggregation/decoder/seq_exec.rs#407-415`

```
cb.condition(
    and::expr([
        s_last_lit_cp_phase_prev.expr(),
        not::expr(is_block_begin.expr()),
    ]),
    |cb| {
        cb.require_equal(
            "phase must keep activated until block end",
            s_last_lit_cp_phase_prev.expr(),
            s_last_lit_cp_phase.expr(),
        );
    },
);
```

Figure D.5: `aggregator/src/aggregation/decoder/seq_exec.rs#442-454`

## IsEqualChip rule

After finding a couple of instances where the `IsEqualChip` gadget could be replaced by the `IsZeroConfig` gadget, we wrote a custom Semgrep rule to find other instances.

```
rules:
- id: is-equal-chip
  message: "The IsEqualChip is used to compare a value to zero, when the IsZeroChip
  could be used instead."
  languages: [rust]
  severity: ERROR
  patterns:
    - pattern: IsEqualChip::configure(
        $X,
        $Y,
        $Z,
        |_| 0.expr(),
    )
```

Figure D.6: Custom Semgrep rule that finds the three instances described in the Code Quality appendix

## D.2. Clippy

The Rust linter `Clippy` can be installed using `rustup` by running the command `rustup component add clippy`. Invoking `cargo clippy -- -W clippy::pedantic` in the root directory of the project runs the tool with the pedantic ruleset.

```
cargo clippy -- -W clippy::pedantic
```

Figure D.7: The invocation command used to run Clippy in the codebase

Converting the output to the SARIF file format (e.g., with `clippy-sarif`) allows easy inspection of the results within an IDE (e.g., using VSCode's `SARIF Explorer` extension).