



# Scroll Euclid Upgrade Phase 1

## Security Assessment

April 4, 2025

*Prepared for:*

**Roy Lou**

Scroll

*Prepared by:* **Filipe Casal and Marc Ilunga**

# Table of Contents

---

<b>Table of Contents</b>	<b>1</b>
<b>Project Summary</b>	<b>2</b>
<b>Executive Summary</b>	<b>3</b>
<b>Project Goals</b>	<b>5</b>
<b>Project Targets</b>	<b>6</b>
<b>Project Coverage</b>	<b>7</b>
<b>Automated Testing</b>	<b>8</b>
<b>Summary of Findings</b>	<b>11</b>
<b>Detailed Findings</b>	<b>12</b>
1. The aggregation circuit does not validate the size of inner proofs	12
2. Insufficient validation of chunk information and blob bytes in the batch circuit	15
3. MPT migration code lacks unit tests	17
4. Incorrect data present on BundleInfo	18
5. Lax parsing of chunk data payload	19
6. Docker release action is vulnerable to cache poisoning	21
7. Unpinned external GitHub CI/CD action versions	23
8. Potential credential persistence in artifacts and stale GitHub action	25
<b>A. Vulnerability Categories</b>	<b>27</b>
<b>B. Code Quality Findings</b>	<b>29</b>
<b>C. Automated Analysis Tool Configuration</b>	<b>33</b>
<b>D. Fix Review Results</b>	<b>35</b>
Detailed Fix Review Results	36
<b>E. Fix Review Status Categories</b>	<b>37</b>
<b>About Trail of Bits</b>	<b>38</b>
<b>Notices and Remarks</b>	<b>39</b>

# Project Summary

---

## Contact Information

The following project manager was associated with this project:

**Jeff Braswell**, Project Manager

[jeff.braswell@trailofbits.com](mailto:jeff.braswell@trailofbits.com)

The following engineering director was associated with this project:

**Jim Miller**, Engineering Director, Cryptography

[james.miller@trailofbits.com](mailto:james.miller@trailofbits.com)

The following consultants were associated with this project:

**Filipe Casal**, Consultant

[filipe.casal@trailofbits.com](mailto:filipe.casal@trailofbits.com)

**Marc Ilunga**, Consultant

[marc.ilunga@trailofbits.com](mailto:marc.ilunga@trailofbits.com)

## Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
February 21, 2025	Pre-project kickoff call
February 28, 2025	Status update meeting #1
March 6, 2025	Delivery of report draft and report readout meeting
April 4, 2025	Delivery of final comprehensive report

# Executive Summary

---

## Engagement Overview

Scroll engaged Trail of Bits to review the security of changes related to the Euclid Upgrade. These changes include OpenVM guest programs for the chunk, batch, and bundle verification circuits, changes to the Scroll smart contracts related to supporting a special state transition due to the change in the computation of the state root, and a Go command-line tool to validate states during the Euclid transition.

A team of two consultants conducted the review from February 24 to March 5, 2025, for a total of three engineer-weeks of effort. Our testing efforts focused on assessing the soundness and correctness of the zero-knowledge circuits, whether the Euclid finalization function works as specified, and whether the Go command-line tool can validate that the state of two nodes is identical at a certain block height. With full access to source code and documentation, we performed static and dynamic testing of the codebase, using automated and manual processes.

## Observations and Impact

We identified two high-severity soundness issues ([TOB-SCREUC-1](#) and [TOB-SCREUC-2](#)) affecting the batch and bundle circuits that would allow obtaining a bundle proof for which a malicious prover could bypass the verification of the inner proofs for the chunk and batch circuits.

We identified two impactful but difficult-to-exploit issues ([TOB-SCREUC-6](#) and [TOB-SCREUC-7](#)) related to the GitHub actions used in the codebases that could undermine the security of the released images to DockerHub. Note that GitHub actions security was not in the audit's scope; we identified these issues using a static analysis tool for GitHub actions.

We did not identify security issues in the smart contract related to the one-off Euclid finalization functionality or in the Go command-line utility. However, we found that the migration checker was tested only with test data and does not have associated unit tests for its subfunctions (see [TOB-SCREUC-3](#)).

## Recommendations

Based on the findings identified during the security review, Trail of Bits recommends that Scroll take the following steps:

- **Remediate the findings disclosed in this report.** These findings should be addressed as part of a direct remediation or any refactor that may occur when addressing other recommendations.

- **Add tests for the migration checker.** The migration checker consists of several functions, each performing a specific verification test. These components are relatively self-contained and, therefore, can be tested against correct and incorrect inputs. Sufficient testing of the migration checker will help increase confidence in its correctness.

## Finding Severities and Categories

The following tables provide the number of findings by severity and category.

### EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	2
Medium	0
Low	2
Informational	4
Undetermined	0

### CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Auditing and Logging	1
Configuration	2
Cryptography	2
Data Validation	2
Testing	1

# Project Goals

---

The engagement was scoped to provide a security assessment of the Scroll Euclid phase 1 code changes. Specifically, we sought to answer the following non-exhaustive list of questions:

- Are there any soundness or correctness issues in the chunk, batch, and bundle verification circuits?
- Is the integration with the OpenVM verifier correctly implemented on the smart contract?
- Does the Go command-line tool correctly validate whether a given MPT and a ZKTrie correspond to the same inherent state?
- Can the Security Council correctly finalize the Euclid transition batch?
- Can the Euclid finalize function be maliciously used to finalize batches other than the transition batch?
- Does batch reverting correctly reset the Euclid batch index?
- Does the `finalizeBundleWithProof` function correctly prevent v4, v5, and v6 batches in the same bundle?

# Project Targets

---

The engagement involved a review and testing of the targets listed below.

## zkvm-prover

Repository      <https://github.com/scroll-tech/zkvm-prover/>  
Version          da9c824e9a0751bd5482d00910a187e86350beee  
Type             Rust, OpenVM

## scroll-contracts

Repository      <https://github.com/scroll-tech/scroll-contracts/>  
Version          c382953479f9a919e850e56115924e5711810dfd  
Type             Solidity

## go-ethereum

Repository      <https://github.com/scroll-tech/go-ethereum>  
Version          642c35f9906cab9001ac77190e198d18e2a044b5  
Type             Go

# Project Coverage

---

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- **Zkvm-prover.** We manually reviewed the chunk, batch, and bundle circuits, focusing on identifying potential soundness and correctness issues.
- **Smart contracts.** We manually reviewed the smart contracts, focusing on determining if the OpenVM verifier is correctly integrated to verify the zero-knowledge proof; if the Security Council can perform the Euclid state transition using the implemented functionality; if this functionality can be maliciously used after the Euclid transition occurs; and if batch reversion correctly updates the `initialEuclidBatchIndex` state variable. We additionally assessed whether the bundle proof and public input are stored in memory according to the memory layout in the documentation.
- **Migration checker.** We manually reviewed the migration checker tool, focusing on assessing whether the code faithfully checks that a zkTrie and a Merkle Patricia Trie both encode the same chain state.



# Automated Testing

---

Trail of Bits uses automated techniques to extensively test the security properties of software. We use both open-source static analysis and fuzzing utilities, along with tools developed in-house, to perform automated testing of source code and compiled software.

## Test Harness Configuration

We used the following tools in the automated testing phase of this project:

Tool	Description	Policy
<a href="#">Semgrep</a>	An open-source static analysis tool for finding bugs and enforcing code standards when editing or committing code and during build time	<a href="#">Appendix C</a>
<a href="#">Clippy</a>	An open-source Rust linter used to catch common mistakes and unidiomatic Rust code	<a href="#">Appendix C</a>
<a href="#">Dylint</a>	A tool for running Rust lints from dynamic libraries	<a href="#">Appendix C</a>
<a href="#">cargo-audit</a>	An open-source tool for checking dependencies against the RustSec advisory database	<a href="#">Appendix C</a>
<a href="#">cargo-edit</a>	A tool for quickly identifying outdated crates	<a href="#">Appendix C</a>
<a href="#">zizmor</a>	A static analysis tool for GitHub Actions	<a href="#">Appendix C</a>
<a href="#">CodeQL</a>	A code analysis engine developed by GitHub to automate security checks	<a href="#">Appendix C</a>

## Areas of Focus

Our automated testing and verification work focused on the following:

- General code quality issues and unidiomatic code patterns
- Security of CI and GitHub Actions
- Identifying dependencies that are outdated or vulnerable to known attacks

## Test Results

The results of this focused testing are detailed below.

### Clippy

Clippy reports warnings for three idiomatic and performance rules: [manual assert](#), [or\\_fun\\_call](#) and [explicit\\_iter\\_loop](#) which we detailed in the [Code Quality appendix](#). We recommend adding these rules to your regular Clippy runs.

### zizmor

Running zizmor on the codebase identified findings [TOB-SCREUC-6](#), [TOB-SCREUC-7](#), and [TOB-SCREUC-8](#). We recommend adding zizmor to your CI/CD pipeline across the organization's repositories.

### cargo-edit

Cargo-edit identified several outdated dependencies, including some that are, in a semantic versioning sense, compatible with the version currently in use, and others that are not. Consider [integrating Dependabot](#) to automatically update dependencies as soon as these are available.

```
$ cargo upgrade --incompatible --dry-run
  Checking virtual workspace's dependencies
name      old req compatible latest new req
====      =====
bitcode    0.6.3    0.6.5      0.6.5  0.6.5
halo2curves-axiom 0.5.3    0.5.3      0.7.0  0.7.0
metrics    0.23.0   0.23.0     0.24.1 0.24.1
metrics-util 0.17     0.17.0     0.19.0 0.19
alloy-serde 0.8      0.8.3      0.11.1 0.11
serde_with 3.11.0   3.12.0     3.12.0 3.12.0
toml       0.8.14   0.8.20     0.8.20 0.8.20
  Checking scroll-zkvm-batch-circuit's dependencies
  Checking scroll-zkvm-build-guest's dependencies
name      old req compatible latest new req
====      =====
cargo_metadata 0.19.1  0.19.2     0.19.2 0.19.2
  Checking scroll-zkvm-bundle-circuit's dependencies
  Checking scroll-zkvm-chunk-circuit's dependencies
  Checking scroll-zkvm-circuit-input-types's dependencies
  Checking scroll-zkvm-integration's dependencies
name  old req compatible latest new req
====  =====
c-kzg 1.0    1.0.3     2.0.0  2.0
  Checking scroll-zkvm-prover's dependencies
name      old req compatible latest new req
====      =====
```

```

git-version 0.3.5 0.3.9 0.3.9 0.3.9
revm 19.0 19.5.0 19.5.0 19.5
    Checking scroll-zkvm-verifier's dependencies
name old req compatible latest new req
==== =====
revm 19.0 19.5.0 19.5.0 19.5

```

*Figure T.1: Result of running cargo-edit on the zkvm-prover codebase*

### **cargo-audit**

cargo-audit did not identify any dependency vulnerable to known security issues.

### **CodeQL**

We ran CodeQL on the migration checker. CodeQL did not identify any vulnerabilities.

## Summary of Findings

The table below summarizes the findings of the review, including details on type and severity.

ID	Title	Type	Severity
1	The aggregation circuit does not validate the size of inner proofs	Cryptography	High
2	Insufficient validation of chunk information and blob bytes in the batch circuit	Cryptography	High
3	MPT migration code lacks unit tests	Testing	Informational
4	Incorrect data present on BundleInfo	Data Validation	Informational
5	Lax parsing of chunk data payload	Data Validation	Informational
6	Docker release action is vulnerable to cache poisoning	Configuration	Low
7	Unpinned external GitHub CI/CD action versions	Auditing and Logging	Low
8	Potential credential persistence in artifacts and stale GitHub action	Configuration	Informational

# Detailed Findings

## 1. The aggregation circuit does not validate the size of inner proofs

Severity: High

Difficulty: Low

Type: Cryptography

Finding ID: TOB-SCREUC-1

Target: `circuits/types/src/lib.rs`, `circuits/batch-circuit/src/main.rs`, `circuits/bundle-circuit/src/main.rs`

### Description

The batch and bundle circuits are aggregation circuits. They aggregate proofs generated by inner circuits. They validate the aggregated public inputs and verify the proofs embedded in the witness. However, due to missing enforcement of the lengths of aggregated proofs, an aggregation circuit would verify an aggregated public input without the associated inner proof.

Figure 1.1 shows the main function of the batch circuit. The function calls `C::verify_proofs(witness)`, which verifies the inner proof embedded in the witness. Subsequently, the aggregated public inputs are validated against the aggregated public input hashes.

```
fn main() {
    // Setup openvm extensions for the circuit.
    C::setup();
    [...]
    // Verify the root proofs being aggregated in the circuit.
    let agg_proofs = C::verify_proofs(witness);

    // Get the public-input values of the proofs being aggregated from witness.
    let agg_pis = C::aggregated_public_inputs(witness);

    // Derive the digests of the public-input values of proofs being aggregated.
    let agg_pi_hashes = C::aggregated_pi_hashes(&agg_proofs);

    // Validate that the pi hashes derived from the root proofs are in fact the
    digests of the
    // public-input values of the previous circuit layer.
    C::validate_aggregated_pi(&agg_pis, &agg_pi_hashes);

    [...]
    // Reveal the public-input values of the current circuit layer.
    C::reveal_pi(&public_inputs);
}
```

```
}
```

*Figure 1.1: The main process of the batch circuit  
([circuits/batch-circuit/src/main.rs#L18-L47](#))*

Figure 1.2 shows that the proof verification function does not check the length of the received vector of proofs and can accept an empty vector. Therefore, the returned aggregated proofs, `agg_proofs`, are also empty.

```
/// Verify the proofs being aggregated.
///
/// Also returns the root proofs being aggregated.
fn verify_proofs(witness: &Self::Witness) -> Vec<proof::AggregationInput> {
    let proofs = witness.get_proofs();

    for proof in proofs.iter() {
        Self::verify_commitments(&proof.commitment);
        proof::verify_proof(&proof.commitment, proof.public_values.as_slice());
    }

    proofs
}
```

*Figure 1.2: Verification of aggregated proofs ([circuits/types/src/lib.rs#L69-L81](#))*

As a consequence, the validation of the aggregated public inputs against the public input hashes—the call to `C::validate_aggregated_pi(&agg_pis, &agg_pi_hashes)` shown in figure 1.1—effectively verifies the aggregated public input against an empty list of hashes. Due to the use of `zip`, the lengths are checked for equality. With an empty vector of hash values, an assertion error would occur, and the circuit would terminate successfully.

```
/// Validate that the public-input values of the aggregated proofs are well-formed.
///
/// - That the public-inputs of contiguous chunks/batches are valid
/// - That the public-input values in fact hash to the pi_hash values from the root
  proofs.
fn validate_aggregated_pi(agg_pis: &[Self::AggregatedPublicInputs], agg_pi_hashes:
&[B256]) {
    // Validation for the contiguous public-input values.
    for w in agg_pis.windows(2) {
        w[1].validate(&w[0]);
    }

    // Validation for public-input values hash being the pi_hash from root proof.
    for (agg_pi, &agg_pi_hash) in agg_pis.iter().zip(agg_pi_hashes.iter()) {
        assert_eq!(
            agg_pi.pi_hash(),
            agg_pi_hash,
            "pi hash mismatch between proofs and witness computed"
        );
    }
}
```

```
    }  
  }  
};
```

*Figure 1.3: Validation of aggregated public input against expected public input hashes  
([circuits/types/src/lib.rs#L89-L107](#))*

### Exploit Scenario

An attacker wishes to produce a wrong aggregated proof for the bundle input. They generate an empty witness and fabricate aggregated inputs, which the circuit verifies.

### Recommendations

Short term, ensure that the aggregated proofs are not empty. Use `zip_eq` when comparing elements expected to have the same size.

Long term, document invariants and assumptions that must hold in every part of the codebase.

## 2. Insufficient validation of chunk information and blob bytes in the batch circuit

Severity: High

Difficulty: Medium

Type: Cryptography

Finding ID: TOB-SCREUC-2

Target: `crates/circuits/batch-circuit/src/builder/v3.rs`

### Description

The witness validation in the batch circuit does not properly validate that the chunk information matches the provided blob bytes. This leads to the circuit returning a batch header hash that uses blob data proof inconsistent with the chunks in the batch.

Specifically, the batch circuit does not sufficiently validate that the chunk information transaction data digest matches the chunk data digest present in the EIP-4844 blob bytes, as it does not compare the length of the `chunks_info` slice with the length of the `payload.chunk_data_digests` slice.

```
// Validate the tx data is match with fields in chunk info
for (chunk_info, &tx_data_digest) in
    chunks_info.iter().zip(payload.chunk_data_digests.iter())
{
    assert_eq!(chunk_info.tx_data_digest, tx_data_digest);
}
```

Figure 2.1: `crates/circuits/batch-circuit/src/builder/v3.rs#L32-L37`

For example, if `chunks_info` contains information about two chunks but the `blob_bytes` has encoded only the first chunk, the `zip` iteration will validate only the first chunk data digest, and the blob consistency check will generate a blob data proof that concerns only the first chunk.

### Exploit Scenario

A malicious prover provides a witness value with a different number of chunks for the `chunk_info` and the `blob_bytes`, resulting in a batch information PI that is inconsistent: it contains a batch hash with the only the chunks in `blob_bytes`, but its state root and withdraw root concern the last chunk in `chunk_info`.

### Recommendations

Short term, ensure that the lengths of `chunks_info` and `payload.chunk_data_digests` match, either with a separate assertion or by using `zip_eq`.



Long term, review all uses of `zip` in the codebase for implicit assumptions on the length of the iterators and enforce those length checks where necessary.

### 3. MPT migration code lacks unit tests

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SCREUC-3

Target: go-ethereum/cmd/migration-checker/main.go

#### Description

As part of the Euclid upgrade, the Scroll team will migrate the chain state from the current zktrie to a Merkle Patricia Tree (MPT). The team has developed a tool in Go that allows the security council to check for a faithful migration of the chain's state. However, there are no unit tests to ensure that the migration checker is implemented correctly.

The migration checker code contains several subfunctions that perform specific verification tasks. Figure 3.1 shows a function used to check that storage nodes are equal across the zktrie and MPT data structures. The function scope is small enough and self-contained to enable unit tests exercising its behavior in both expected and adversarial scenarios.

```
func checkStorageEquality(label string, _ *dbs, zkStorageBytes, mptStorageBytes
[]byte, _ bool) {
    zkValue := common.BytesToHash(zkStorageBytes)
    _, content, _, err := rlp.Split(mptStorageBytes)
    panicOnError(err, label, "failed to decode mpt storage")
    mptValue := common.BytesToHash(content)
    if !bytes.Equal(zkValue[:], mptValue[:]) {
        panic(fmt.Sprintf("%s storage mismatch: zk: %s, mpt: %s", label,
zkValue.Hex(), mptValue.Hex()))
    }
}
```

Figure 3.1: (go-ethereum/cmd/migration-checker/main.go#158-166)

However, the Scroll team has tested the migration checker on some of its test nodes. Though this test was successful, it may not necessarily guarantee the robustness of the code against incorrect inputs.

#### Recommendations

Short term, write unit tests for the migration checker before deployment. Ensure that both correctness and robustness are tested appropriately.

Long term, ensure all critical software used for irreversible tasks is appropriately tested.

#### 4. Incorrect data present on BundleInfo

Severity: Informational

Difficulty: N/A

Type: Data Validation

Finding ID: TOB-SCREUC-4

Target: `crates/circuits/bundle-circuit/src/circuit.rs`

#### Description

The `BundleCircuit::validate` function returns an incorrect `prev_state_root` value in the `BundleInfo` data. Instead of `first_batch.parent_batch_hash`, the correct value is `first_batch.parent_state_root`.

```
let prev_state_root = first_batch.parent_batch_hash.into();  
let prev_batch_hash = first_batch.parent_batch_hash.into();
```

*Figure 4.1: `crates/circuits/bundle-circuit/src/circuit.rs#L57-L58`*

Given that the smart contract rederives the public inputs from independent data, incorrectly setting this value causes a correctness issue when verifying the proof on the smart contract side. The Scroll team has also independently identified this issue.

#### Recommendations

Short term, modify the code to use the correct field for `prev_state_root`.

Long term, document how each layer's public inputs are used. Add tests verifying the correctness of the returned public inputs. Enforce end-to-end testing from proving to smart contract verification to ensure end-to-end correctness.

## 5. Lax parsing of chunk data payload

Severity: Informational

Difficulty: N/A

Type: Data Validation

Finding ID: TOB-SCREUC-5

Target: crates/circuits/batch-circuit/src/payload.rs

### Description

The `Payload::from_payload` function does not fully validate the metadata chunk sizes when less than `N_MAX_CHUNKS` are present.

The function parses the payload by 1) obtaining the number of chunks present `valid_chunks`, 2) obtaining the chunk size of each chunk up to `N_MAX_CHUNKS` and considering only the first `valid_chunks` lengths, and finally, obtaining the chunk bytes based on the metadata lengths while validating that no extra bytes remain at the end of the payload.

However, this procedure does not validate that if, for example, only two chunks are present, all the remaining `N_MAX_CHUNKS-2` chunk lengths equal zero. This means that different metadata bytes can lead to the same chunk data digests, as the payload provider could state that the chunk lengths after chunk number two are arbitrary values, not changing the behavior of the rest of the parsing.

```
// Decoded batch bytes require segmentation based on chunk length
let valid_chunks = metadata_bytes[..N_BYTES_NUM_CHUNKS]
    .iter()
    .fold(0usize, |acc, &d| acc * 256usize + d as usize);

let chunk_size_bytes = metadata_bytes[N_BYTES_NUM_CHUNKS..]
    .iter()
    .chunks(N_BYTES_CHUNK_SIZE);
let chunk_lens = chunk_size_bytes
    .into_iter()
    .map(|chunk_bytes| chunk_bytes.fold(0usize, |acc, &d| acc * 256usize + d as
    usize))
    .take(valid_chunks);

// reconstruct segments
let (segmented_batch_data, final_bytes) = chunk_lens.fold(
    (Vec::new(), batch_bytes),
    |(mut datas, rest_bytes), size| {
        datas.push(Vec::from(&rest_bytes[..size]));
        (datas, &rest_bytes[size..])
    },
);
```

```
);  
  
assert!(  
    final_bytes.is_empty(),  
    "chunk segmentation len must add up to the correct value"  
);
```

Figure 5.1: [crates/circuits/batch-circuit/src/payload.rs#L37-L62](#)

Given that both the metadata digest and chunk digests are used in the computation of the challenge digest, we could not identify a way of further exploiting this issue.

### Recommendations

Short term, validate that the size of the chunks past the number specified in `valid_chunks` equals zero, ensuring that each unique payload corresponds to a unique pair of metadata digests and batch data digests.

## 6. Docker release action is vulnerable to cache poisoning

Severity: Low

Difficulty: High

Type: Configuration

Finding ID: TOB-SCREUC-6

Target: go-ethereum/.github/workflows/{docker-arm64.yaml, docker.yaml}, scroll-contracts/.github/workflows/docker-release.yaml

### Description

The Docker release GitHub workflows use caching to store the buildx binary. An attacker who gains write access to the GitHub actions cache (e.g., through a valid GitHub token) can poison the cache and compromise the build process to maliciously alter build artifacts.

```
- name: Set up Docker Buildx
  id: buildx
  uses: docker/setup-buildx-action@v2
- name: Login to Docker Hub
  uses: docker/login-action@v2
  with:
    username: ${ secrets.DOCKERHUB_USERNAME }
    password: ${ secrets.DOCKERHUB_TOKEN }
- name: Build docker image
  uses: docker/build-push-action@v2
  with:
    platforms: linux/arm64
    context: .
    file: Dockerfile
    push: true
    tags: scrolltech/l2geth:${ inputs.tag }-arm64
```

Figure 6.1: *.github/workflows/docker-arm64.yaml#26-41*

```
jobs:
  build-and-push:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v2
      - name: Set up QEMU
        uses: docker/setup-qemu-action@v2
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v2
```

Figure 6.2: *.github/workflows/docker.yaml#10-19*

```
- name: Set up Docker Buildx
  uses: docker/setup-buildx-action@v2

- name: Login to Dockerhub
  uses: docker/login-action@v2
  with:
    username: ${ secrets.DOCKERHUB_USERNAME }
    password: ${ secrets.DOCKERHUB_TOKEN }
```

Figure 6.3: *.github/workflows/docker-release.yml#31-39*

## Exploit Scenario

An attacker gains unauthorized access to a GitHub token with write permissions to the GitHub actions cache. She uses this to poison the cache to inject a malicious `buildx` binary. When the Docker container is built, the final image is modified to download and run attacker-controlled software. Since the codebase is not tampered with directly, the attack goes unnoticed by the Scroll team.

## Recommendations

Short term, in general, avoid using previously cached CI state within workflows intended to publish build artifacts such as Docker containers. Caching can be disabled by **setting the `cache-binary` option to `false`**.

Long term, add **`zizmor`** to the CI/CD pipeline to detect GitHub action-related issues.

## 7. Unpinned external GitHub CI/CD action versions

Severity: Low

Difficulty: High

Type: Auditing and Logging

Finding ID: TOB-SCREUC-7

Target: `.github/workflows/{docker-arm64.yaml, docker.yaml}`,  
`.github/workflows/l2geth_ci.yaml`

### Description

Several GitHub actions workflows in the `go-ethereum` and `scroll-contracts` repositories use third-party actions pinned to a tag or branch name instead of a full commit SHA as **recommended by GitHub**. This configuration enables repository owners to silently modify the actions. A malicious actor could use this ability to tamper with an application release or leak secrets.

The following actions are owned by GitHub organizations or individuals that are not affiliated directly with Scroll:

- `foundry-rs/foundry-toolchain@v1`
- `hrishikesh-kadam/setup-lcov@v1`
- `codecov/codecov-action@v3`
- `docker/setup-qemu-action@v2`
- `docker/setup-buildx-action@v2`
- `docker/login-action@v2`
- `docker/build-push-action@v3`
- `docker/metadata-action@v3`
- `actions-rs/toolchain@v1`

```
jobs:
  foundry:
    if: github.event.pull_request.draft == false
    runs-on: ubuntu-latest

    steps:
      - name: Checkout sources
```



```
uses: actions/checkout@v4
with:
  submodules: recursive

- name: Install Foundry
  uses: foundry-rs/foundry-toolchain@v1
  with:
    version: nightly

- name: Setup LCOV
  uses: hrishikesh-kadam/setup-lcov@v1
```

Figure 7.1: *.github/workflows/contracts.yml#24-41*

## Exploit Scenario

An attacker gains unauthorized access to the account of a GitHub action owner. The attacker manipulates the action's code to steal GitHub credentials and compromise the repository.

## Recommendations

Short term, pin each third-party action to a specific full-length commit SHA, as **recommended by GitHub**. Additionally, **configure Dependabot** to update the actions' commit SHAs after reviewing their available updates.

Long term, add **zizmor** to the CI/CD pipeline.

## 8. Potential credential persistence in artifacts and stale GitHub action

Severity: Informational

Difficulty: N/A

Type: Configuration

Finding ID: TOB-SCREUC-8

Target: scroll-contracts/.github/workflows/{contracts.yml, docker-release.yml}, go-ethereum/.github/workflows/\*

### Description

The default behavior of the actions/checkout GitHub action is to persist credentials, which means that the GitHub token is written to the local repository directory. This allows the action to execute Git commands that require authentication. The credentials are stored in .git/config files in the local repository directory, where they could be inadvertently included in workflow artifacts or accessed by subsequent workflow steps.

```
name: Docker

on:
  push:
    tags:
      - '*'
  release:
    types: [published]

jobs:
  build-and-push:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v2
```

Figure 8.1: .github/workflows/docker.yml#1-15

The current Dockerfile does not appear to copy the .git/ folder to the image, meaning that this is not currently a security risk. However, if the Dockerfile is altered, credentials could be exposed in the Docker image.

We also observed that scroll-contracts/.github/workflows/docker-release.yml file references a Dockerfile that no longer exists in the repository, so it can be removed from the codebase.

### Recommendations

Short term, add persist-credentials: false to checkout actions that do not require privileged Git operations.

Long term, add **zizmor** to the CI/CD pipeline.

## References

- [ArtiPACKED: Hacking Giants Through a Race Condition in GitHub Actions Artifacts](#)
- [zizmor ArtiPACKED documentation](#)

## A. Vulnerability Categories

---

The following tables describe the vulnerability categories, severity, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

## B. Code Quality Findings

We identified the following code quality issues through manual and automatic code review. These issues are not associated with any specific vulnerabilities. However, they will enhance code readability and may prevent the introduction of vulnerabilities in the future.

- **Manual assert construction.** The following if statement plus panic statement is equivalent to an assert statement. Use [Clippy's manual assert](#) rule to find and fix all instances of this pattern in the codebase.

```
fn verify_commitments(commitment: &ProgramCommitment) {  
    if commitment.exe != CHUNK_EXE_COMMIT {  
        panic!(  
            "mismatch chunk-proof exe commitment: expected={:?}, got={:?}",  
            CHUNK_EXE_COMMIT, commitment.exe,  
        );  
    }  
    if commitment.leaf != CHUNK_LEAF_COMMIT {  
        panic!(  
            "mismatch chunk-proof leaf commitment: expected={:?}, got={:?}",  
            CHUNK_EXE_COMMIT, commitment.leaf,  
        );  
    }  
}
```

Figure B.1: [crates/circuits/batch-circuit/src/circuit.rs#46-59](#)

- **Use of ok\_or with function calls.** Arguments passed to ok\_or are eagerly evaluated and may lead to unnecessary computation. Since, in this case, a function is evaluated, we recommend using ok\_or\_else instead, as it will evaluate this function only if needed, preventing unnecessary computation. The codebase contains 16 instances of this pattern. Use [Clippy's or\\_fun\\_call](#) rule to find and fix all instances of this pattern.

```
let sbv_chunk_info =  
    SbvChunkInfo::from_blocks(chain_id, first_block.pre_state_root(),  
    &blocks);  
  
let chain_spec =  
    get_chain_spec(Chain::from_id(sbv_chunk_info.chain_id())).ok_or(  
        Error::GenProof(format!("{err_prefix}: failed to get chain spec")),  
    )?;
```

Figure B.2: [crates/prover/src/prover/chunk.rs#83-88](#)

- **Loop code could be simplified.** Some loops explicitly call `.iter()` when this is not needed. Figure B.3 shows a case where the `.iter()` call can be removed. Use [Clippy's explicit\\_iter\\_loop](#) rule to find and fix all instances of this pattern.

```

for block in blocks.iter() {
    let output = ManuallyDrop::new(
        EvmExecutor::new(chain_spec.clone(), &db, block)
            .execute()
            .expect("failed to execute block"),
    );
    db.update(&nodes_provider, output.state.state.iter())
        .expect("failed to update db");
}

```

Figure B.3: *crates/circuits/chunk-circuit/src/execute.rs#56-64*

- **Incorrect documentation on the BatchHeaderV3Codec contract.** The BatchHeaderV3Codec contract has two instances of incorrect documentation. The lastBlockTimestamp documentation mentions the L1 skipped messages, and the storeBlobDataProof documentation refers the timestamp:

```

/// * lastBlockTimestamp      8      uint64      121      A bitmap to
indicate which L1 messages are skipped in the batch
/// * blobDataProof           64      bytes64      129      The blob data
proof: z (32), y (32)

```

Figure B.4: *src/libraries/codec/BatchHeaderV3Codec.sol#17-18*

```

/// @notice Store the last block timestamp of batch header.
/// @param batchPtr The start memory offset of the batch header in memory.
/// @param blobDataProof The blob data proof: z (32), y (32)
function storeBlobDataProof(uint256 batchPtr, bytes calldata blobDataProof)
internal pure {

```

Figure B.5: *src/libraries/codec/BatchHeaderV3Codec.sol#66-69*

- **Imprecise documentation on the ScrollChain contract.** The error is raised when the parent batch hash of the genesis block is non-zero, instead of the state zero:

```

/// @dev Thrown when the parent batch hash in genesis batch is zero.
error ErrorGenesisParentBatchHashIsNonZero();

```

Figure B.6: *src/L1/rollup/ScrollChain.sol#60-61*

- **Imprecise documentation of \_loadBatchHeader.** The documentation states that \_totalL1MessagesPoppedOverall is a parameter, but it is a return value.

```

/// @return _batchIndex The index of this batch.
/// @param _totalL1MessagesPoppedOverall The number of L1 messages popped
after this batch.
function _loadBatchHeader(bytes calldata _batchHeader)
internal
view
virtual

```

```

returns (
    uint256 batchPtr,
    bytes32 _batchHash,
    uint256 _batchIndex,
    uint256 _totalL1MessagesPoppedOverall
)

```

Figure B.7: `src/L1/rollup/ScrollChain.sol#775-786`

- **Panic messages use incorrect data.** The panic messages incorrectly reference BATCH\_EXE\_COMMIT and CHUNK\_EXE\_COMMIT when compared against BATCH\_LEAF\_COMMIT and CHUNK\_LEAF\_COMMIT.

```

if commitment.leaf != BATCH_LEAF_COMMIT {
    panic!(
        "mismatch batch-proof leaf commitment: expected={:?}, got={:?}",
        BATCH_EXE_COMMIT, commitment.leaf,
    );
}

```

Figure B.8: `crates/circuits/bundle-circuit/src/circuit.rs#L85-L90`

```

if commitment.leaf != CHUNK_LEAF_COMMIT {
    panic!(
        "mismatch chunk-proof leaf commitment: expected={:?}, got={:?}",
        CHUNK_EXE_COMMIT, commitment.leaf,
    );
}

```

Figure B.9: `crates/circuits/batch-circuit/src/circuit.rs#L53-L58`

- **Redundant batch index validation on the commitBatchWithBlobProof function.** The commitBatchWithBlobProof function has a redundant check that `batchIndex > euclidForkBatchIndex`. The check appears to exist to allow committing v4 batches at an earlier index than the euclidForkBatchIndex (after this index has been set). However, given that batches need to be committed sequentially, it is not possible to commit a v4 batch before the already-committed v5 batch.

```

// Don't allow to commit version 4 after Euclid upgrade.
// This check is to avoid sequencer committing wrong batch due to human
// error.
// And This check won't introduce much gas overhead (likely less than 100).
if (_version == 4) {
    uint256 euclidForkBatchIndex = initialEuclidBatchIndex;
    if (euclidForkBatchIndex > 0 && batchIndex > euclidForkBatchIndex)
        revert ErrorEuclidForkEnabled();
}

```

Figure B.10: `src/L1/rollup/ScrollChain.sol#319-324`

- **The parallel branch loadMPT could be simplified.** The loadMPT function uses



mutexes in parallel mode to ensure thread safety. However, the code used to ensure that threads are properly synchronized implicitly relies on checking a map entry to safely unlock a mutex without resulting in a runtime error when no parallelism is used. Although the code is correct, it could be easier to understand and review if the expected behavior is grouped in if branches depending on the parallelism mode.

```
for trieIt.Next() {
    if parallel {
        mptLeafMutex.Lock()

        if _, ok := mptLeafMap[string(trieIt.Key)]; ok {
            mptLeafMutex.Unlock()
            break
        }

        mptLeafMap[string(dup(trieIt.Key))] = dup(trieIt.Value)

        if parallel {
            mptLeafMutex.Unlock()
        }

        if parallel && len(mptLeafMap)%10000 == 0 {
            fmt.Println("MPT Accounts Loaded:", len(mptLeafMap))
        }
    }
}
```

Figure B.11: ([go-ethereum/cmd/migration-checker/main.go#187-206](#))

## C. Automated Analysis Tool Configuration

---

We used the following tools to perform automated testing of the codebase.

### C.1. Clippy

The Rust linter **Clippy** can be installed using `rustup` by running the command `rustup component add clippy`. Invoking `cargo clippy -- -W clippy::pedantic` in the root directory of the project runs the tool with the pedantic ruleset.

Clippy results from the regular set of rules should always be fixed.

```
# run clippy in regular and pedantic mode and output to SARIF
cargo clippy --message-format=json | clippy-sarif > clippy.sarif
cargo clippy --message-format=json -- -W clippy::pedantic -Wclippy::or_fun_call |
clippy-sarif > clippy_pedantic.sarif
```

*Figure C.1: The invocation commands used to run Clippy in the codebase*

Converting the output to the SARIF file format (e.g., with **clippy-sarif**) allows easy inspection of the results within an IDE (e.g., using VSCode's **SARIF Explorer** extension).

### C.2. Dylint

**Dylint** is a tool for running Rust lints from dynamic libraries, similar to Clippy. We ran the general and supplementary rulesets against the codebase finding no relevant issues.

```
cargo dylint --no-deps --git https://github.com/trailofbits/dylint --pattern
examples/general -- --message-format=json | clippy-sarif > general.sarif

cargo dylint --no-deps --git https://github.com/trailofbits/dylint --pattern
examples/supplementary -- --message-format=json | clippy-sarif > supplementary.sarif
```

*Figure C.2: The command used to run Dylint on the project*

### C.3. cargo-edit

**cargo-edit** allows developers to quickly find outdated Rust crates. The tool can be installed with the `cargo install cargo-edit` command, and the `cargo upgrade --incompatible --dry-run` command can be used to find outdated crates.

### C.4. cargo-audit

The **cargo-audit** Cargo plugin identifies known vulnerable dependencies in Rust projects. It can be installed using `cargo install cargo-audit`. To run the tool, run `cargo audit` in the crate root directory.

## C.5. zizmor

**zizmor** is a static analysis tool for GitHub Actions. It can be installed using `cargo install --locked zizmor`. We ran zizmor with the following commands:

```
zizmor . --format sarif > zizmor.sarif
zizmor --persona pedantic . --format sarif > zizmor_pedantic.sarif
```

*Figure C.3: Commands used to run zizmor and output to SARIF*

## C.6. CodeQL

We analyzed the Go codebase with public and Trail of Bits private CodeQL queries.

```
# Create the go database
codeql database create codeql.db --language=go

# Run all go queries
codeql database analyze codeql.db --additional-packs ~/.codeql/codeql-repo
--format=sarif-latest --output=codeql_tob_all.sarif -- tob-go-all
```

*Figure C.4: Commands used to run CodeQL*

## D. Fix Review Results

When undertaking a fix review, Trail of Bits reviews the fixes implemented for issues identified in the original report. This work involves a review of specific areas of the source code and system configuration, not comprehensive analysis of the system.

On March 24, 2025, Trail of Bits reviewed the fixes and mitigations implemented by the Scroll team for the issues identified in this report. We reviewed each fix to determine its effectiveness in resolving the associated issue.

In summary, of the eight issues described in this report, Scroll has resolved seven issues and has not resolved one issue. The unresolved finding relates to missing unit tests in the migration code. For additional information, please see the Detailed Fix Review Results below.

ID	Title	Severity	Status
1	The aggregation circuit does not validate the size of inner proofs	High	Resolved
2	Insufficient validation of chunk information and blob bytes in the batch circuit	High	Resolved
3	MPT migration code lacks unit tests	Informational	Unresolved
4	Incorrect data present on BundleInfo	Informational	Resolved
5	Lax parsing of chunk data payload	Informational	Resolved
6	Docker release action is vulnerable to cache poisoning	Low	Resolved
7	Unpinned external GitHub CI/CD action versions	Low	Resolved
8	Potential credential persistence in artifacts and stale GitHub action	Informational	Resolved

## Detailed Fix Review Results

### **TOB-SCREUC-1: The aggregation circuit does not validate the size of inner proofs**

Resolved in [PR#85](#) by using the `zip_eq` iterator, which checks that its inputs have the same length.

### **TOB-SCREUC-2: Insufficient validation of chunk information and blob bytes in the batch circuit validate**

Resolved in [PR#85](#) by using the `zip_eq` iterator, which checks that its inputs have the same length.

### **TOB-SCREUC-3: MPT migration code does not have unit tests**

Unresolved. The Scroll team decided not to add any unit tests to the migration code, providing the following context:

*[...] while we consider the current manual testing sufficient, we will consider adding additional unit tests before the upgrade.*

### **TOB-SCREUC-4: Incorrect data present on BundleInfo**

Resolved in the [51847 commit](#) by setting the public input to the correct value.

### **TOB-SCREUC-5: Lax parsing of chunk data payload**

Resolved in commits [9d9ed](#) and [c58cc](#) by ensuring that the unused chunk lengths are always zero.

### **TOB-SCREUC-6: Docker release action is vulnerable to cache poisoning**

Resolved in [PR#1137](#) and [PR#83](#) by disabling binary cache persistence in release GitHub actions.

### **TOB-SCREUC-7: Unpinned external GitHub CI/CD action versions**

Resolved in [PR#1138](#) and [PR#84](#) by pinning all GitHub action versions to hashes, except those from the actions repository (e.g., `actions/checkout`).

### **TOB-SCREUC-8: Potential credential persistence in artifacts and stale GitHub action**

Resolved in [PR#1139](#) and [PR#85](#) by disabling credential persistence when calling the `actions/checkout` GitHub action.

## E. Fix Review Status Categories

---

The following table describes the statuses used to indicate whether an issue has been sufficiently addressed.

Fix Status	
Status	Description
Undetermined	The status of the issue was not determined during this engagement.
Unresolved	The issue persists and has not been resolved.
Partially Resolved	The issue persists but has been partially resolved.
Resolved	The issue has been sufficiently resolved.

# About Trail of Bits

---

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries and government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on X and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact> or email us at [info@trailofbits.com](mailto:info@trailofbits.com).

## Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

[info@trailofbits.com](mailto:info@trailofbits.com)

# Notices and Remarks

---

## Copyright and Distribution

© 2025 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

Trail of Bits considers this report public information; it is licensed to Scroll under the terms of the project statement of work and has been made public at Scroll's request. Material within this report may not be reproduced or distributed in part or in whole without Trail of Bits' express written permission.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through sources other than that page may have been modified and should not be considered authentic.

## Test Coverage Disclaimer

Trail of Bits performed all activities associated with this project in accordance with a statement of work and an agreed-upon project plan.

Security assessment projects are time-boxed and often rely on information provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test software controls and security properties. These techniques augment our manual security review work, but each has its limitations. For example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. A project's time and resource constraints also limit their use.